

**Elements Of Data Science - F2022**

# **Week 6: Intro to Machine Learning Models Continued**

**10/12/2021**

# TODOs

- Readings:
  - PDSH 05.03 Hyperparameters and Model Validation
  - Recommended: PML Chapter 6 (Except for Pipelines) and sklearn - model selection
  - Reference: PML Chapter Chap 3, 7, and sklearn - supervised learning
- Quiz 6, Due Tues Oct 18th, 11:59pm
- HW2 out end of *next* week
- **Midterm**
  - Review sheet in github repo
  - Online via gradescope, open-book, open-note, open-python
  - Released *Wednesday Oct 19th 11:59pm*
  - Due *Friday Oct 21st 11:59pm ET*
  - Have maximum of 24hrs after starting to finish
  - 30-40 questions (fill in the blank/multiple choice/short answer)
  - Questions asked/answered **privately** via Ed

# Don't Use Old Course Materials

- Do not use old course materials including quizzes, homeworks or exams
- *If there is reason to believe you are submitting work copied from an old course, you will get a 0 on that assignment*
- Using the current semester slides and readings is acceptable (though I recommend not just copying)

# Quiz 4 Common Mistakes

- include instructions with submission
- `.info()` not `.info`
- Pandas `.sample()` default `n=1`: need to set `n=` or `frac=`

I'll start taking off points for errors I point out in class.

# Today

- Review Linear Models
- Distance Based: kNN
- Tree Based: Decision Tree
- Ensembles: Bagging, Boosting, Stacking
- Multiclass/Multilabel and One Vs. Rest Classification
- Model Review

Questions?

# Environment Setup

# Environment Setup

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

from mlxtend.plotting import plot_decision_regions

from sklearn.linear_model import LinearRegression, LogisticRegression

sns.set_style('darkgrid')
%matplotlib inline
```

# Environment Setup

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

from mlxtend.plotting import plot_decision_regions

from sklearn.linear_model import LinearRegression, LogisticRegression

sns.set_style('darkgrid')
%matplotlib inline
```

```
In [2]: def my_plot_decision_regions(X, y, model, figsize=(8,8), ax=None):
        '''Plot classifier decision regions, classification predictions and training data'''
        if not ax:
            fig, ax = plt.subplots(1,1,figsize=figsize)
            # use mlxtend plot_decision_regions
            model = model.fit(X.values, y.values)
            plot_decision_regions(X.values, y.values, model, ax=ax)
            ax.set_xlabel(X.columns[0]); ax.set_ylabel(X.columns[1]);

def my_plot_regression(X, y, model, label='yhat', figsize=(8,8), ax=None):
    '''Plot regression predictions and training data'''
    # generate test data and make predictions
    X_test = np.linspace(X.iloc[:,0].min(), X.iloc[:,0].max(), 1000).reshape(-1,1)
    model = model.fit(X.values, y.values)
    y_hat = model.predict(X_test)
    fig, ax = plt.subplots(1,1,figsize=figsize)
    ax.scatter(X, y, s=20, edgecolor="black", c="darkorange", label="data")
    ax.plot(X_test, y_hat, color="cornflowerblue", label=label, linewidth=2)
    ax.set_xlabel(X.columns[0]); ax.set_ylabel(y.name); ax.legend();
```



# Linear Models (Review)

- Simple/Multiple Linear Regression
- Logistic Regression
- SVM
- Perceptron, Multi-Layer Perceptron

# Wine as Binary Classification

# Wine as Binary Classification

```
In [3]: df_wine = pd.read_csv('../data/wine_dataset.csv', usecols=['alcohol', 'ash', 'proline', 'hue', 'class'])
# rename 'class' as 'target', since class is a reserved python word
df_wine = df_wine.rename({'class': 'target'}, axis=1)

df_wine_2c = df_wine[df_wine.target < 2] # only keep classes 0 and 1

X_2c = df_wine_2c[['proline', 'hue']]
y_2c = df_wine_2c['target']

zscore = lambda x: (x - x.mean()) / x.std()

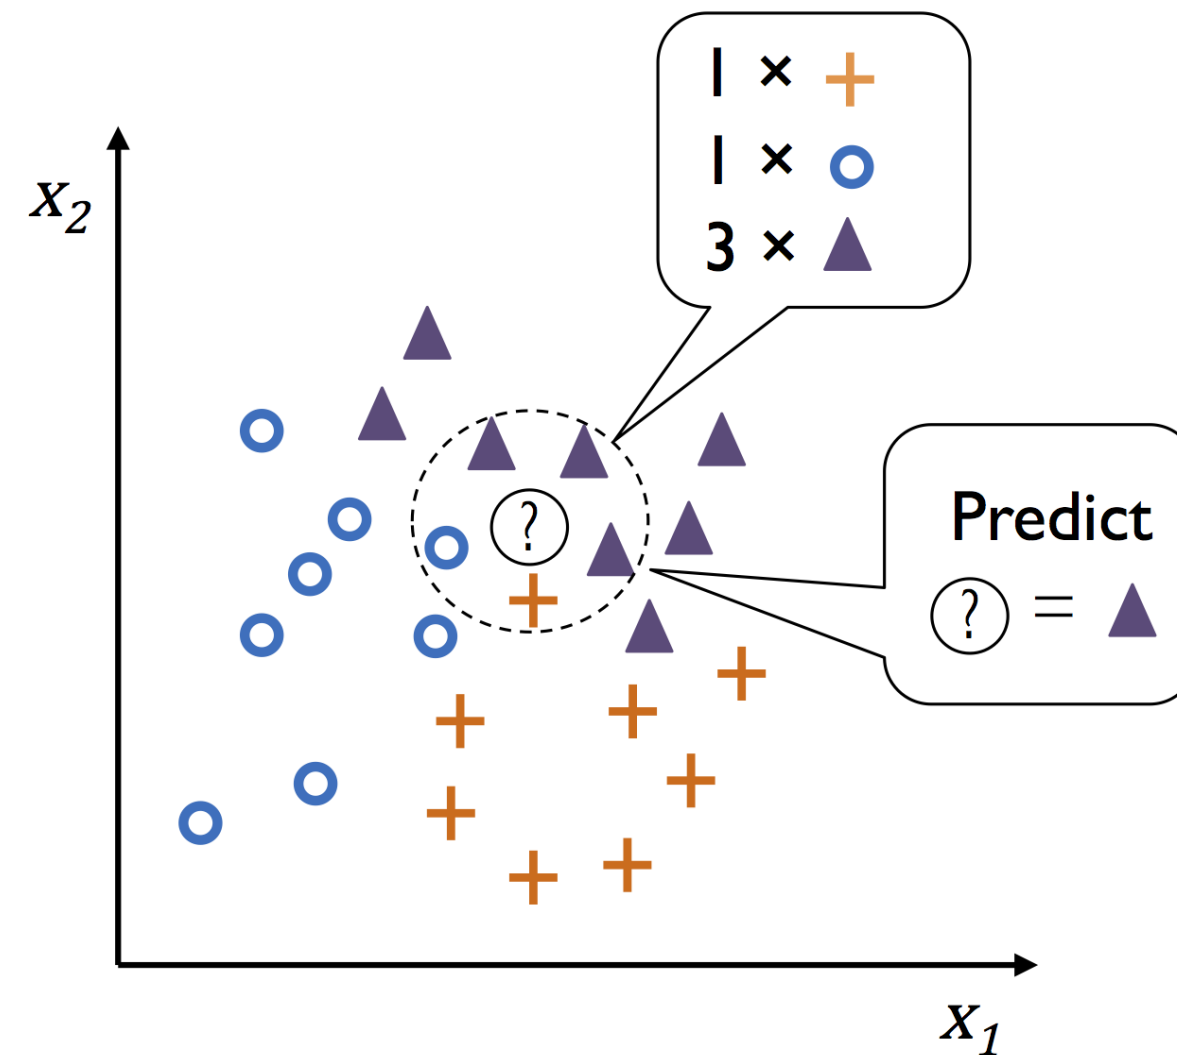
X_2c_zscore = X_2c.apply(zscore, axis=0)
alcohol_2c_zscore = zscore(df_wine_2c.alcohol)

y_2c.value_counts().sort_index()
```

```
Out[3]: 0    59
        1    71
        Name: target, dtype: int64
```

# Distance Based: k-Nearest Neighbor (kNN)

- What category do most of the  $k$  nearest neighbors belong to?



From PML

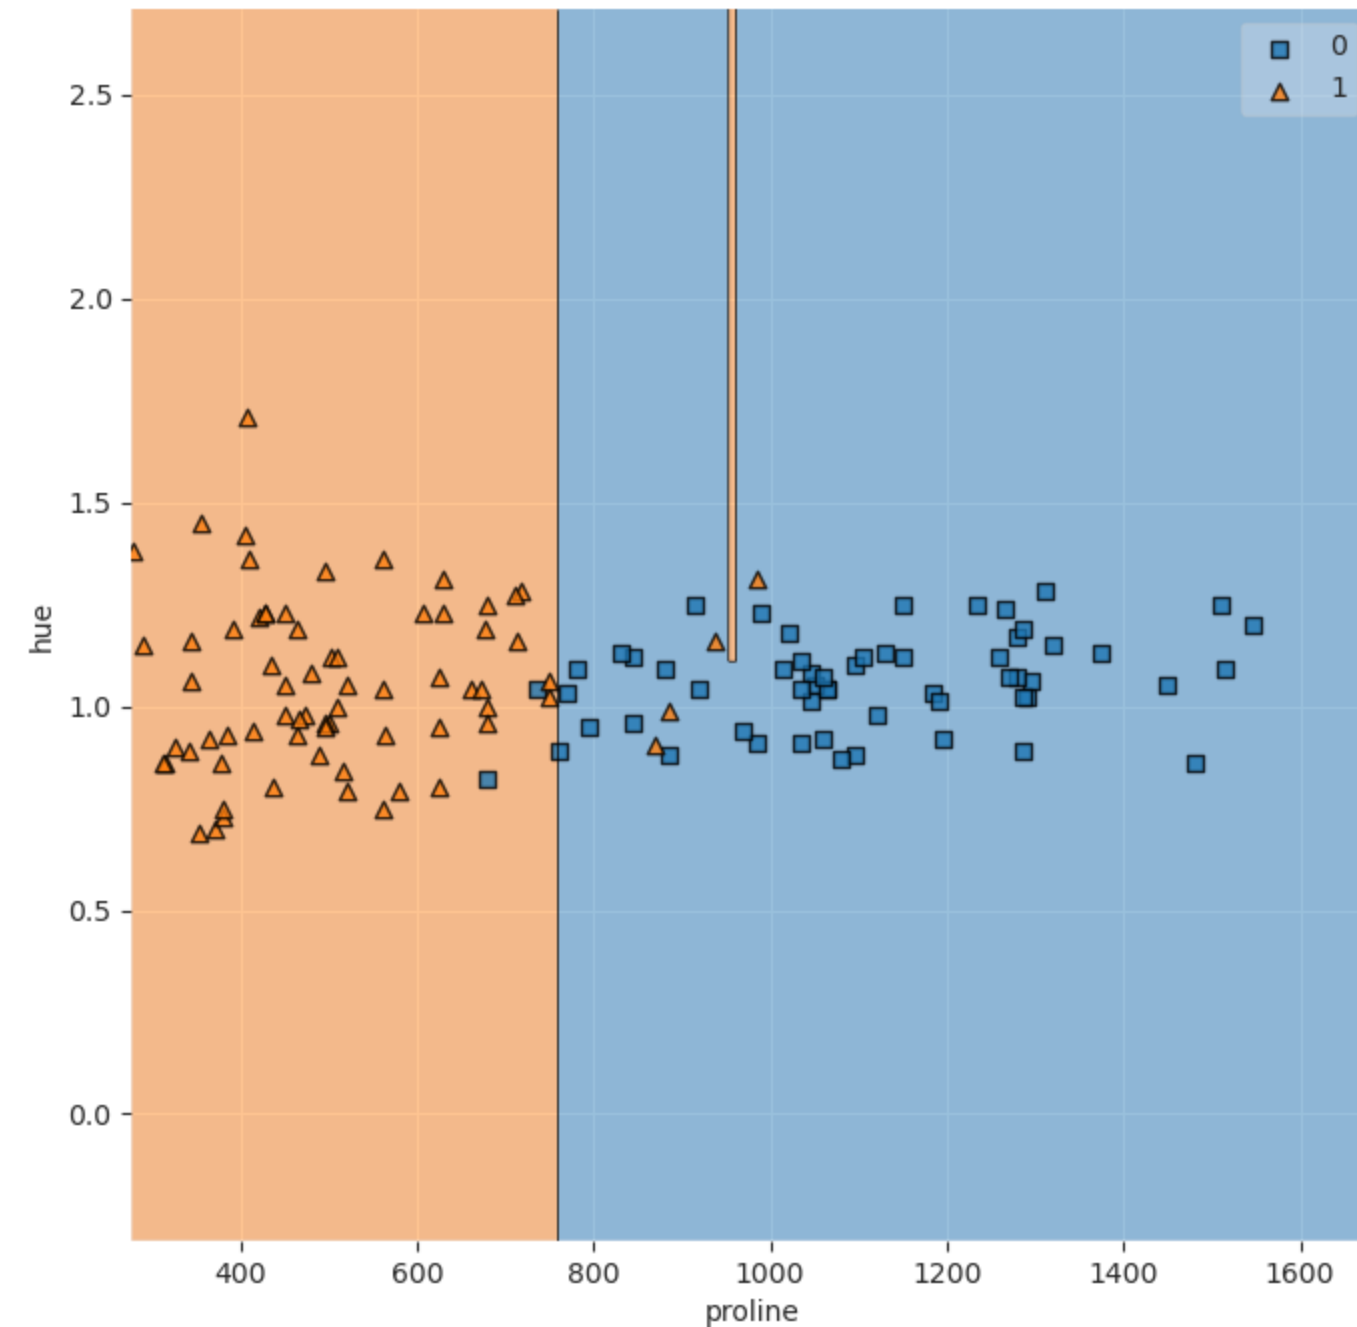
# KNN in sklearn

# KNN in sklearn

```
In [4]: from sklearn.neighbors import KNeighborsClassifier
```

```
knn = KNeighborsClassifier(n_neighbors=3)  
knn.fit(X_2c, y_2c)
```

```
my_plot_decision_regions(X_2c, y_2c, knn)
```

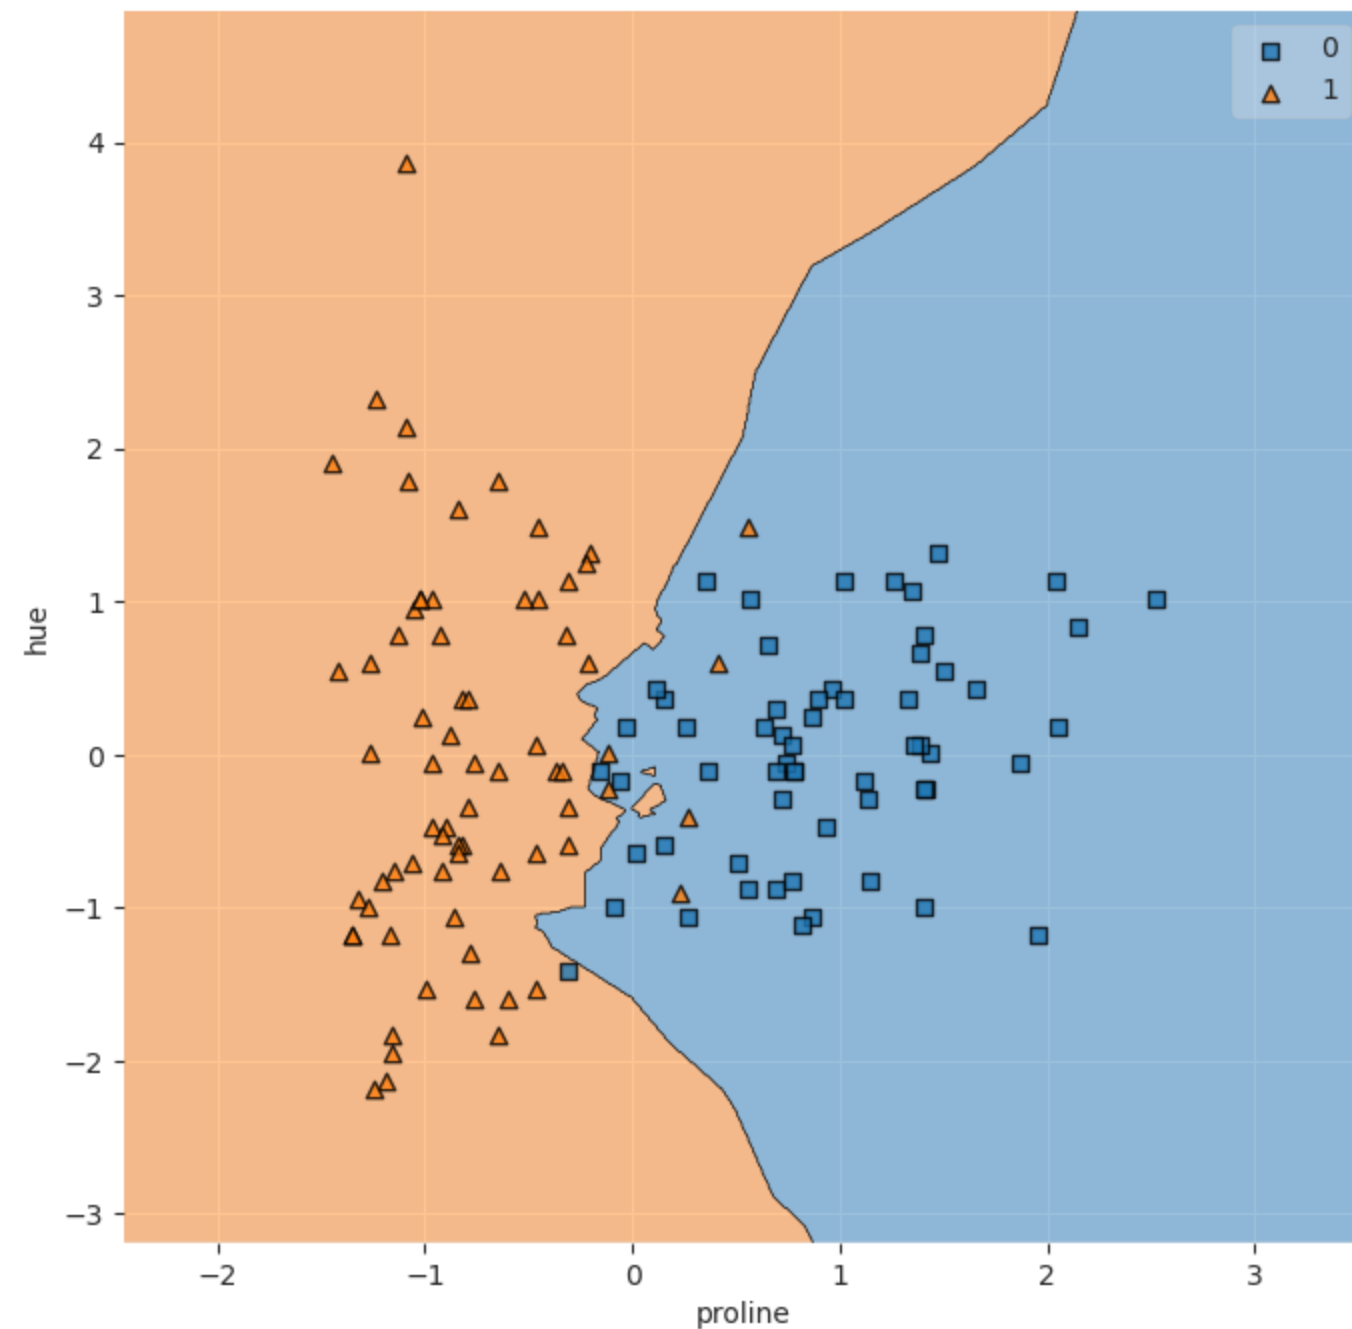


# Effects of Standardization on Distance Based Methods

# Effects of Standardization on Distance Based Methods

```
In [5]: knn_z = KNeighborsClassifier(n_neighbors=3)
knn_z.fit(X_2c_zscore, y_2c)

my_plot_decision_regions(X_2c_zscore, y_2c, knn_z)
```





# Curse of Dimensionality

The more dimensions, the less likely points are "close" to each other.

# Curse of Dimensionality

The more dimensions, the less likely points are "close" to each other.

```
In [6]: # From Data Science From Scratch by Joel Grus

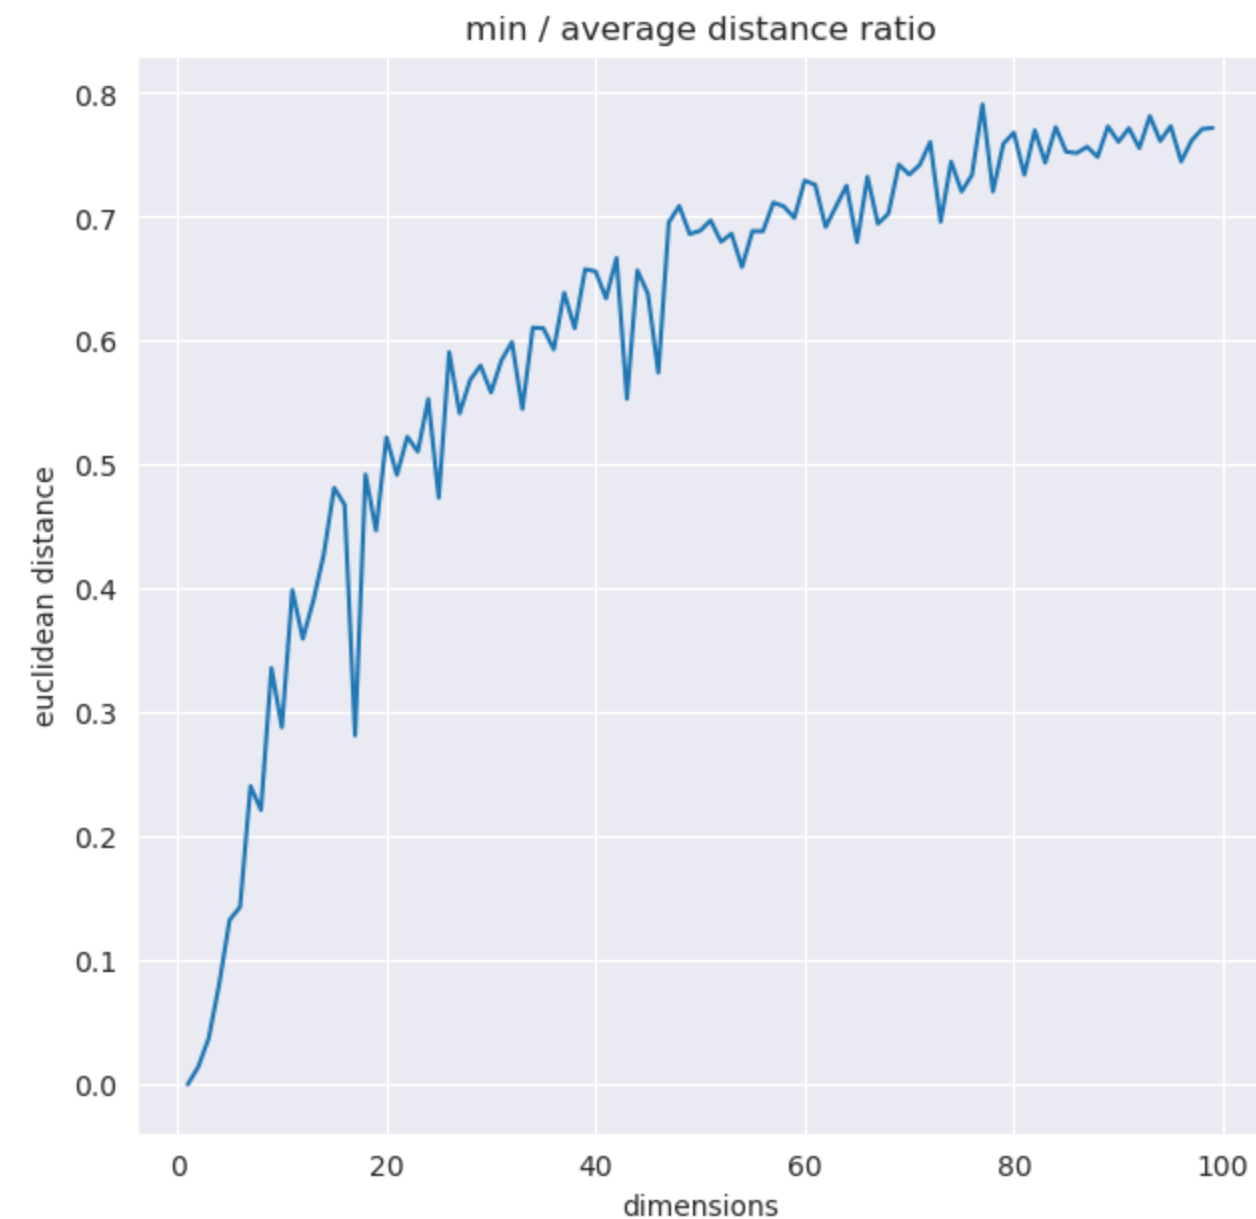
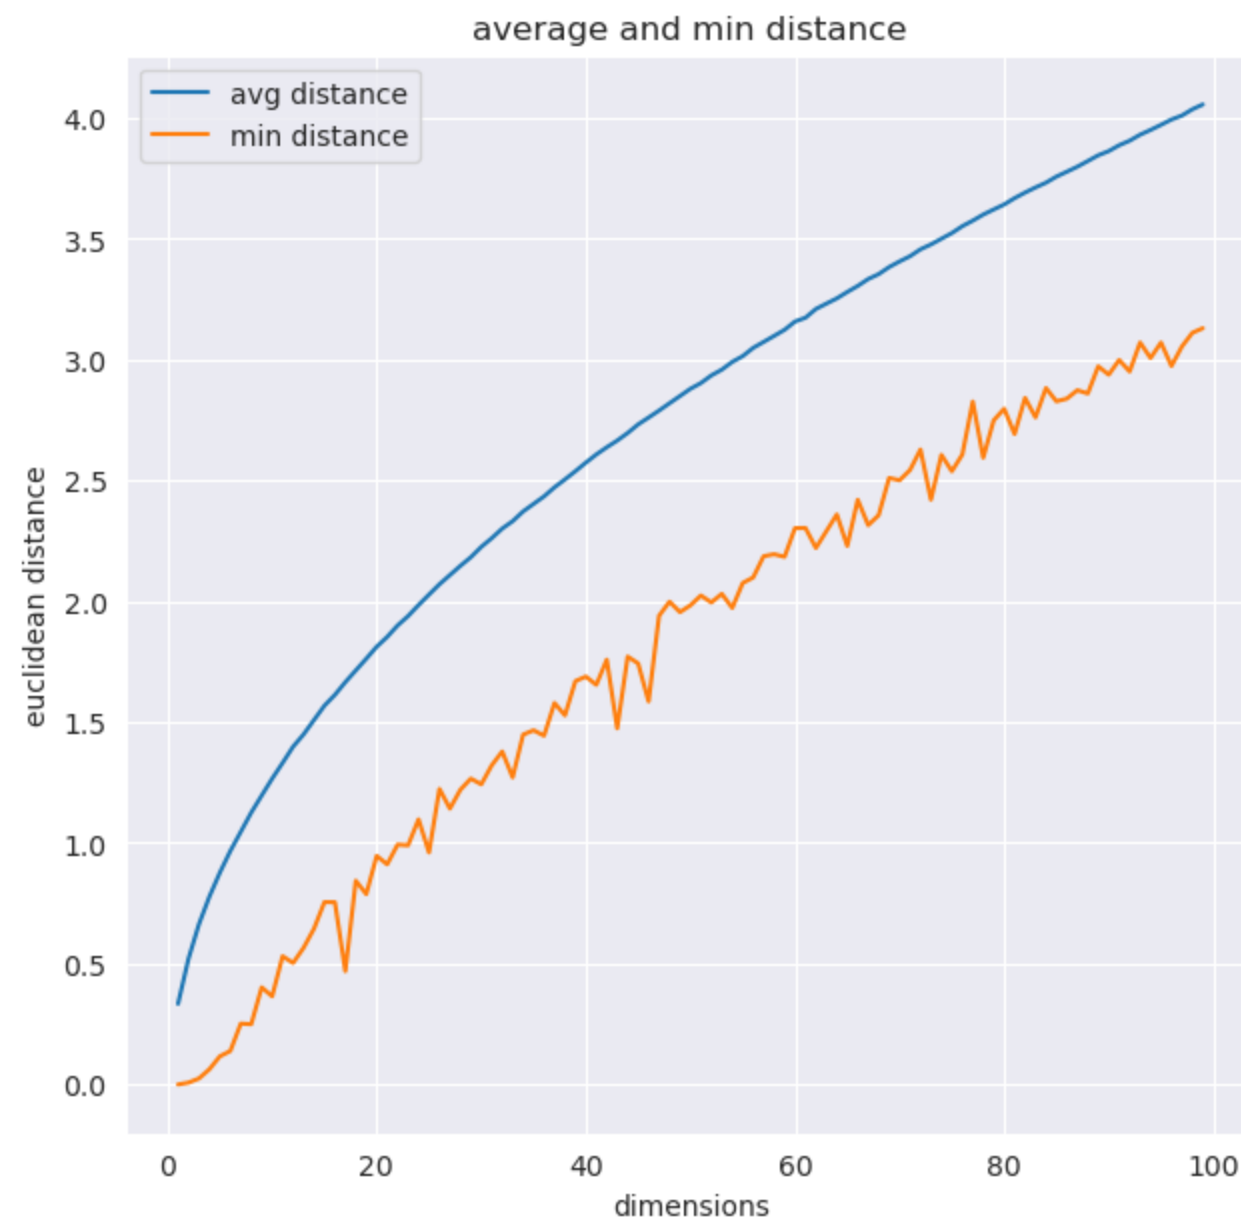
def random_distances(dim,num_pairs=10_000):
    return np.sqrt(np.square(np.random.rand(num_pairs,dim) - np.random.rand(num_pairs,dim)).sum(axis=1))

# calculate average and minimum distance for 1 to 100 dimensions
dimensions = range(1,100)
avg_distances = []
min_distances = []
min_avg_ratio = []
np.random.seed(0)
for d in dimensions:
    distances = random_distances(d)
    avg_distances.append(distances.mean())
    min_distances.append(distances.min())
    min_avg_ratio.append(distances.min() / distances.mean())
```

# Curse of Dimensionality Cont.

# Curse of Dimensionality Cont.

```
In [7]: fig, ax = plt.subplots(1, 2, figsize=(16, 7))
ax[0].plot(dimensions, avg_distances, label='avg distance');
ax[0].plot(dimensions, min_distances, label='min distance');
ax[0].legend()
ax[0].set_title('average and min distance'); ax[0].set_xlabel('dimensions'); ax[0].set_ylabel('euclidean distance');
ax[1].plot(dimensions, min_avg_ratio)
ax[1].set_title('min / average distance ratio'); ax[1].set_xlabel('dimensions'); ax[1].set_ylabel('euclidean distance');
```



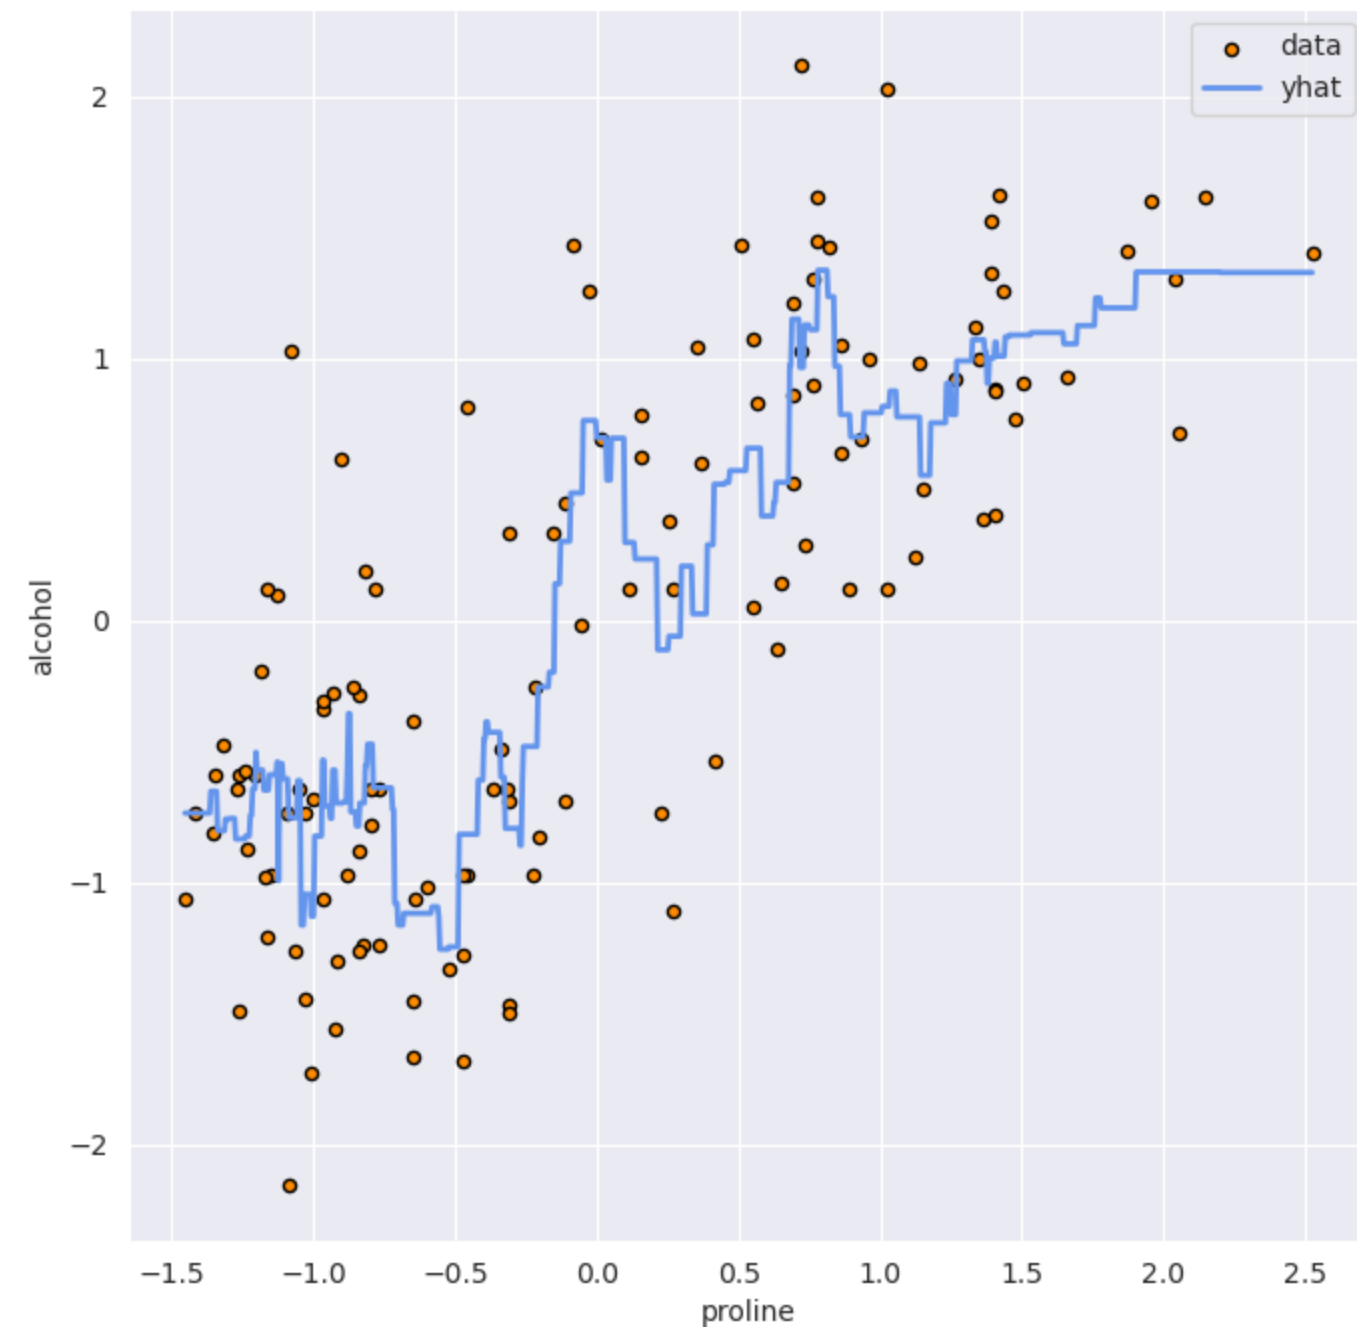
# Regression with kNN

# Regression with kNN

```
In [8]: from sklearn.neighbors import KNeighborsRegressor

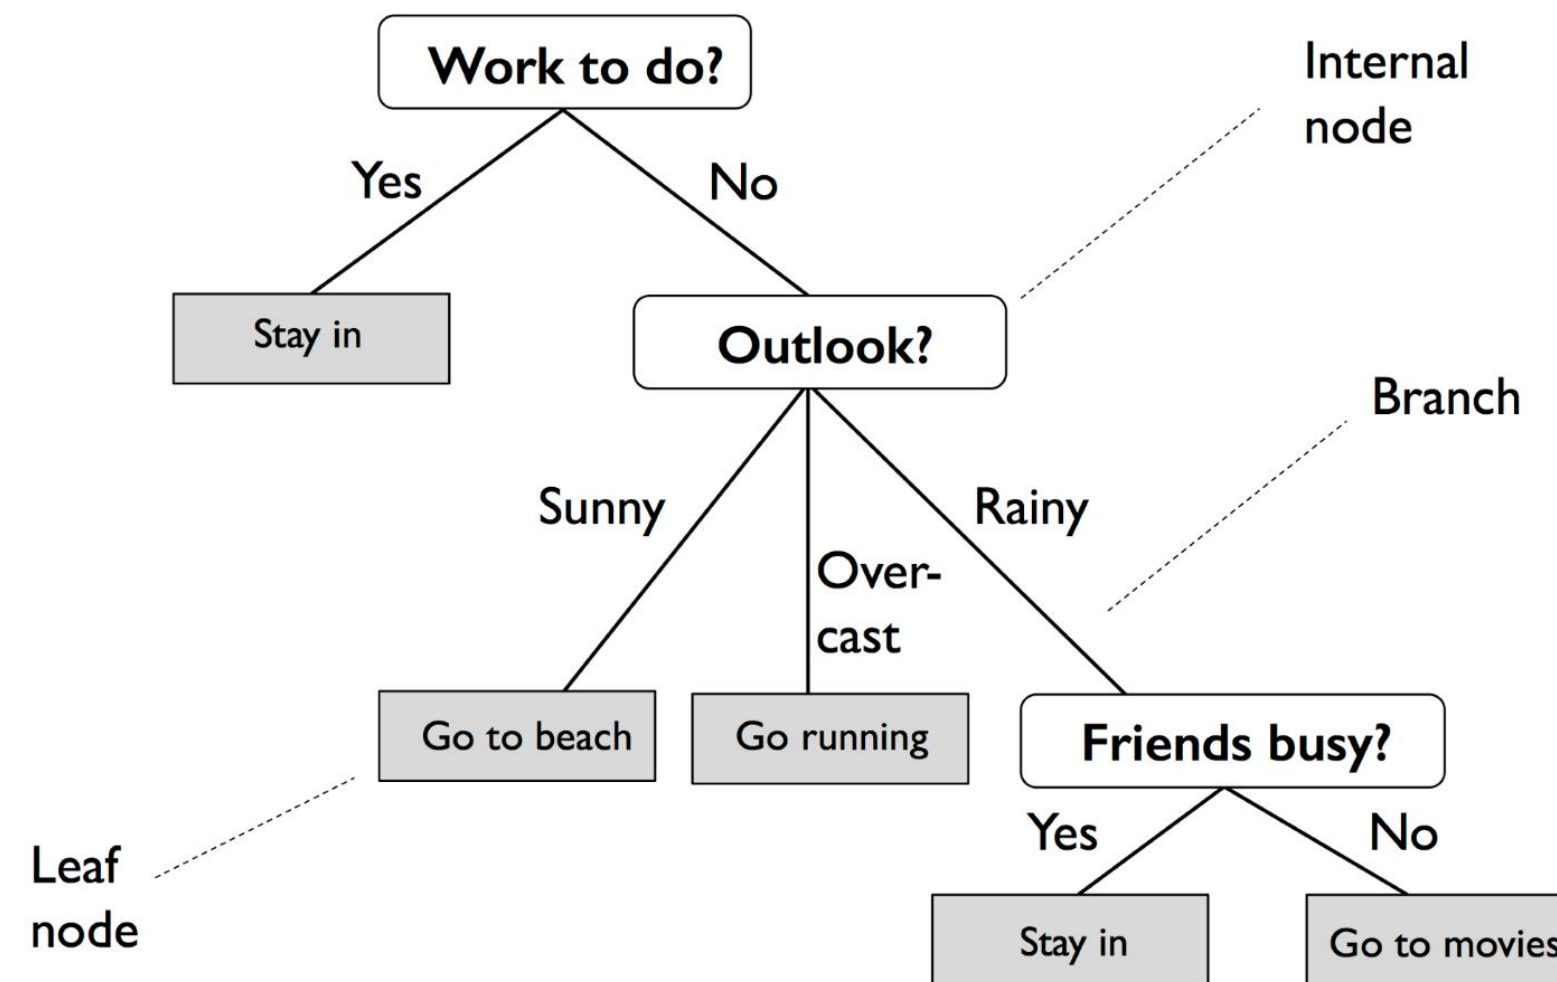
knnr = KNeighborsRegressor(n_neighbors=5)
knnr.fit(X_2c_zscore[['proline']], alcohol_2c_zscore)

my_plot_regression(X_2c_zscore[['proline']], alcohol_2c_zscore, knnr)
```



# Decision Tree

- What answer does a series of yes/no questions lead us to?



From PML

# Decision Tree Classifier in sklearn

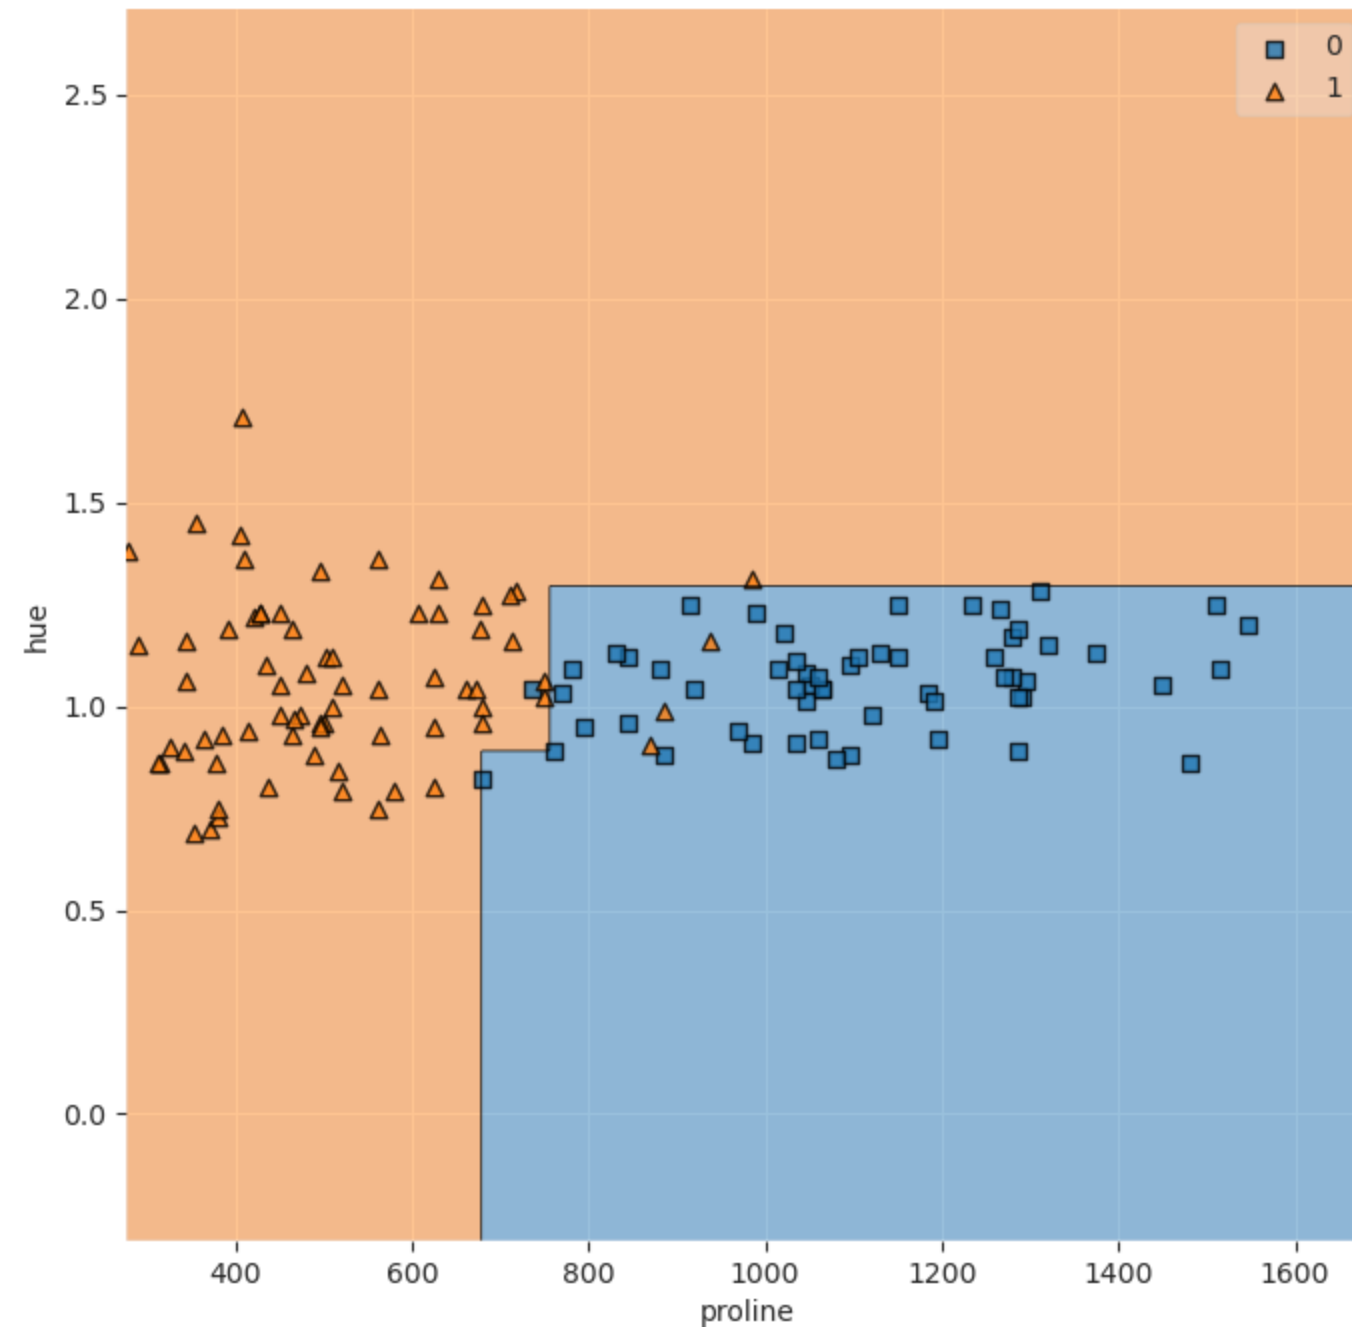


# Decision Tree Classifier in sklearn

```
In [10]: from sklearn.tree import DecisionTreeClassifier

dtc_md3 = DecisionTreeClassifier(max_depth=3) # max_depth: max number of questions
dtc_md3.fit(X_2c,y_2c)

my_plot_decision_regions(X_2c,y_2c,dtc_md3)
```

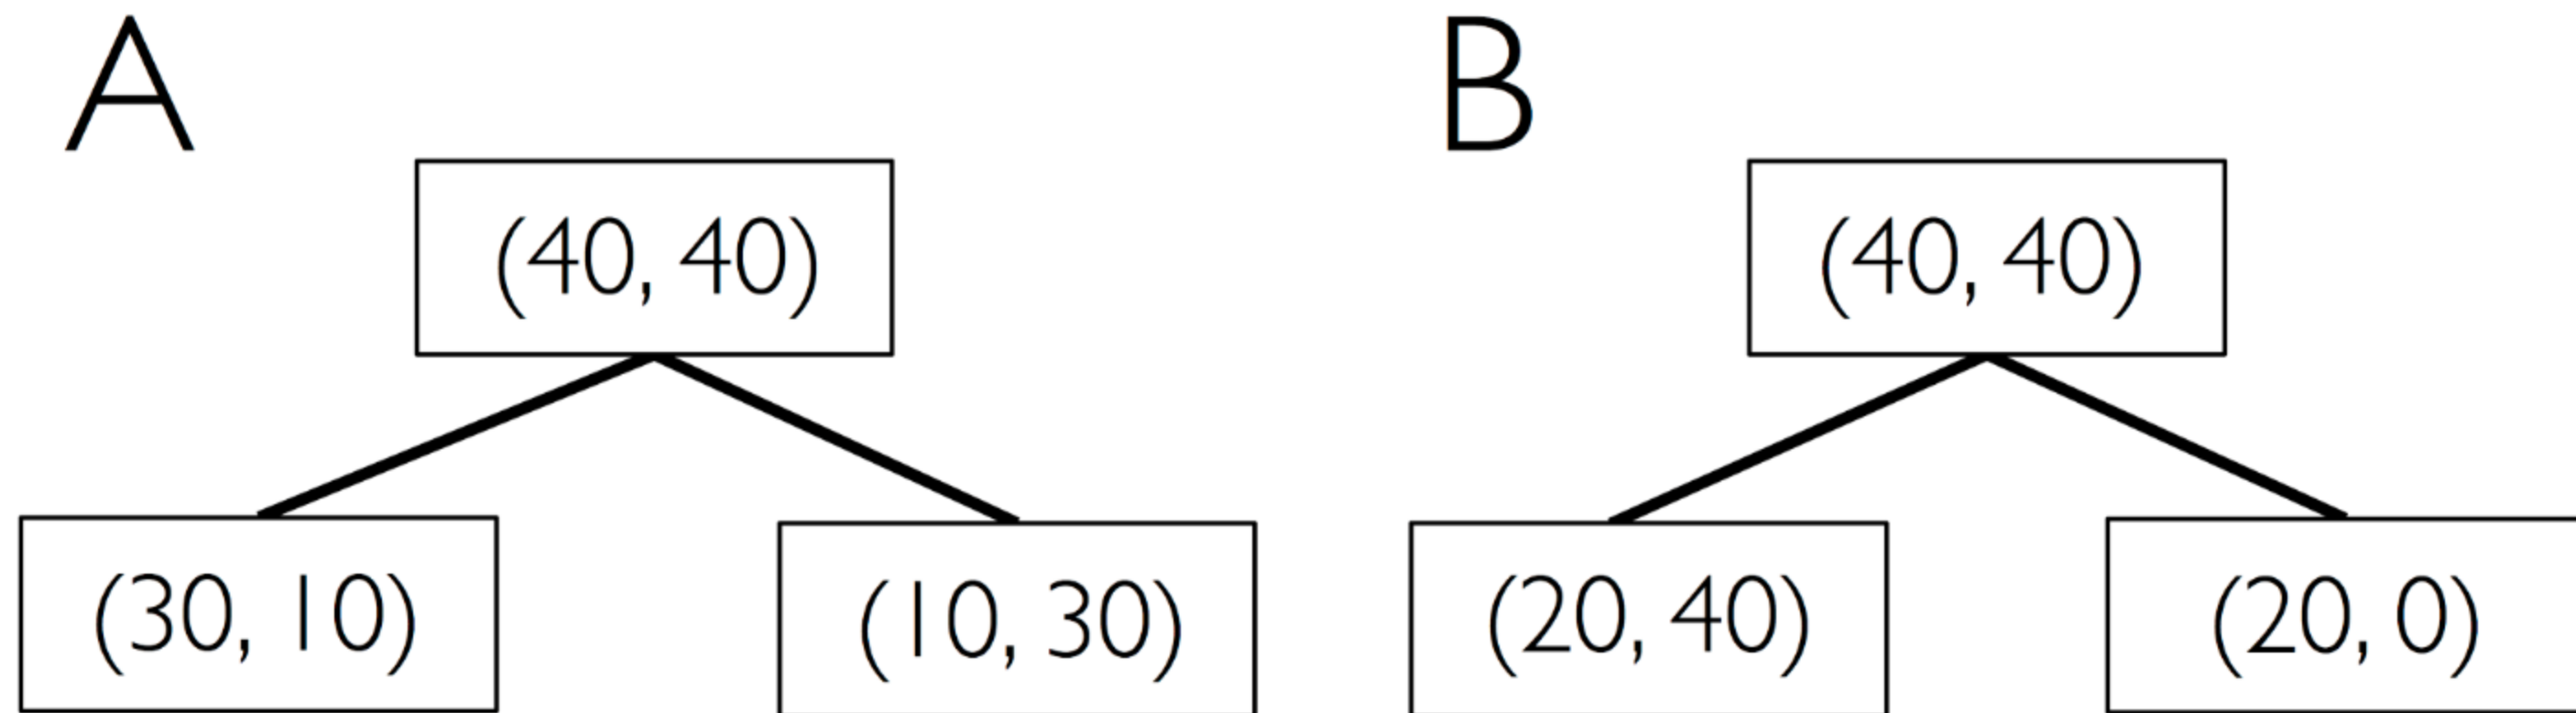


# Building a Decision Tree

- How to decide which question to choose (eg. Should I choose question A or B)?
- **Reduce Impurity**

# Building a Decision Tree

- How to decide which question to choose (eg. Should I choose question A or B)?
- **Reduce Impurity**



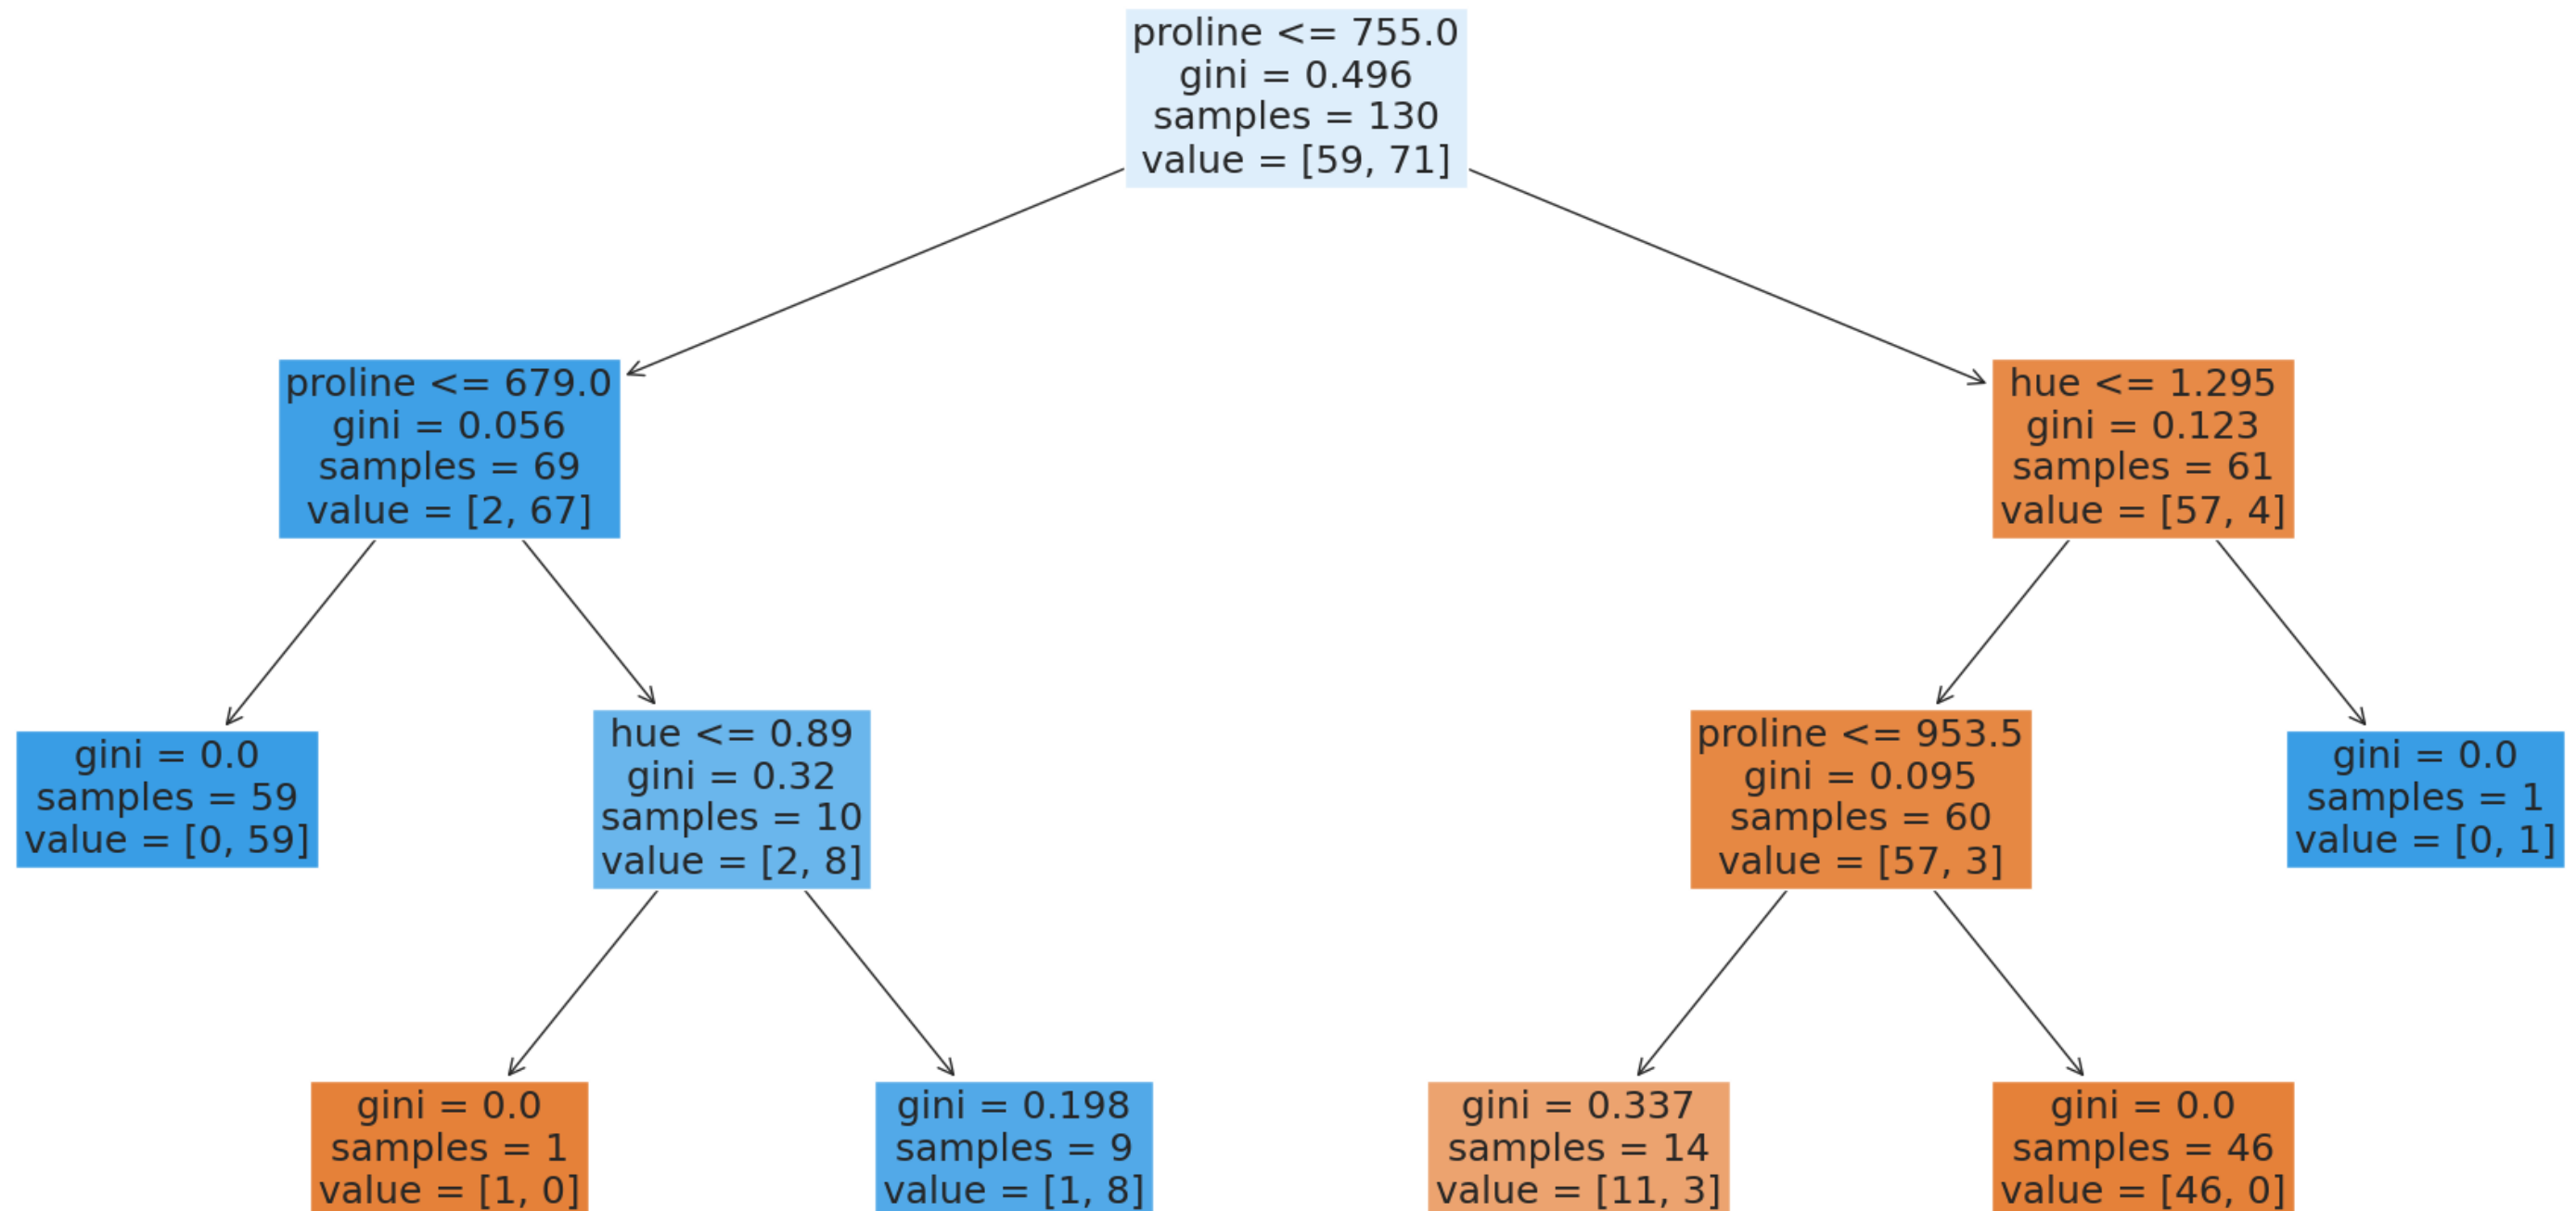
From PML

- Information Gain: Tie, *Gini*: B, Entropy: B

# Plot Learned Decision Tree Using sklearn

# Plot Learned Decision Tree Using sklearn

```
In [12]: from sklearn.tree import plot_tree
fig, ax = plt.subplots(1, 1, figsize=(24, 12))
plot_tree(dtc_md3, ax=ax, fontsize=18, feature_names=X_2c.columns, filled=True);
```

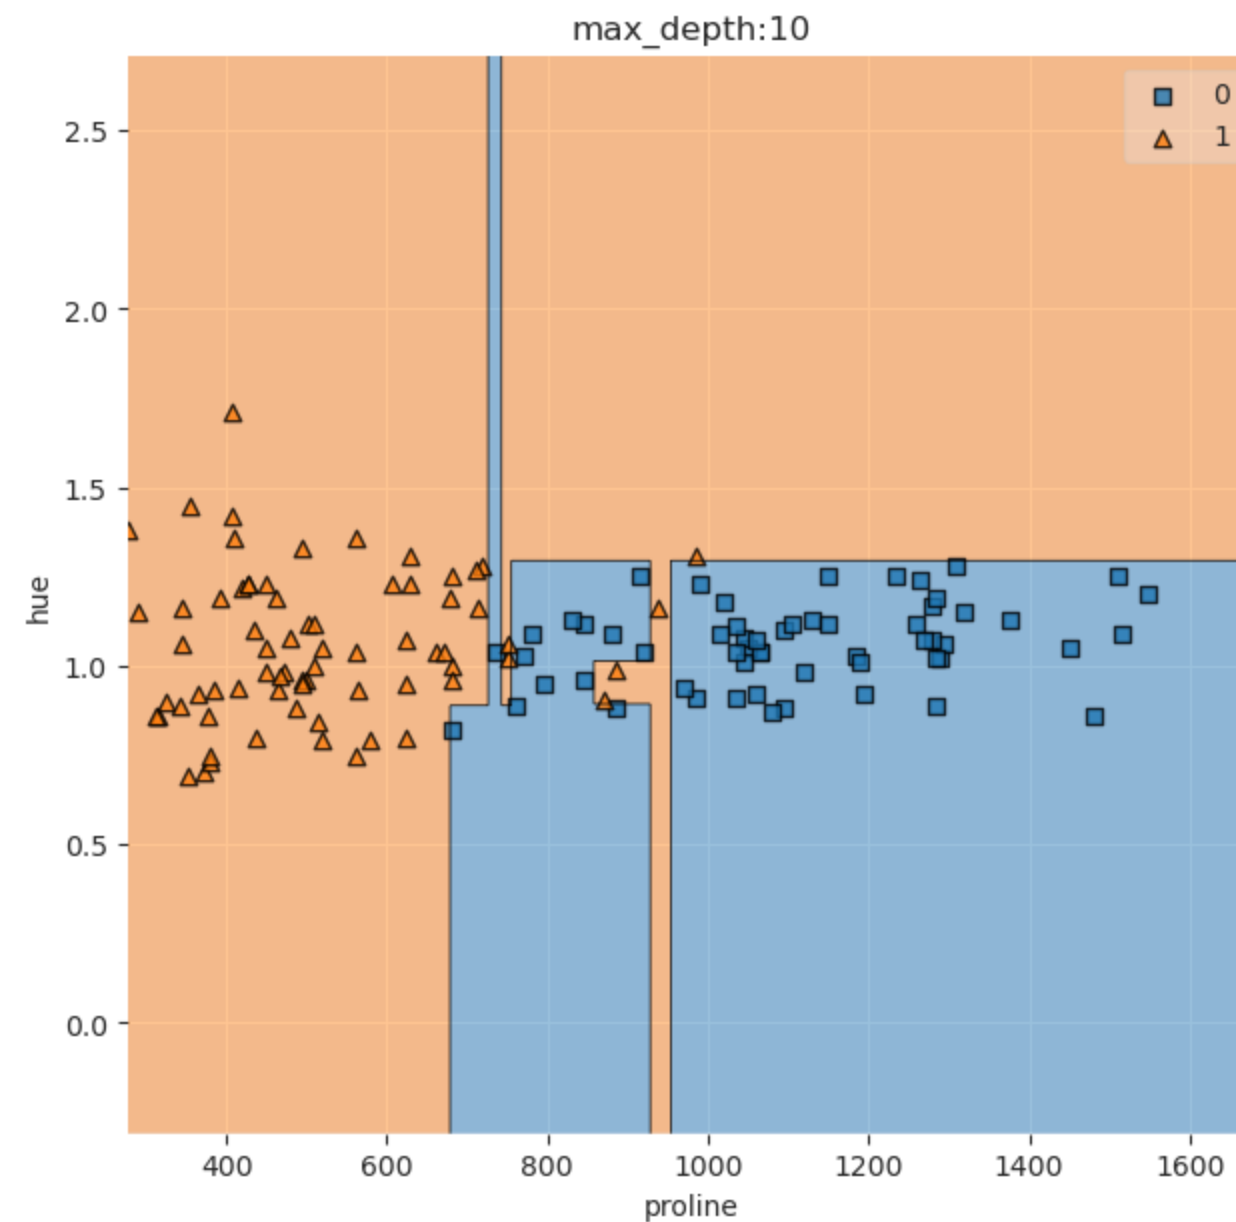
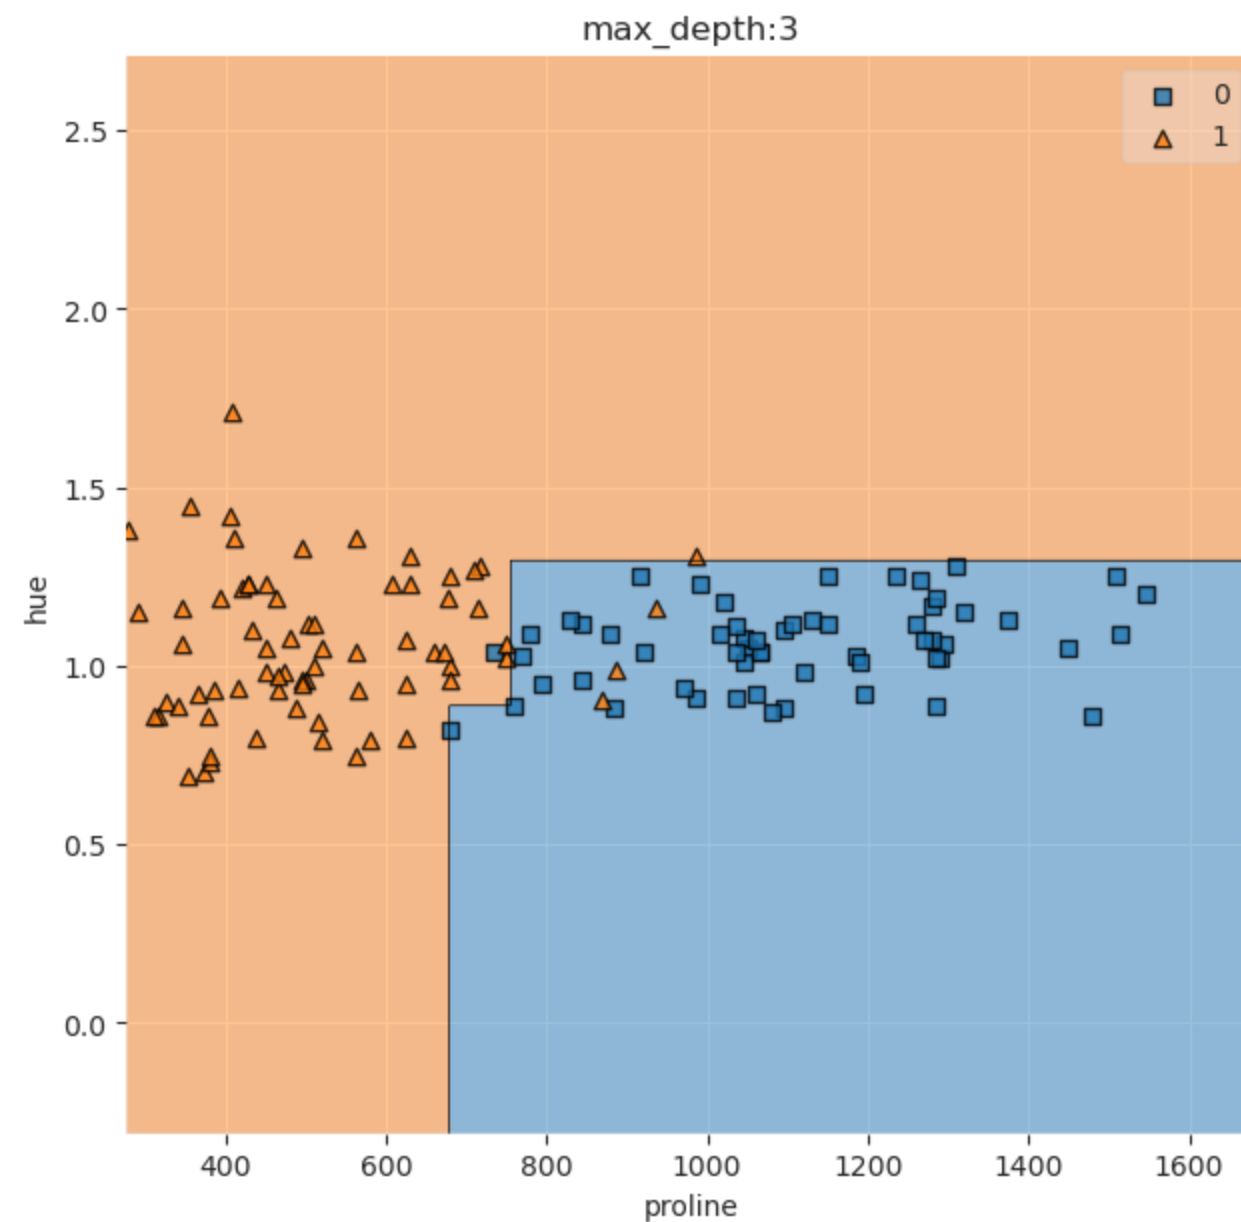


# Decision Tree: Increase Maximum Depth

# Decision Tree: Increase Maximum Depth

```
In [13]: dtc_md10 = DecisionTreeClassifier(max_depth=10)
dtc_md10.fit(X_2c,y_2c)

fig,ax = plt.subplots(1,2,figsize=(16,7))
my_plot_decision_regions(X_2c, y_2c, model=dtc_md3, ax=ax[0]);
my_plot_decision_regions(X_2c, y_2c, model=dtc_md10, ax=ax[1]);
ax[0].set_title('max_depth:3');ax[1].set_title('max_depth:10');
```



# Plot Learned Decision Tree Using sklearn

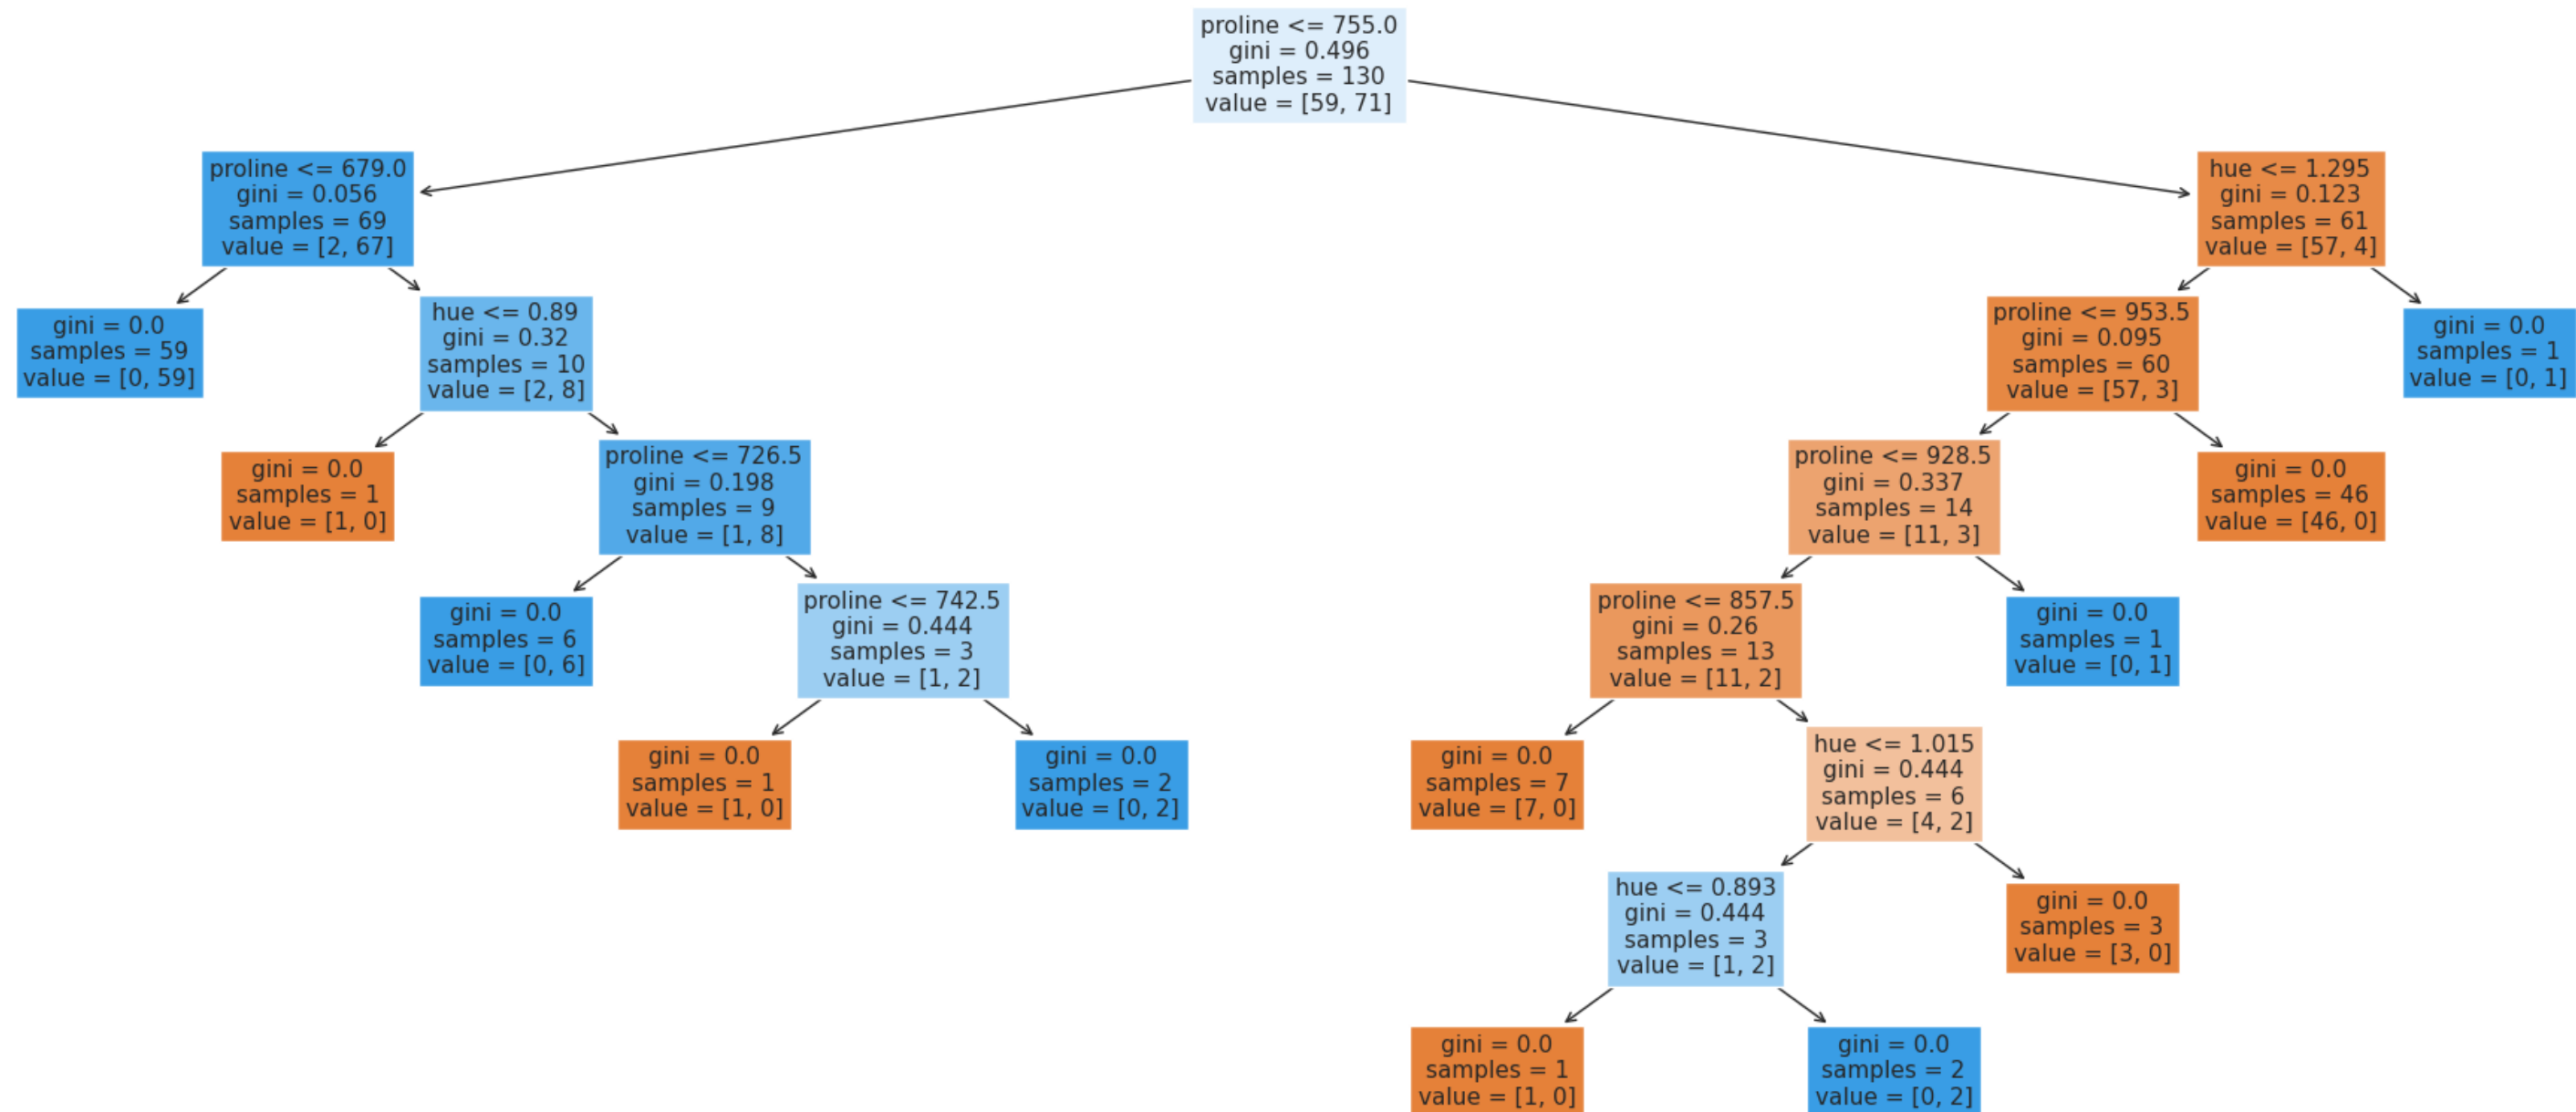
- For tree with max\_depth=10



# Plot Learned Decision Tree Using sklearn

- For tree with max\_depth=10

```
In [14]: fig, ax = plt.subplots(1, 1, figsize=(24, 10))
plot_tree(dtc_md10, ax=ax, fontsize=11, feature_names=X_2c.columns, filled=True);
```



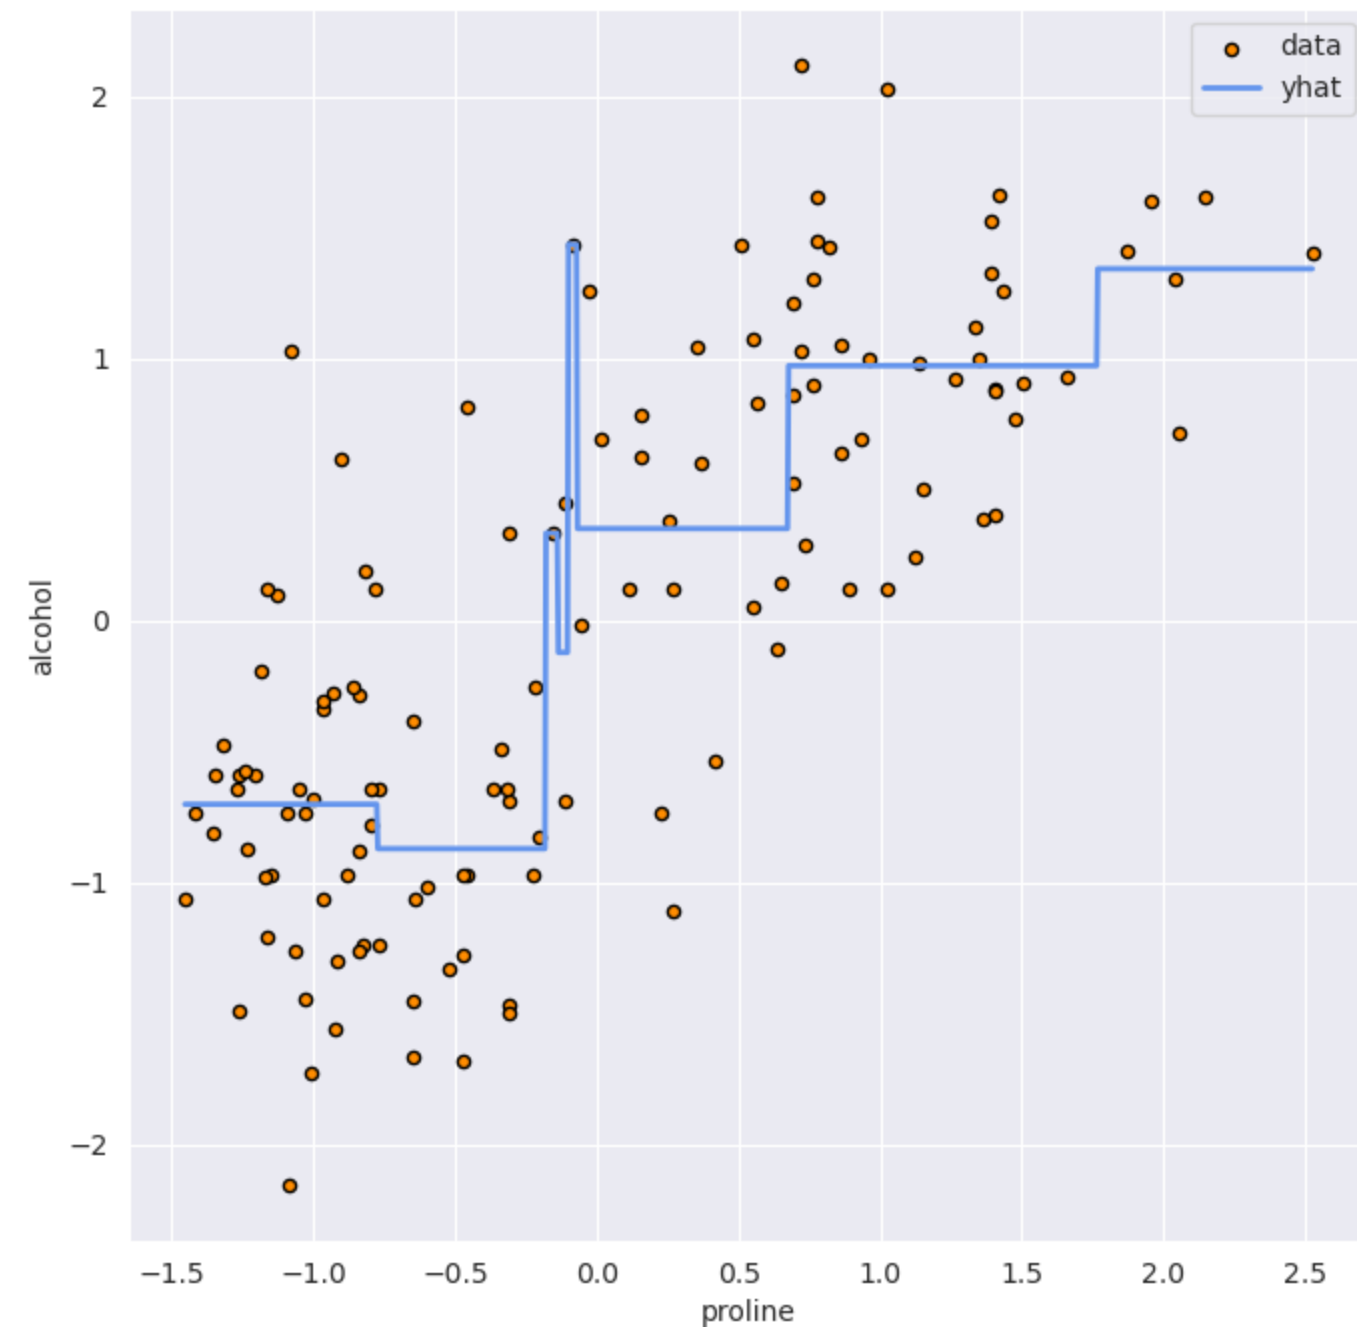
# Regression with Decision Trees

# Regression with Decision Trees

```
In [15]: from sklearn.tree import DecisionTreeRegressor

dtr = DecisionTreeRegressor(max_depth=3)
dtr.fit(X_2c_zscore[['proline']], alcohol_2c_zscore)

my_plot_regression(X_2c_zscore[['proline']], alcohol_2c_zscore, dtr)
```



# Ensemble Methods

- "Wisdom of the crowd"
- Can often achieve better performance with collection of learners
- Often use shallow trees as base learners

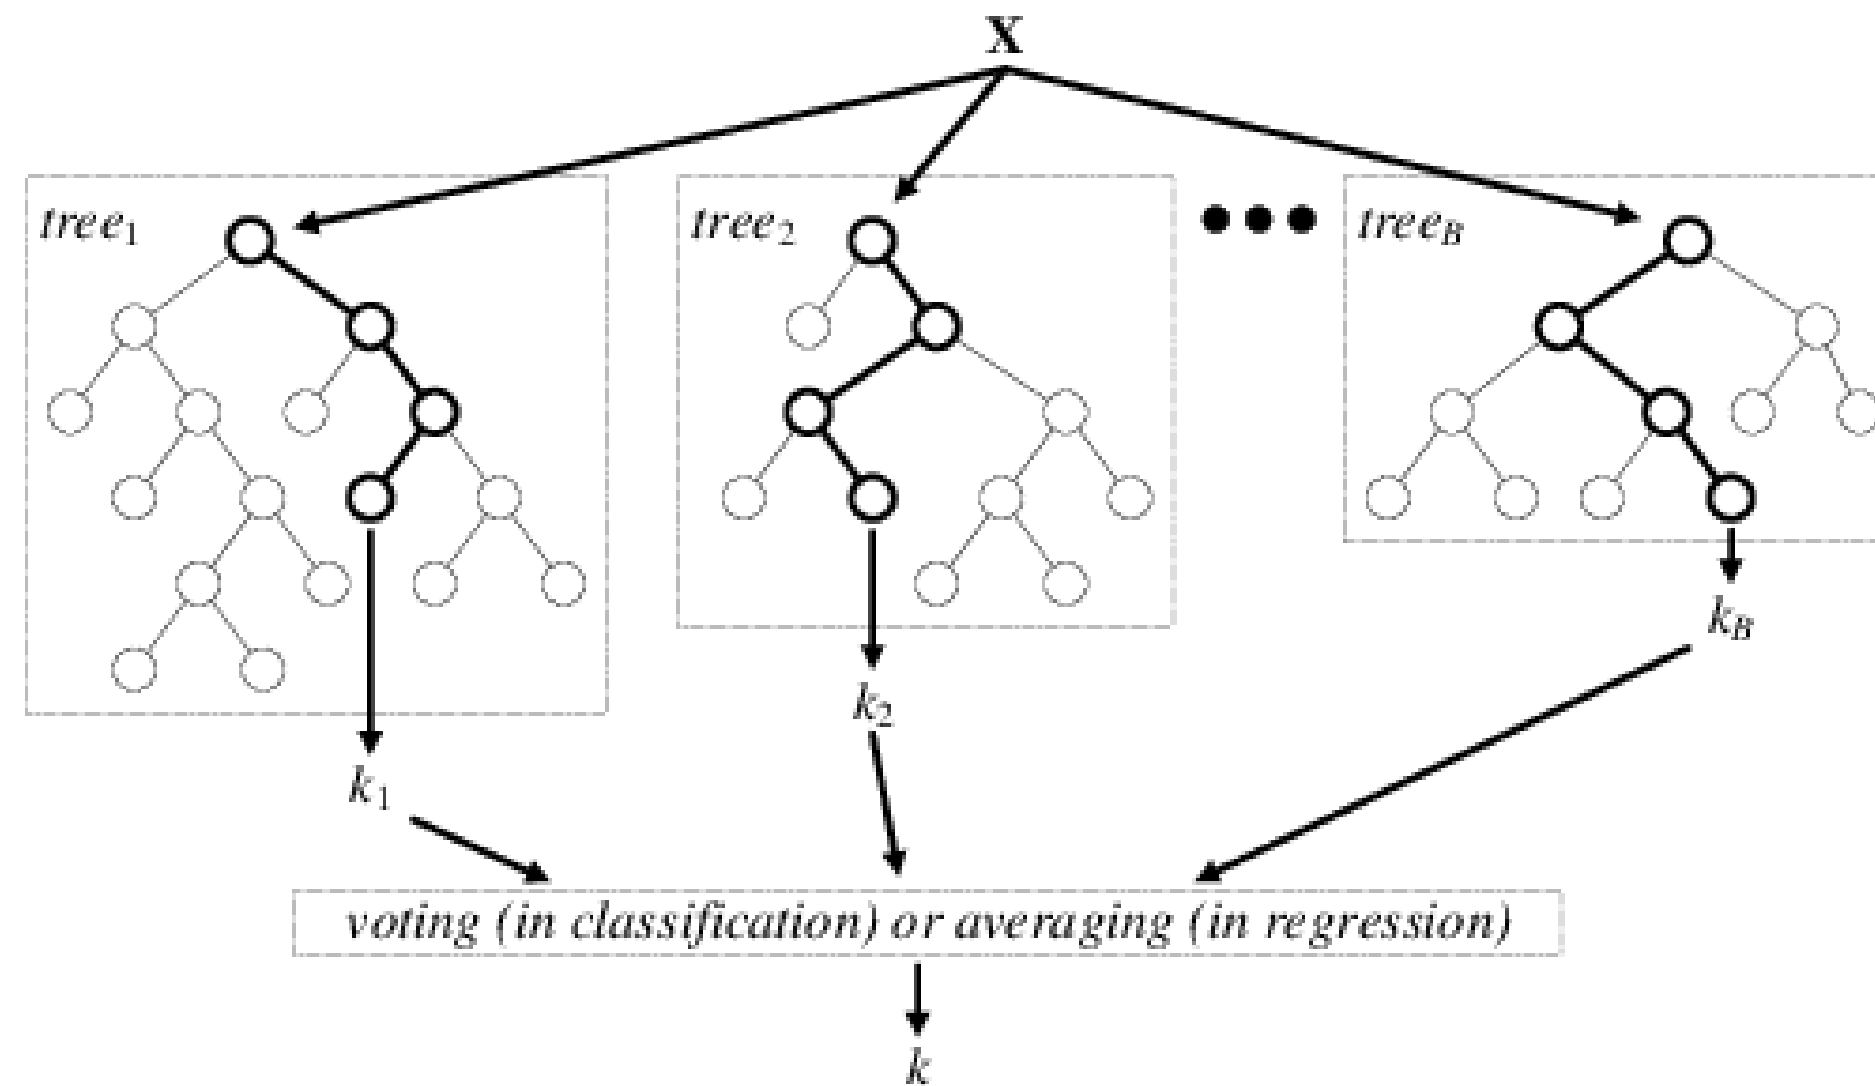
# Ensemble Methods

- "Wisdom of the crowd"
- Can often achieve better performance with collection of learners
- Often use shallow trees as base learners

Common methods for generating ensembles:

- **Bagging** (Bootstrap Aggregation)
  - Random Forest
- **Boosting**
  - Gradient Boosting
- **Stacking**

# Random Forest and Gradient Boosted Trees

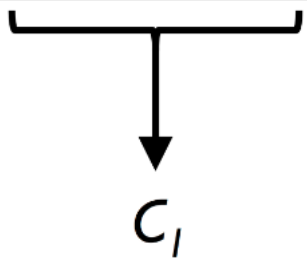


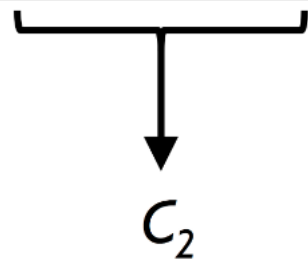
From [https://www.researchgate.net/publication/301638643\\_Electromyographic\\_Patterns\\_during\\_Golf\\_Swing\\_Activation\\_Sequence\\_Profiling\\_and\\_Prediction\\_of\\_Shot\\_Effectiveness](https://www.researchgate.net/publication/301638643_Electromyographic_Patterns_during_Golf_Swing_Activation_Sequence_Profiling_and_Prediction_of_Shot_Effectiveness)

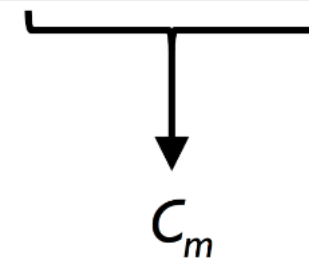
# Bagging with Random Forests

- Trees built with bootstrap samples and subsets of features
- Achieve variation with random selection of observations and features

Sample indices	Bagging round 1	Bagging round 2	...
1	2	7	...
2	2	3	...
3	1	2	...
4	3	1	...
5	7	1	...
6	2	7	...
7	4	7	...

 $C_1$

 $C_2$

 $C_m$

# Random Forests with sklearn

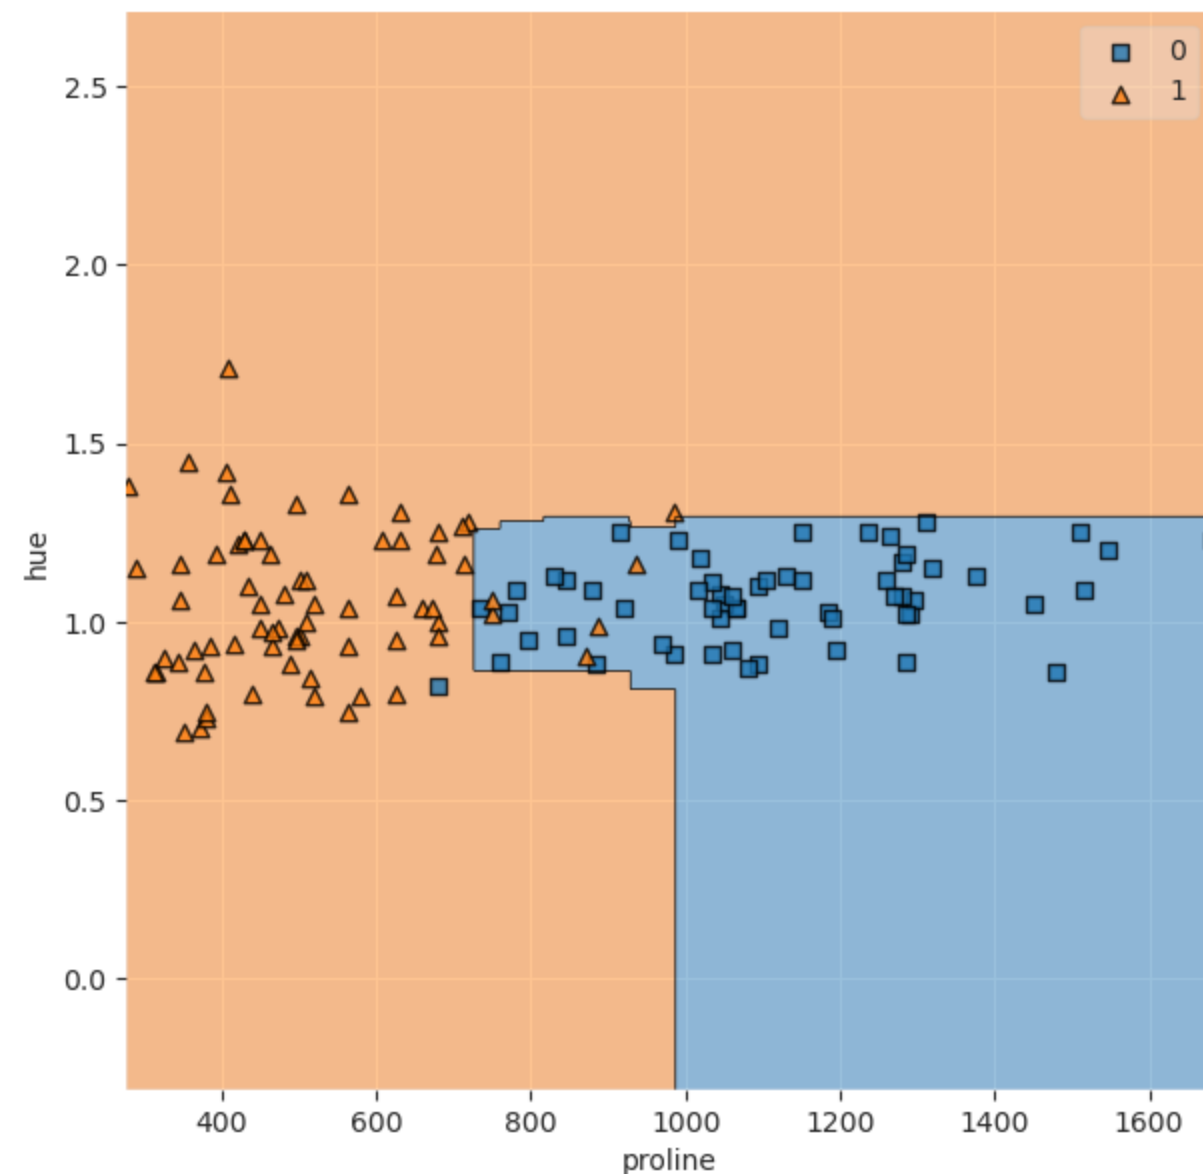


# Random Forests with sklearn

```
In [16]: from sklearn.ensemble import RandomForestClassifier

rfc = RandomForestClassifier(n_estimators=10, # number of trees in ensemble
                             max_depth=3,   # same as decision trees
                             n_jobs=-1,     # parallelize using all available cores, default: None=1
                             random_state=0) # for demonstration only

rfc.fit(X_2c, y_2c)
my_plot_decision_regions(X_2c, y_2c, rfc, figsize=(7, 7))
```



# Regression with RandomForest

# Regression with RandomForest

```
In [17]: from sklearn.ensemble import RandomForestRegressor

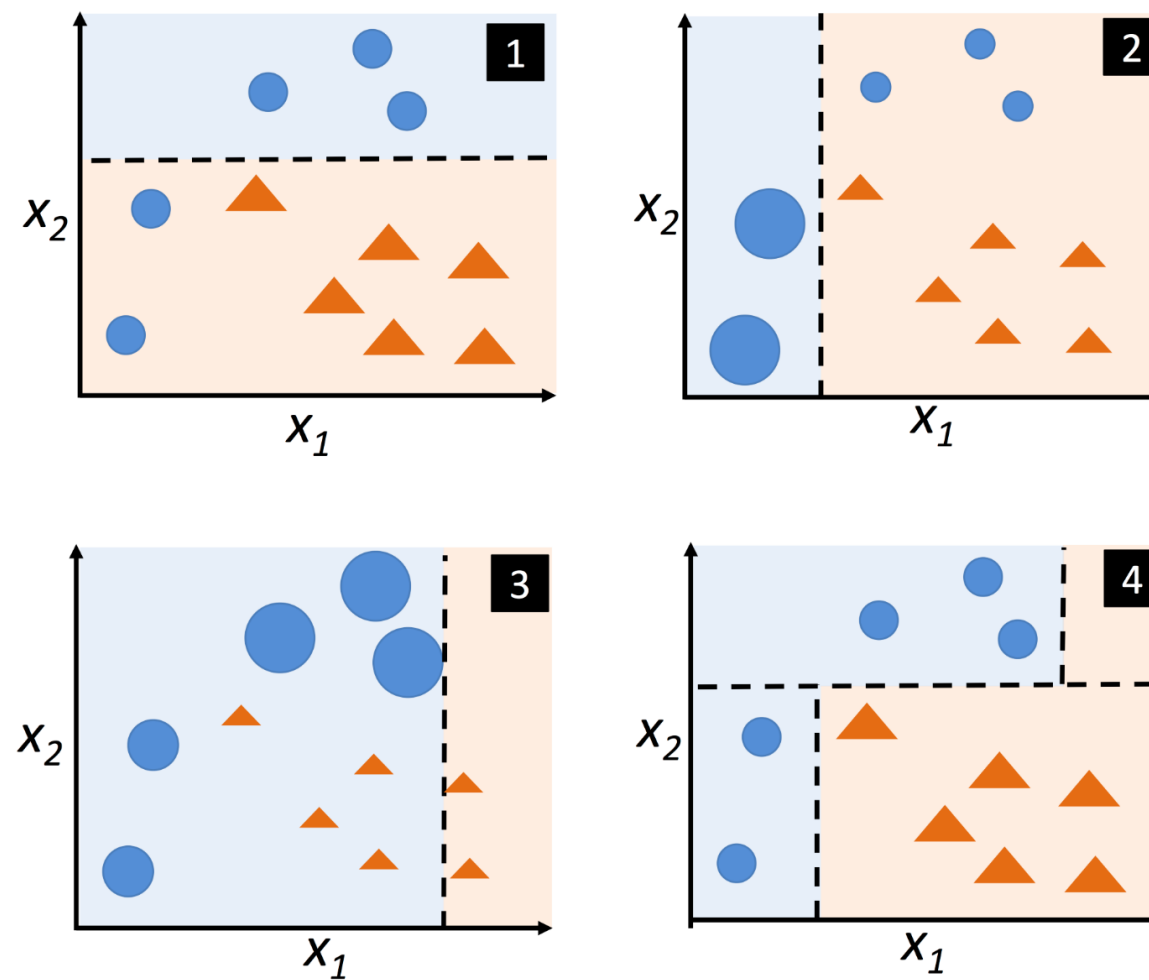
rfr = RandomForestRegressor(n_estimators=3, n_jobs=-1)
rfr.fit(df_wine[['proline']], df_wine.alcohol)

my_plot_regression(X_2c_zscore[['proline']], alcohol_2c_zscore, rfr)
```



# Gradient Boosted Trees

- Trees built by adding weight to mis-classification
- Achieve variation due to changes in weights on observations



From PML

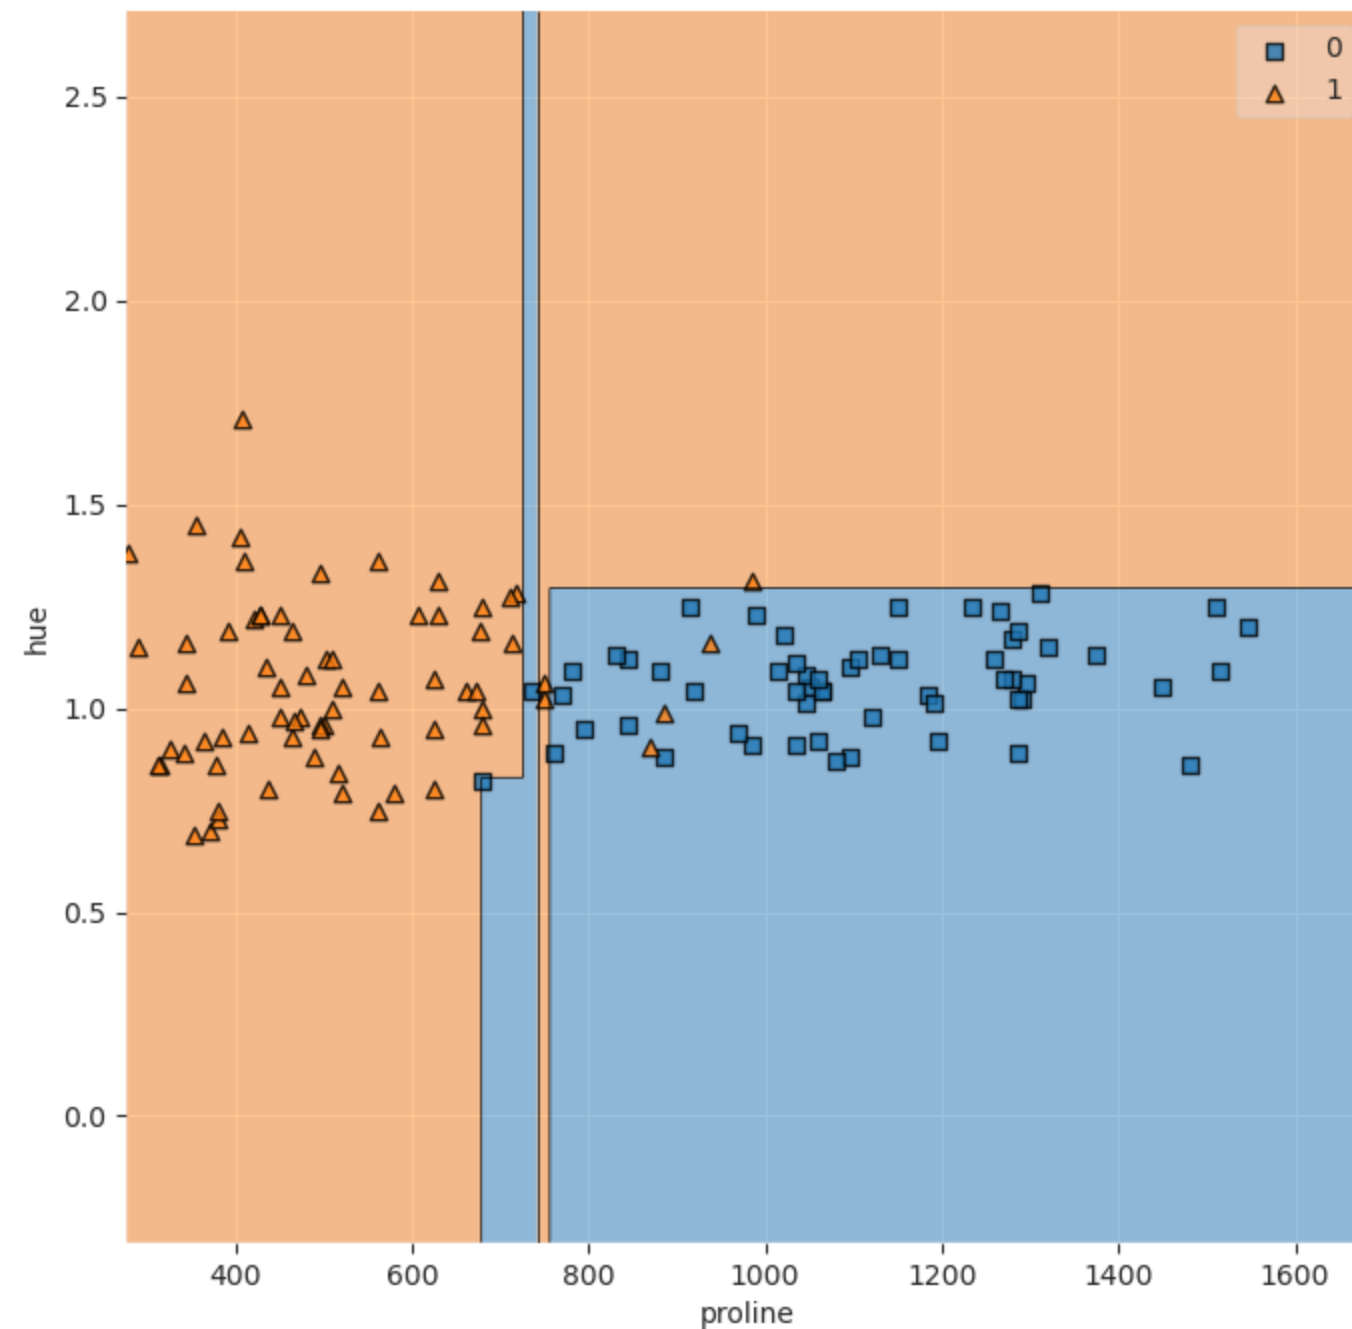
# Gradient Boosted Trees in sklearn

# Gradient Boosted Trees in sklearn

```
In [18]: from sklearn.ensemble import GradientBoostingClassifier

gbc = GradientBoostingClassifier(n_estimators=10)
gbc.fit(X_2c, y_2c)

my_plot_decision_regions(X_2c, y_2c, gbc)
```



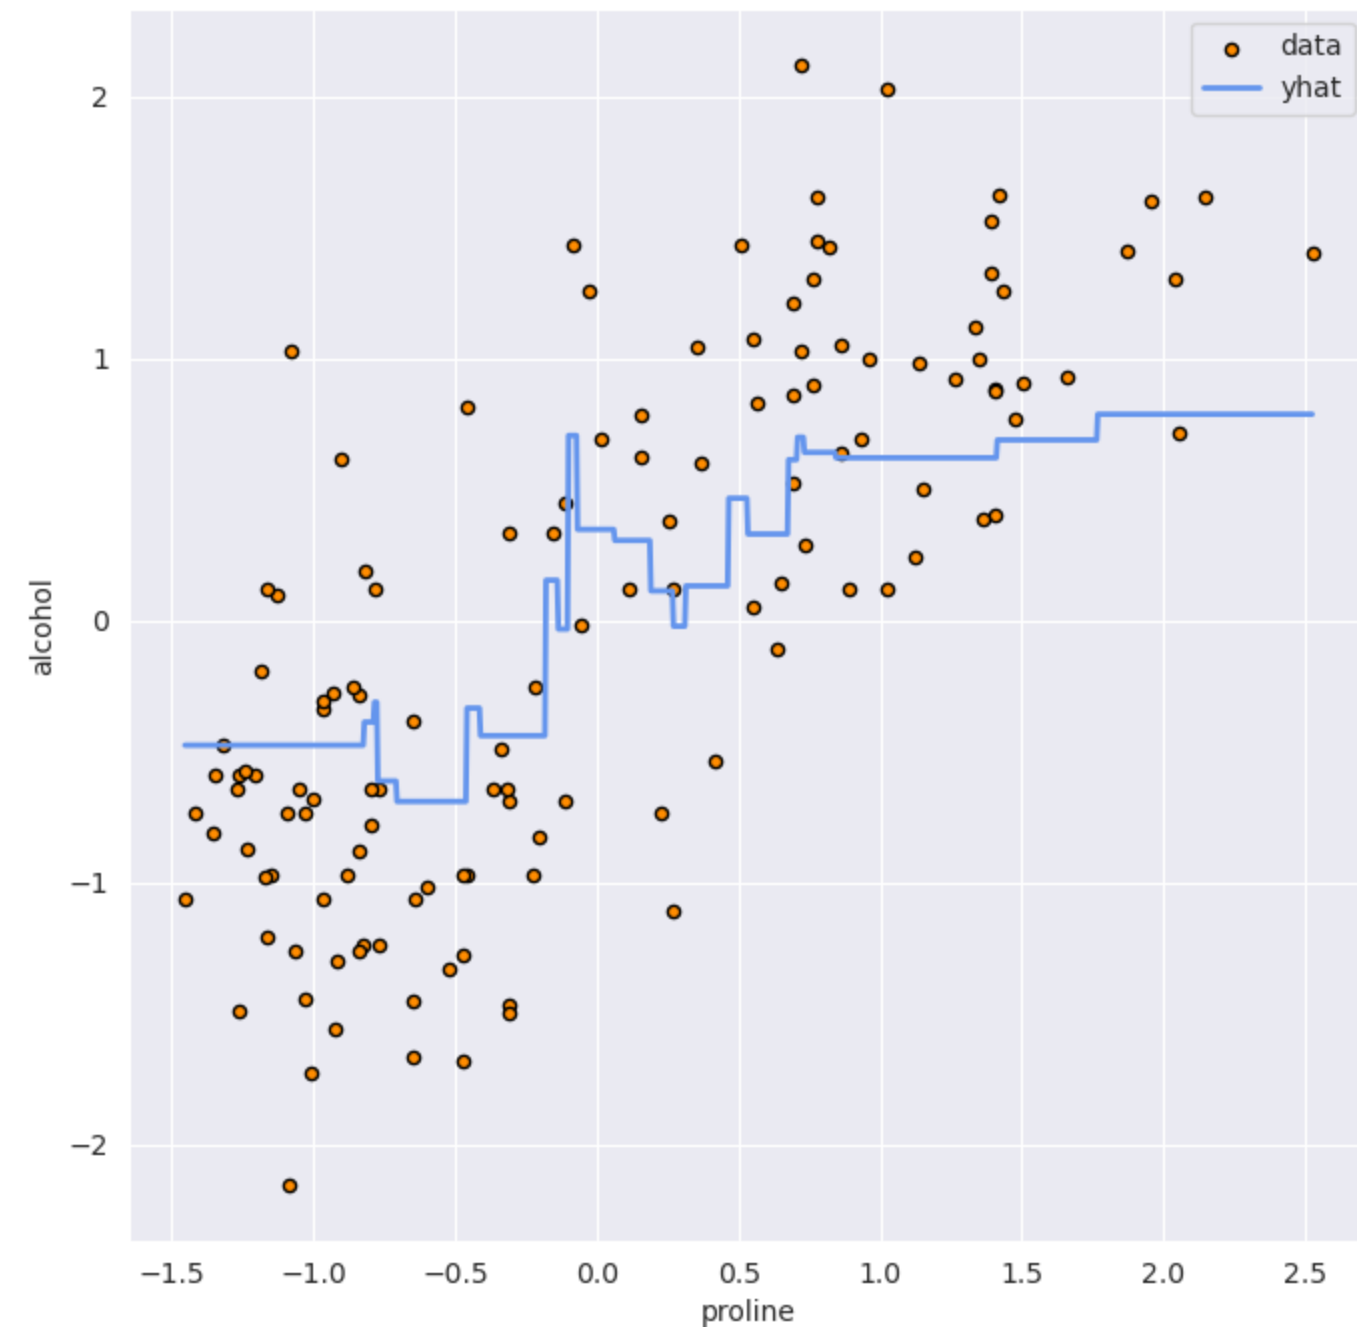
# Regression with Gradient Boosted Trees

# Regression with Gradient Boosted Trees

```
In [19]: from sklearn.ensemble import GradientBoostingRegressor

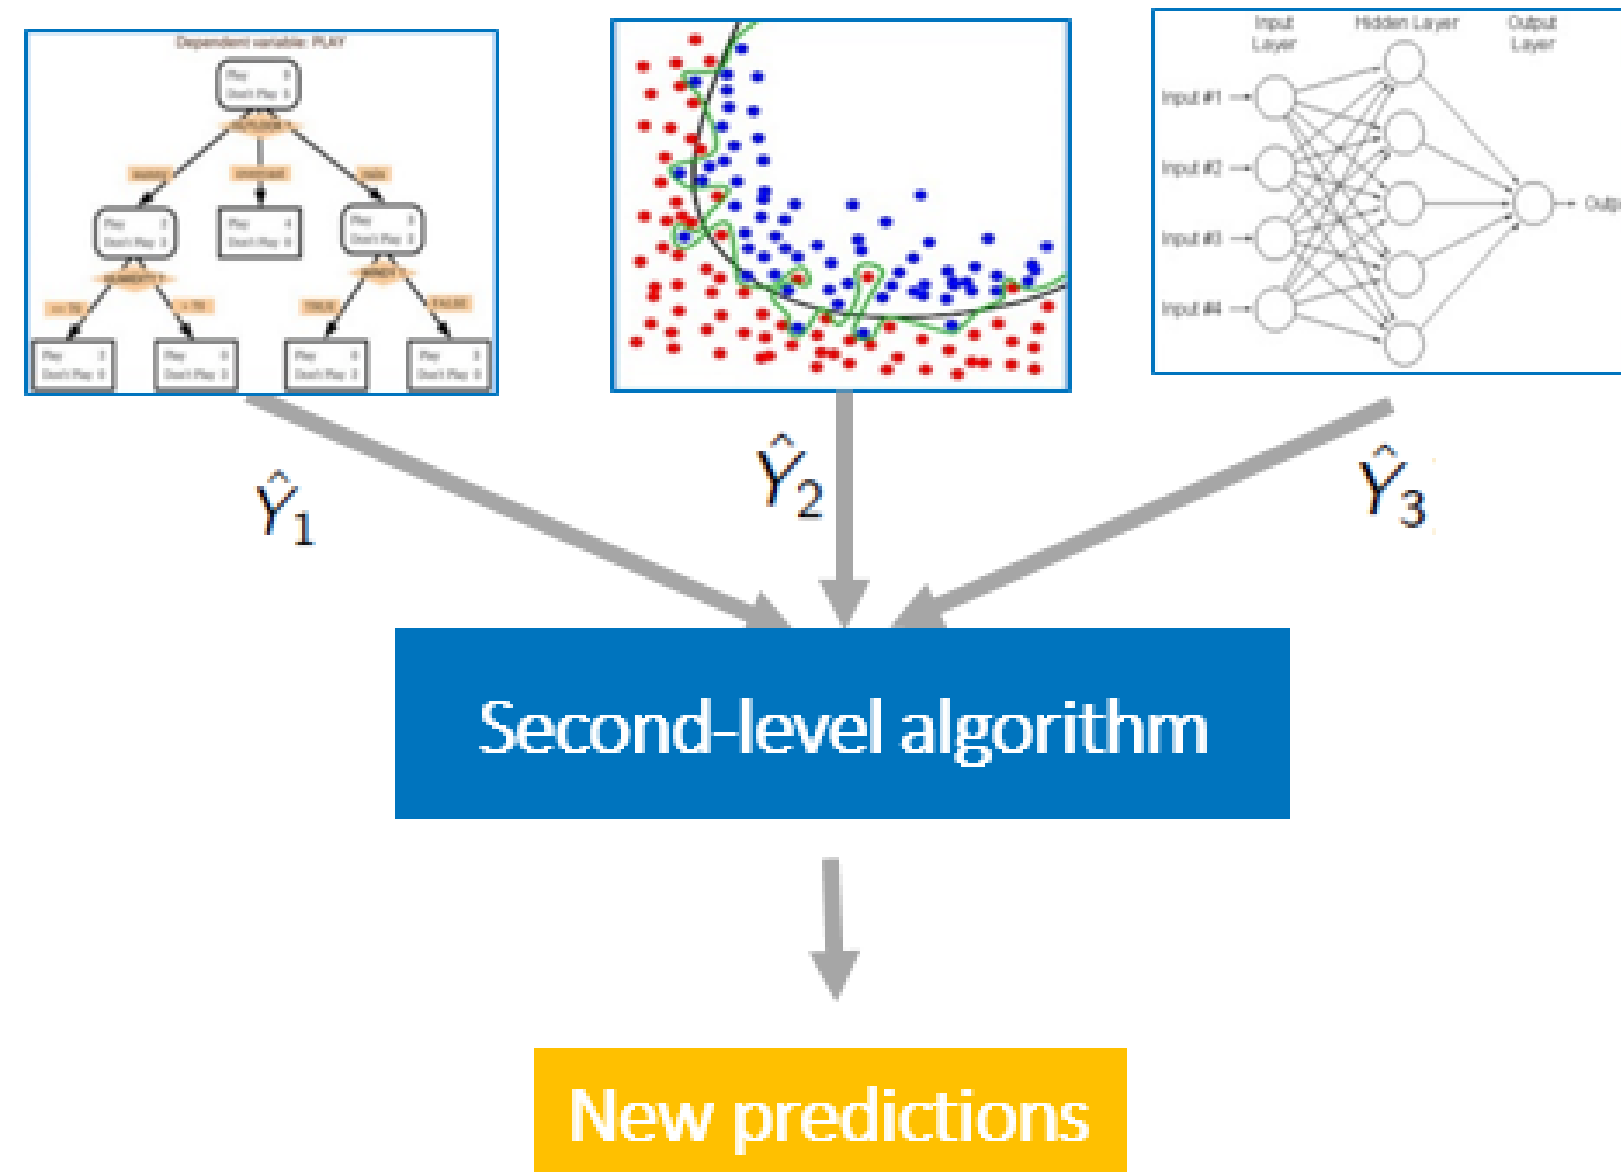
gbr = GradientBoostingRegressor(n_estimators=10)
gbr.fit(X_2c_zscore[['proline']], alcohol_2c_zscore)

my_plot_regression(X_2c_zscore[['proline']], alcohol_2c_zscore, gbr)
```





# Stacking



From <https://blogs.sas.com/content/subconsciousmusings/2017/05/18/stacked-ensemble-models-win-data-science-competitions/>

# Stacking for Classification

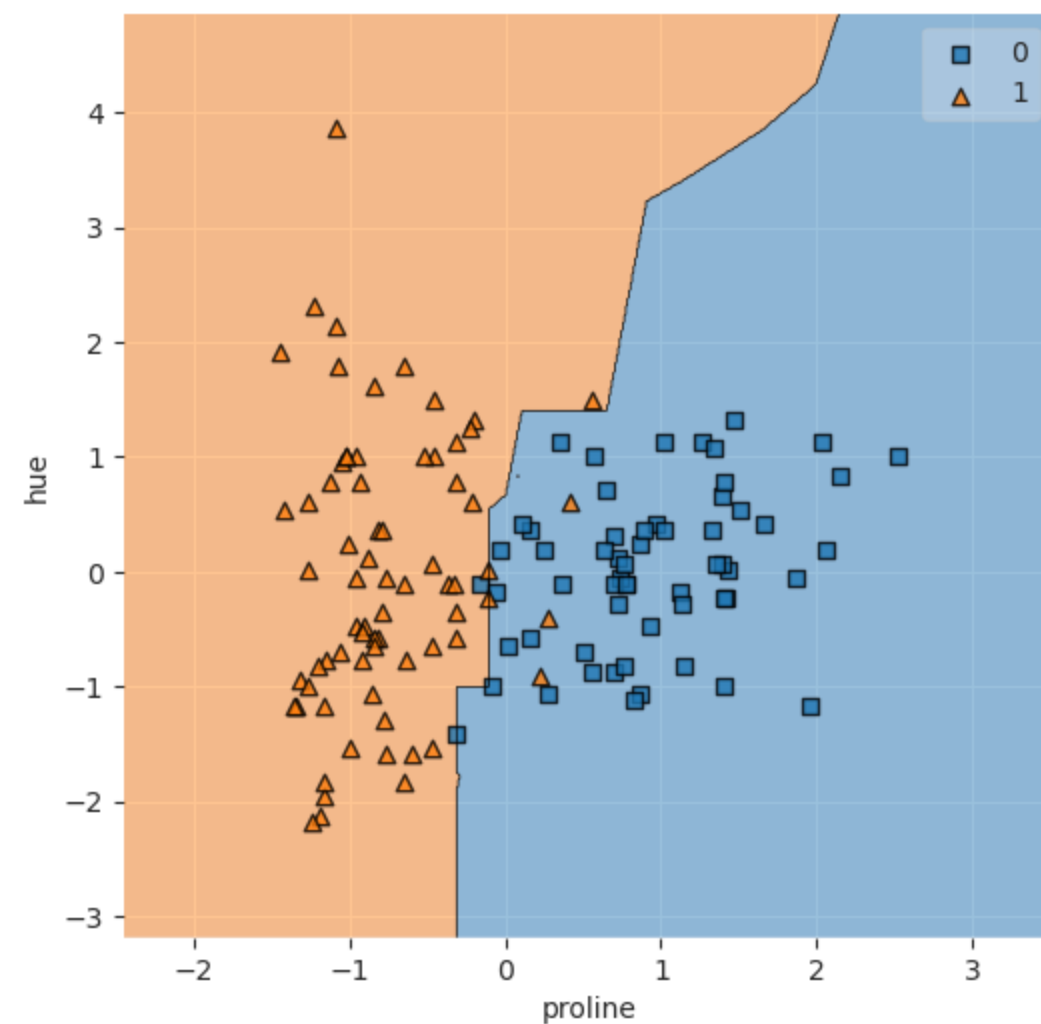
# Stacking for Classification

```
In [20]: from sklearn.ensemble import StackingClassifier

ensemble = [('lr', LogisticRegression(max_iter=1000)),
            ('dt', DecisionTreeClassifier(max_depth=3)),
            ('knn', KNeighborsClassifier(n_neighbors=3))]

stackc = StackingClassifier(estimators=ensemble,
                           final_estimator=LogisticRegression())
stackc.fit(X_2c_zscore, y_2c)

my_plot_decision_regions(X_2c_zscore, y_2c, stackc, figsize=(6,6))
```



# Stacking for Regression

# Stacking for Regression

```
In [21]: from sklearn.ensemble import StackingRegressor

ensemble = [('lr', LinearRegression()),
            ('dt', DecisionTreeRegressor(max_depth=3)),
            ('knn', KNeighborsRegressor(n_neighbors=6))]

stackr = StackingRegressor(estimators=ensemble,
                           final_estimator=LinearRegression())
stackr.fit(X_2c_zscore[['proline']], alcohol_2c_zscore)

my_plot_regression(X_2c_zscore[['proline']], alcohol_2c_zscore, stackr, figsize=(6, 6))
```



# Wine as Multi-Class Classification

# Wine as Multi-Class Classification

```
In [22]: X_mc = df_wine[['proline', 'hue']]
         y_mc = df_wine.target

         X_mc_zscore = X_mc.apply(zscore,axis=0)
         alcohol_mc_zscore = zscore(df_wine.alcohol)

         y_mc.value_counts().sort_index()
```

```
Out[22]: 0    59
         1    71
         2    48
         Name: target, dtype: int64
```

# Multiclass and Multilabel

- **Multiclass Classification** : more than two categories/classes
  - red/green/blue, flower type, integer 0-10
- **Multilabel Classification** : can assign more than one category to an instance
  - paper topics, entities in image
- **Multiclass-Multilabel/Multitask Classification** : >1 one property with >2 one categories
  - type of fruit AND color of fruit
- **Multioutput Regression** : more than one numeric targets
  - temperature AND humidity

See sklearn docs (<https://scikit-learn.org/stable/modules/multiclass.html#>)



# Sklearn Inherantly Multiclass

- `LogisticRegression(multi_class='multinomial')`
- `KNeighborsClassifier`
- `DecisionTreeClassifier`
- `RandomForestClassifier`

# Sklearn Inherantly Multiclass

- `LogisticRegression(multi_class='multinomial')`
- `KNeighborsClassifier`
- `DecisionTreeClassifier`
- `RandomForestClassifier`

```
In [23]: dt_mc = DecisionTreeClassifier().fit(X_mc_zscore,y_mc) # fit on multiclass

# generate 3 predictions
y_hats = dt_mc.predict(X_mc_zscore.iloc[[82,15,166]])

# display target and prediction
pd.DataFrame({'y':y_mc.iloc[[82,15,166]], 'y_hat':y_hats})
```

Out[23]:

	y	y_hat
82	1	1
15	0	0
166	2	2

# One Vs. Rest (OvR) Classification For Multiclass

# One Vs. Rest (OvR) Classification For Multiclass

- What about other models (eg Perceptron)?
- Can use any binary classifier for Multiclass classification by training multiple models:
  - model 1 : class 1 vs (class 2 and class 3)
  - model 2 : class 2 vs (class 1 and class 3)
  - model 3 : class 3 vs (class 1 and class 2)
- Then
  - Predict  $\hat{y}$  using the model with highest  $P(y = \hat{y} \mid x)$ , or distance from boundary, or ...

# Sklearn OvR for Multiclass

- `LogisticRegression(multi_class="ovr")`
- `GradientBoostingClassifier`
- `Perceptron`

# OvR For Logistic Regression

# OvR For Logistic Regression

```
In [24]: logr_mc = LogisticRegression(multi_class='ovr', # default
                                     max_iter=1000)      # to avoid timeout errors

logr_mc.fit(X_mc_zscore, y_mc)
y_hats = logr_mc.predict(X_mc_zscore.iloc[[82,15,166]]) # generate 3 predictions
y_prob = logr_mc.predict_proba(X_mc_zscore.iloc[[82,15,166]])
pd.DataFrame({'y_hat':y_hats, 'p_c1':y_prob[:,0], 'p_c2':y_prob[:,1], 'p_c3':y_prob[:,2]})\
    .style.background_gradient(subset=['p_c1', 'p_c2', 'p_c3']).format('{:.2f}', subset=['p_c1', 'p_c2', 'p_c3'])
```

Out[24]:

	y_hat	p_c1	p_c2	p_c3
0	1	0.15	0.85	0.00
1	0	0.97	0.03	0.00
2	2	0.18	0.34	0.48

# OvR For Logistic Regression

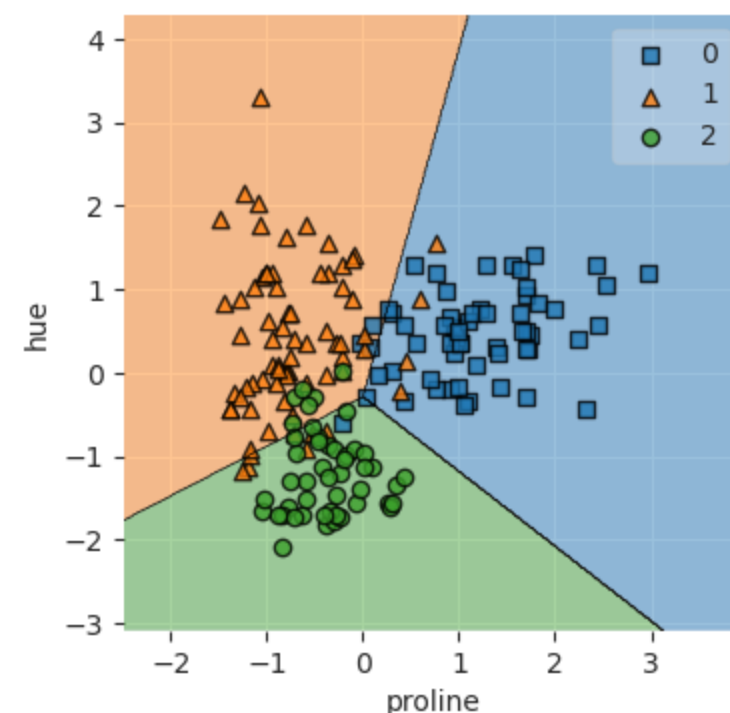
```
In [24]: logr_mc = LogisticRegression(multi_class='ovr', # default
                                     max_iter=1000)      # to avoid timeout errors

logr_mc.fit(X_mc_zscore,y_mc)
y_hats = logr_mc.predict(X_mc_zscore.iloc[[82,15,166]]) # generate 3 predictions
y_prob = logr_mc.predict_proba(X_mc_zscore.iloc[[82,15,166]])
pd.DataFrame({'y_hat':y_hats, 'p_c1':y_prob[:,0], 'p_c2':y_prob[:,1], 'p_c3':y_prob[:,2]})\
.style.background_gradient(subset=['p_c1', 'p_c2', 'p_c3']).format('{:.2f}',subset=['p_c1', 'p_c2', 'p_c3'])
```

Out[24]:

	y_hat	p_c1	p_c2	p_c3
0	1	0.15	0.85	0.00
1	0	0.97	0.03	0.00
2	2	0.18	0.34	0.48

```
In [25]: my_plot_decision_regions(X_mc_zscore,y_mc,logr_mc,figsize=(4,4))
```





# One vs. One Classification

- Train one classifier for each pair-wise comparison of classes
- SVC

# Inherantly Multilabel (aka Multioutput)

- `KNeighborsClassifier`
- `DecisionTreeClassifier`
- `MLPClassifier`
- `RandomForestClassifier`

# Inherantly Multilabel (aka Multioutput)

- KNeighborsClassifier
- DecisionTreeClassifier
- MLPClassifier
- RandomForestClassifier

```
In [26]: X_ml = X_mc[['proline']].iloc[-3:] # get 3 samples
         y_ml = np.array([[1, 0, 1], [0, 1, 1], [0, 0, 0]]) # generate 3 sets of random multi-labels for demonstration

         dt_ml = DecisionTreeClassifier(max_depth=1).fit(X_ml,y_ml) # fit multilabel
         dt_ml.predict(X_ml)
```

```
Out[26]: array([[0, 0, 1],
                [0, 0, 1],
                [0, 0, 0]])
```

# Sklearn `MultiOutputClassifier` meta-estimator

- fits one classifier per target (One vs. Rest)

# Sklearn `MultiOutputClassifier` meta-estimator

- fits one classifier per target (One vs. Rest)

```
In [27]: from sklearn.multioutput import MultiOutputClassifier

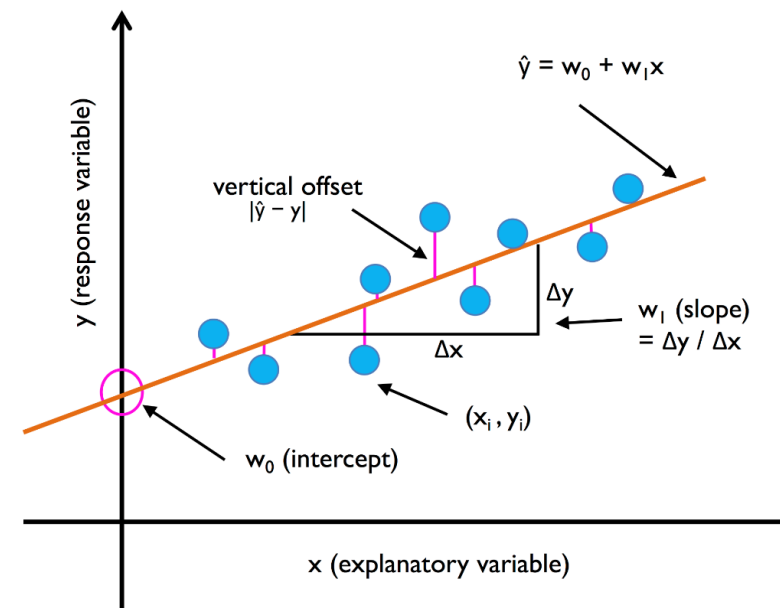
         # wrap your classifier with MultiOutputClassifier
         mc_logr = MultiOutputClassifier(LogisticRegression())
         mc_logr.fit(X_ml, y_ml)
         mc_logr.predict(X_ml)
```

```
Out[27]: array([[0, 0, 1],
                [1, 1, 1],
                [0, 0, 0]])
```

# Review of Models

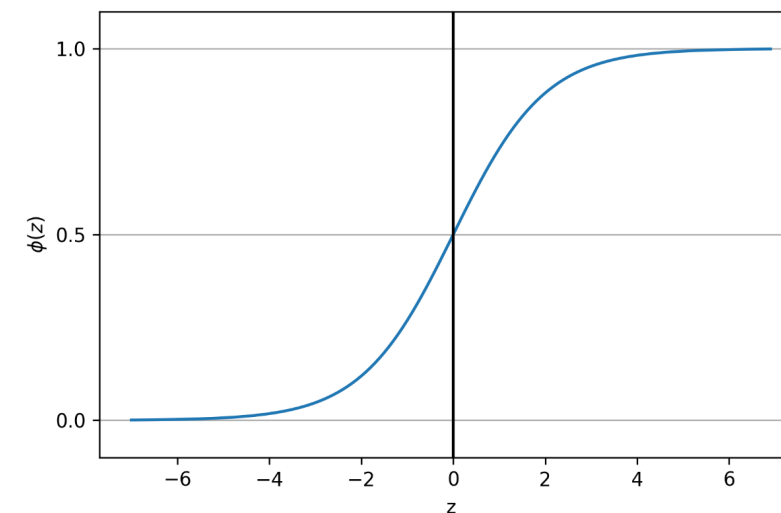
# Model Review: Simple/Multiple Linear Regression

- Use for: Regression
- Pros:
  - fast to train
  - interpretable coefficients
- Cons:
  - assumes linear relationship
  - depends on removing colinear features



# Model Review: Logistic Regression

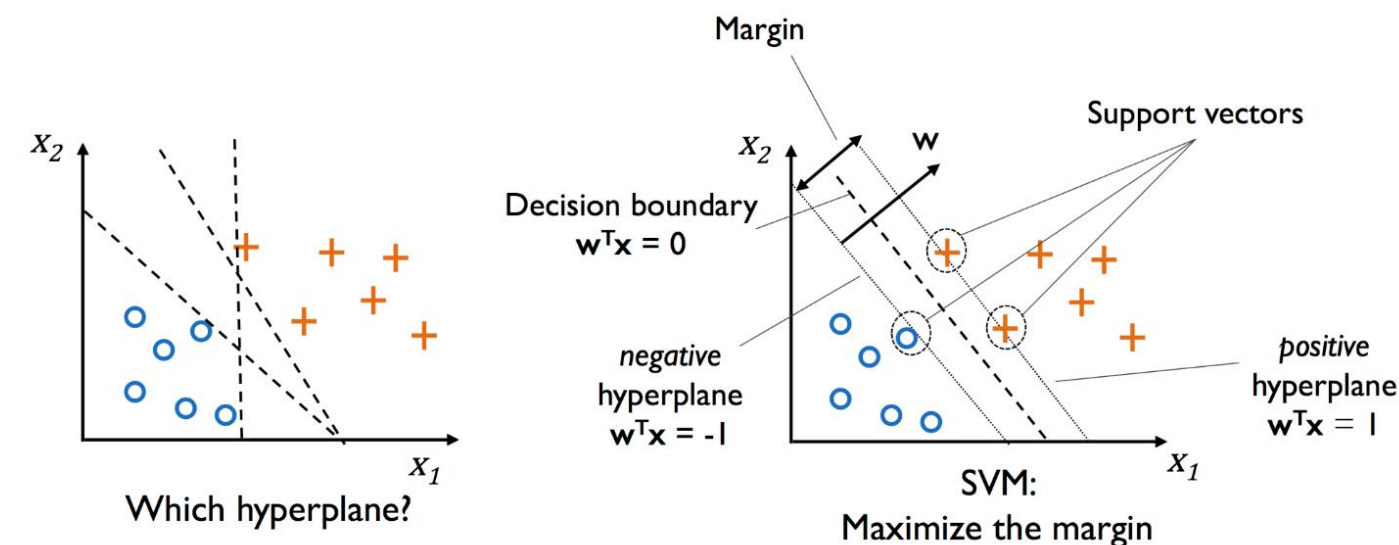
- Use for: Classification
- Pros:
  - fast to train
  - interpretable coefficients (log odds)
- Cons:
  - assumes linear boundary
  - depends on removing colinear features





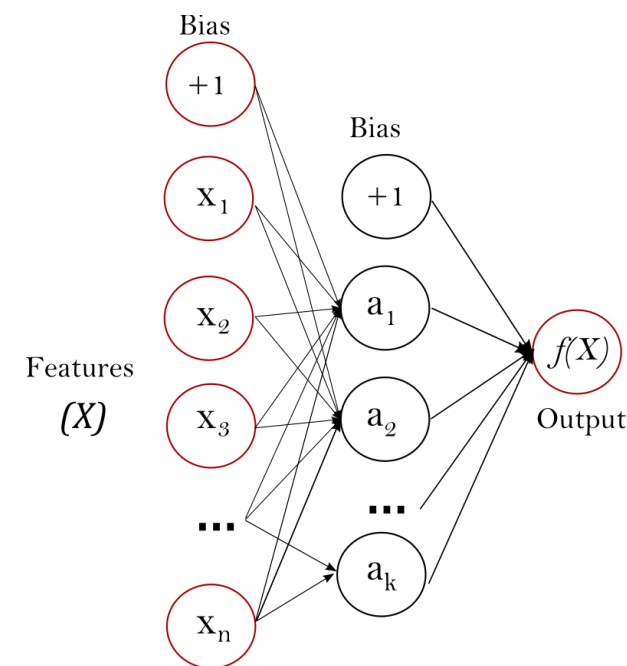
# Model Review: Support Vector Machine (SVM)

- Use for: Classification and Regression
- Pros:
  - fast to evaluate
  - can use kernel trick to learn non-linear functions
- Cons:
  - slow to train
  - can fail to converge on very large datasets



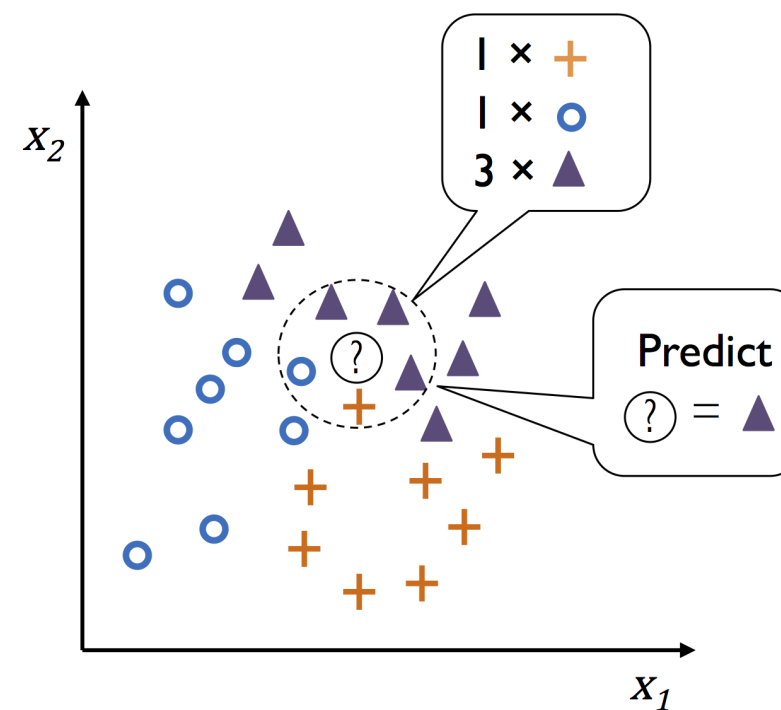
# Model Review: Multi-Layer Perceptron

- Use for Classification or Regression
- Pros:
  - non-linear boundary
- Cons:
  - non-convex loss function (sensitive to initial weights)
  - sensitive to feature scaling
  - no GPU support in sklearn: use tensorflow or pytorch



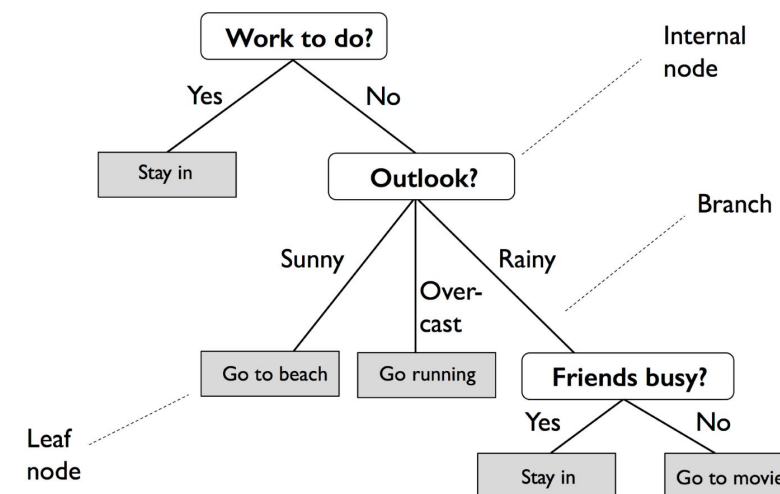
# Model Review: k Nearest Neighbor (kNN)

- Use for: Classification or Regression
- Pros:
  - fast to train
  - non-linear boundary
- Cons:
  - potentially slow to predict
  - curse of dimensionality



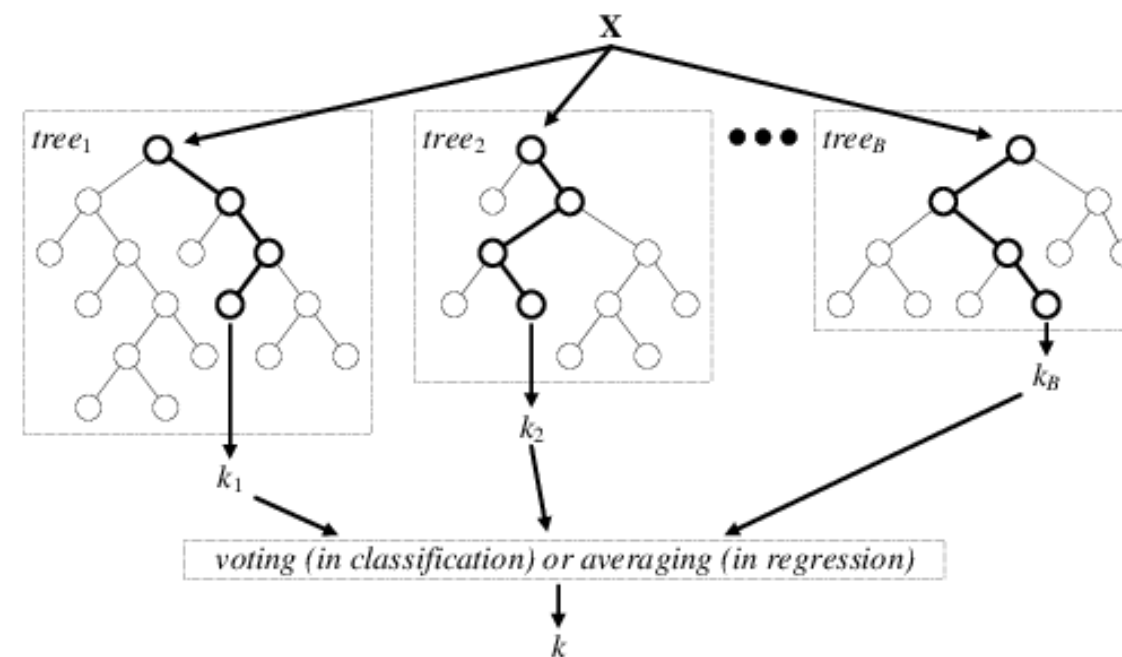
# Model Review: Decision Tree

- Use for: Classification or Regression
- Pros:
  - very interpretable
  - quick to predict
  - can handle numeric and categorical variables without transformation
- Cons:
  - tendency to overfit (learn training set too well, more next class!)



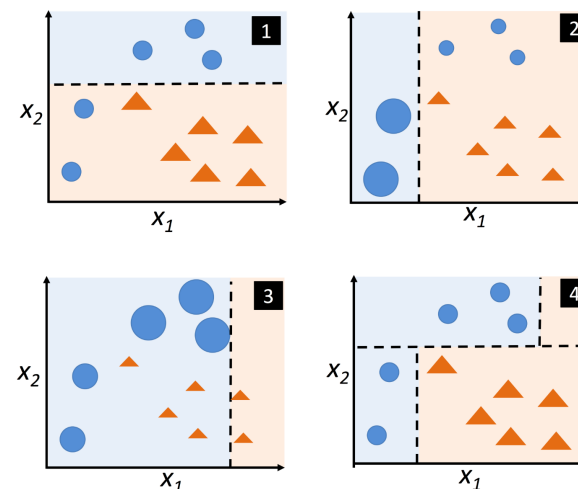
# Model Review: Random Forest (Ensemble via Bagging)

- Use for: Classification or Regression
- Pros:
  - less likely to overfit than decision tree
  - quick to train (through parallelization, quick to predict)
- Cons:
  - less interpretable, though still possible



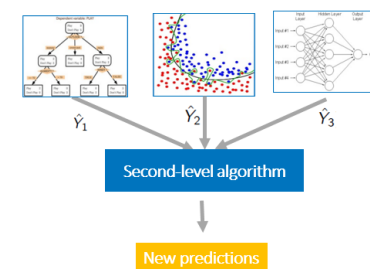
# Model Review: Gradient Boosted Trees (Ensemble via Boosting)

- Use for: Classification or Regression
- Pros:
  - pays more attention to difficult decision regions
  - quick to predict
  - tends to work well on difficult tasks
- Cons:
  - slow to train (parallelization not possible)
  - less interpretable, though still possible



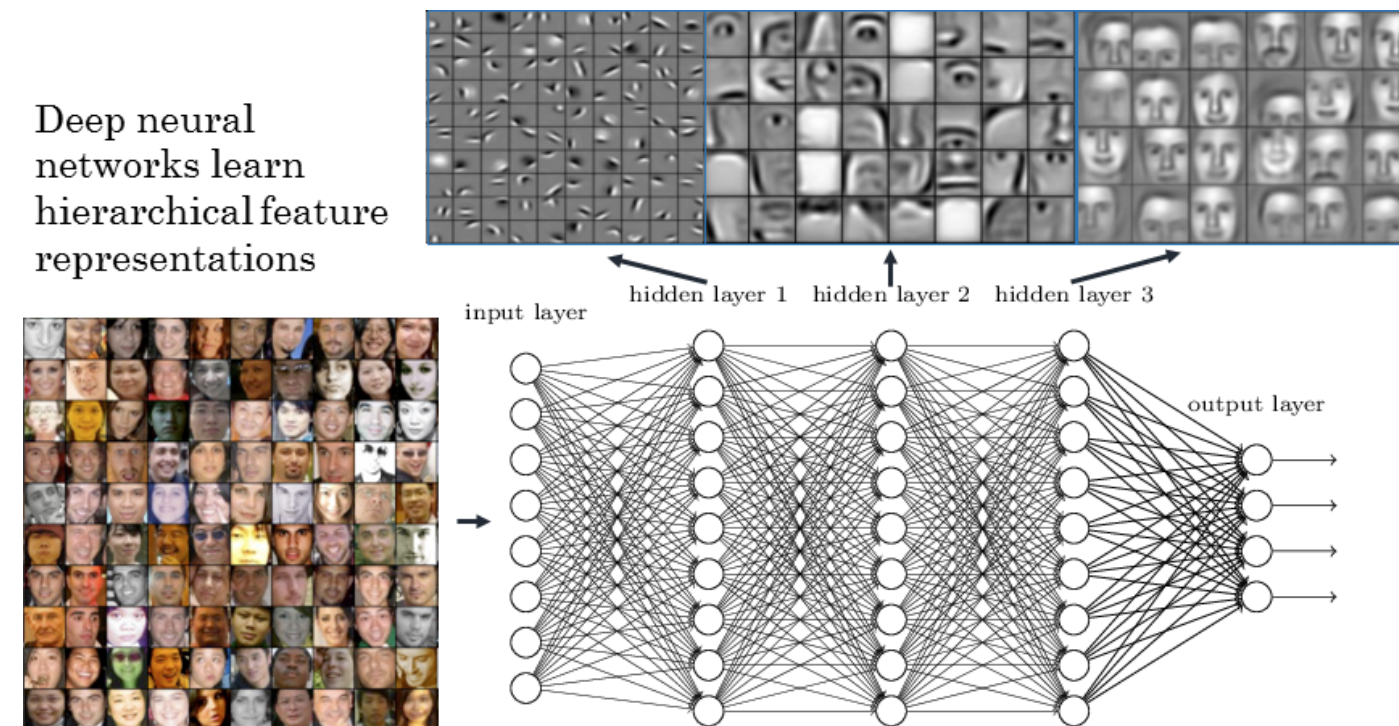
# Model Review: Ensemble via Stacking

- Use for: Classification or Regression
- Pros:
  - combines benefits of multiple learning types
  - easy to implement
  - tends to win competitions
- Cons:
  - difficult to interpret
  - training/prediction time depends on component models



# Neural Networks (aka Deep Learning)

- Pros and Cons of Deep Learning
  - sensitive to initialization and structure
  - high complexity -> needs more data
  - low interpretability
  - can learn complex interactions
  - performs well on tasks involving complex signals (ex images, sound, etc)





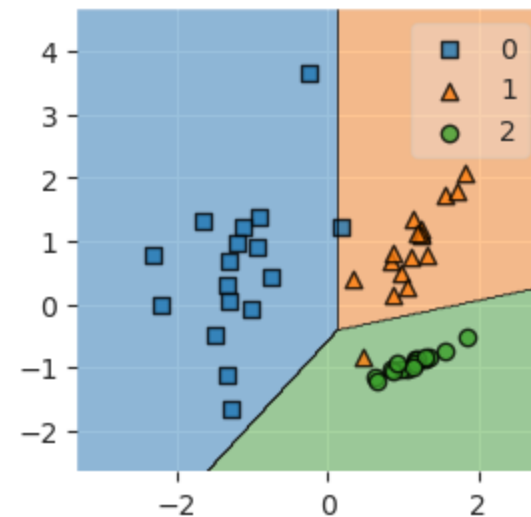
# Playing with synthetic classification datasets

# Playing with synthetic classification datasets

```
In [28]: from sklearn.datasets import make_classification, make_multilabel_classification

X_syn,y_syn = make_classification(n_samples=50,
                                n_features=2,
                                n_informative=2,
                                n_redundant=0,
                                n_clusters_per_class=1,
                                class_sep=1,
                                n_classes=3,
                                random_state=0,
                                )

fig,ax = plt.subplots(1,1,figsize=(3,3))
plot_decision_regions(X_syn,y_syn,LogisticRegression().fit(X_syn,y_syn));
```



# Playing with synthetic classification datasets - multilabel

# Playing with synthetic classification datasets - multilabel

```
In [32]: X_syn_ml, y_syn_ml = make_multilabel_classification(n_samples=100,  
                                                            n_features=2,  
                                                            n_classes=5,  
                                                            random_state=0  
                                                            )  
  
print(X_syn_ml[:10])  
print()  
print(y_syn_ml[:10])
```

```
[[24. 25.]  
 [38. 15.]  
 [39. 14.]  
 [23. 20.]  
 [26. 29.]  
 [30. 16.]  
 [22. 30.]  
 [25. 22.]  
 [29. 12.]  
 [25. 21.]]
```

```
[[0 0 0 0 0]  
 [0 0 1 0 1]  
 [0 1 1 1 1]  
 [1 1 1 1 1]  
 [1 0 0 1 0]  
 [0 1 0 1 1]  
 [1 0 0 0 0]  
 [1 1 1 1 0]  
 [0 0 0 0 0]  
 [0 0 1 1 0]]
```

**Questions?**