

Elements Of Data Science - F2022

Week 11: Clustering and Recommendation Systems

11/16/2022

TODOs

- Readings:
 - PDSH: Chap 3.11 Working with Time Series
 - PDSH: Chap 5.06 Example: Predicting Bicycle Traffic
 - Optional: Python for Data Analysis: Chap 11: Time Series
 - Optional: PML: Chap 9: Embedding a Machine Learning Model into a Web Application
- HW3: due Friday Nov 18th 11:59pm ET
- Quiz 11: due Tuesday Nov 22nd, 11:59pm ET
- HW4: out Friday night, due Friday Dec 2nd 11:59pm ET

Quiz Common Mistakes (points off)

- don't remove instructions from quiz/homework
- `.info()` not `.info`: make sure function/method calls are made with `()`
- Pandas `.sample()` default `n=1`: need to set `n=` or `frac=`
- LinearRegression (regression) vs LogisticRegression (classification)
- using a model "with default settings" means `Model()` or just a subset of parameters set
- Be careful which dataset you're training/evaluating on: `X_train` vs `X_test`
- Make sure all plotting settings get used (eg `hue=`)

Today

- Clustering
- Recommendation Systems
- Imbalanced Data

Questions?

Environment Setup

Environment Setup

```
In [1]: import numpy
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

from mlxtend.plotting import plot_decision_regions

import warnings
warnings.filterwarnings('ignore')

sns.set_style('darkgrid')
%matplotlib inline
```

Clustering

- Can we group our data based on the features alone?
- **Unsupervised:** There is no label/target y
- Use similarity to group X into k clusters

Why do Clustering?

- Exploratory data analysis
- Group media: images, music, news articles,...
- Group people: social network
- Science applications: gene families, psychological groups,...
- Image segmentation: group pixels, regions, ...
- ...

Clustering Methods

- k-Means
- **Heirarchical Agglomerative Clustering**
- Spectral Clustering
- DBScan
- ...

Clustering: k -Means

- Not to be confused with k-NN!
- Idea:
 - Finds k points in space as cluster centers (means)
 - Assigns datapoints to their closest cluster mean
- Need to specify the number of clusters k up front
- sklearn uses euclidean distance to judge similarity

k -Means: How it works

FIRST: choose initial k points (means)

A: fix means \rightarrow assign all datapoints to their closest mean

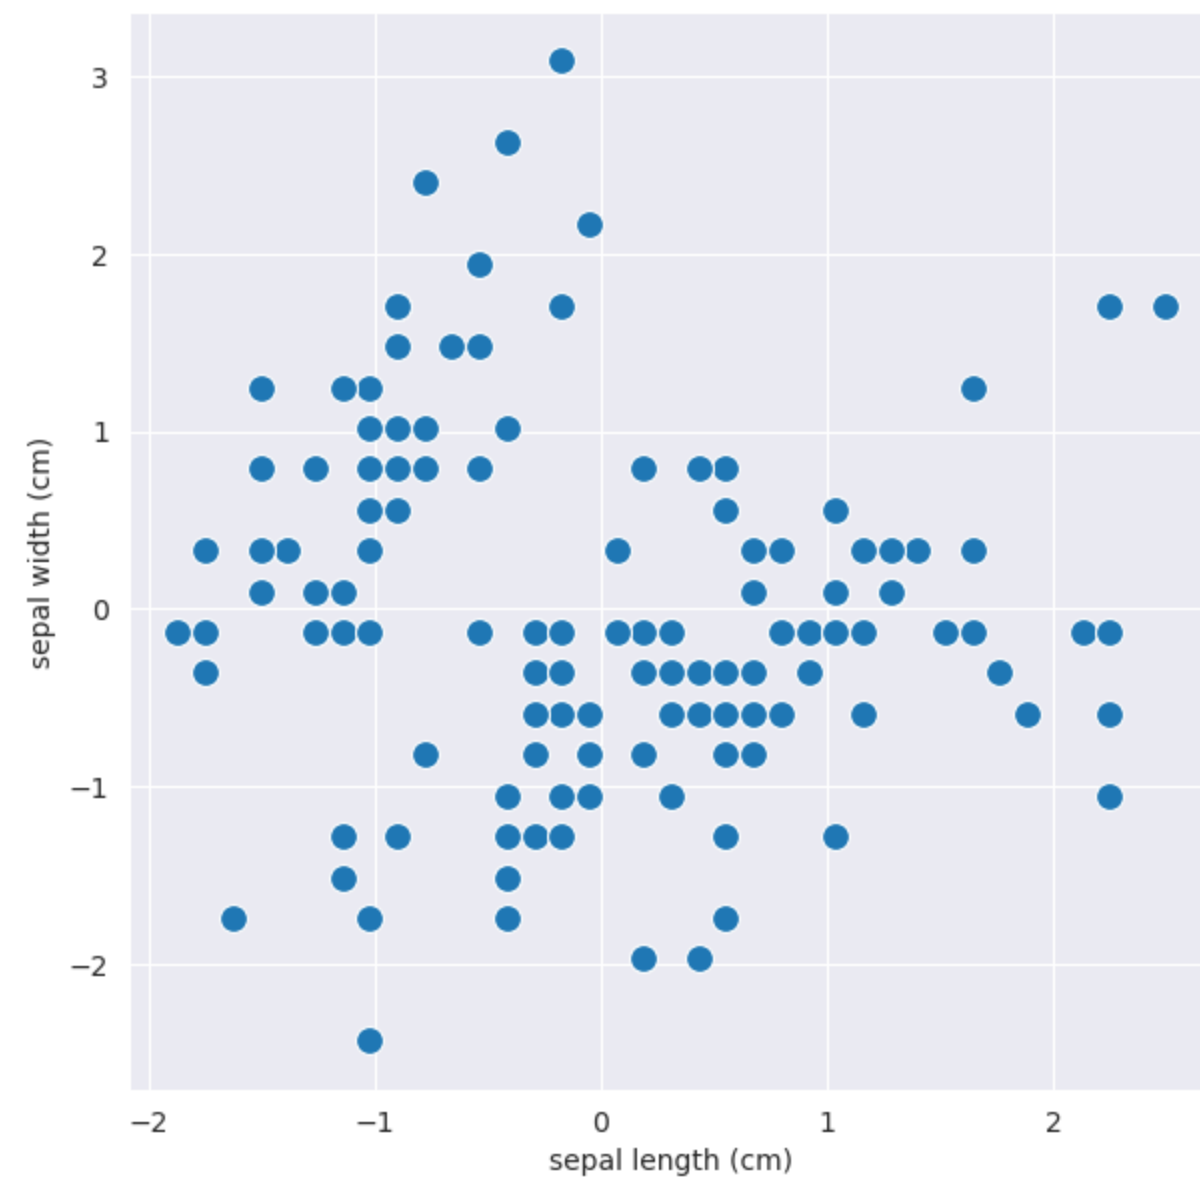
B: fix cluster assignments \rightarrow recalculate means

RETURN TO A and Repeat until convergence!

Load Example Data

```
In [2]: from sklearn.datasets import load_iris
from sklearn.preprocessing import StandardScaler
iris = load_iris()
X_iris = StandardScaler().fit_transform(iris.data[:, :2])
X_iris = pd.DataFrame(X_iris, columns=iris.feature_names[:2])

fig, ax = plt.subplots(1, 1, figsize=(7, 7))
sns.scatterplot(x='sepal length (cm)', y='sepal width (cm)', data=X_iris, s=100);
```



KMeans in sklearn

KMeans in sklearn

```
In [3]: from sklearn.cluster import KMeans  
  
km = KMeans(n_clusters=2, init='random', random_state=0) # default init=k-means++  
  
c = km.fit_predict(X_iris)
```

KMeans in sklearn

```
In [3]: from sklearn.cluster import KMeans

km = KMeans(n_clusters=2, init='random', random_state=0) # default init=k-means++

c = km.fit_predict(X_iris)
```

```
In [4]: # cluster assignments
tmp = X_iris.copy()
tmp['cluster_assignments'] = c
tmp.sample(5, random_state=0).round(2)
```

Out[4]:

| | sepal length (cm) | sepal width (cm) | cluster_assignments |
|-----|-------------------|------------------|---------------------|
| 114 | -0.05 | -0.59 | 1 |
| 62 | 0.19 | -1.97 | 1 |
| 33 | -0.42 | 2.63 | 0 |
| 107 | 1.77 | -0.36 | 1 |
| 7 | -1.02 | 0.79 | 0 |

KMeans in sklearn

```
In [3]: from sklearn.cluster import KMeans

km = KMeans(n_clusters=2, init='random', random_state=0) # default init=k-means++

c = km.fit_predict(X_iris)
```

```
In [4]: # cluster assignments
tmp = X_iris.copy()
tmp['cluster_assignments'] = c
tmp.sample(5, random_state=0).round(2)
```

Out[4]:

| | sepal length (cm) | sepal width (cm) | cluster_assignments |
|-----|-------------------|------------------|---------------------|
| 114 | -0.05 | -0.59 | 1 |
| 62 | 0.19 | -1.97 | 1 |
| 33 | -0.42 | 2.63 | 0 |
| 107 | 1.77 | -0.36 | 1 |
| 7 | -1.02 | 0.79 | 0 |

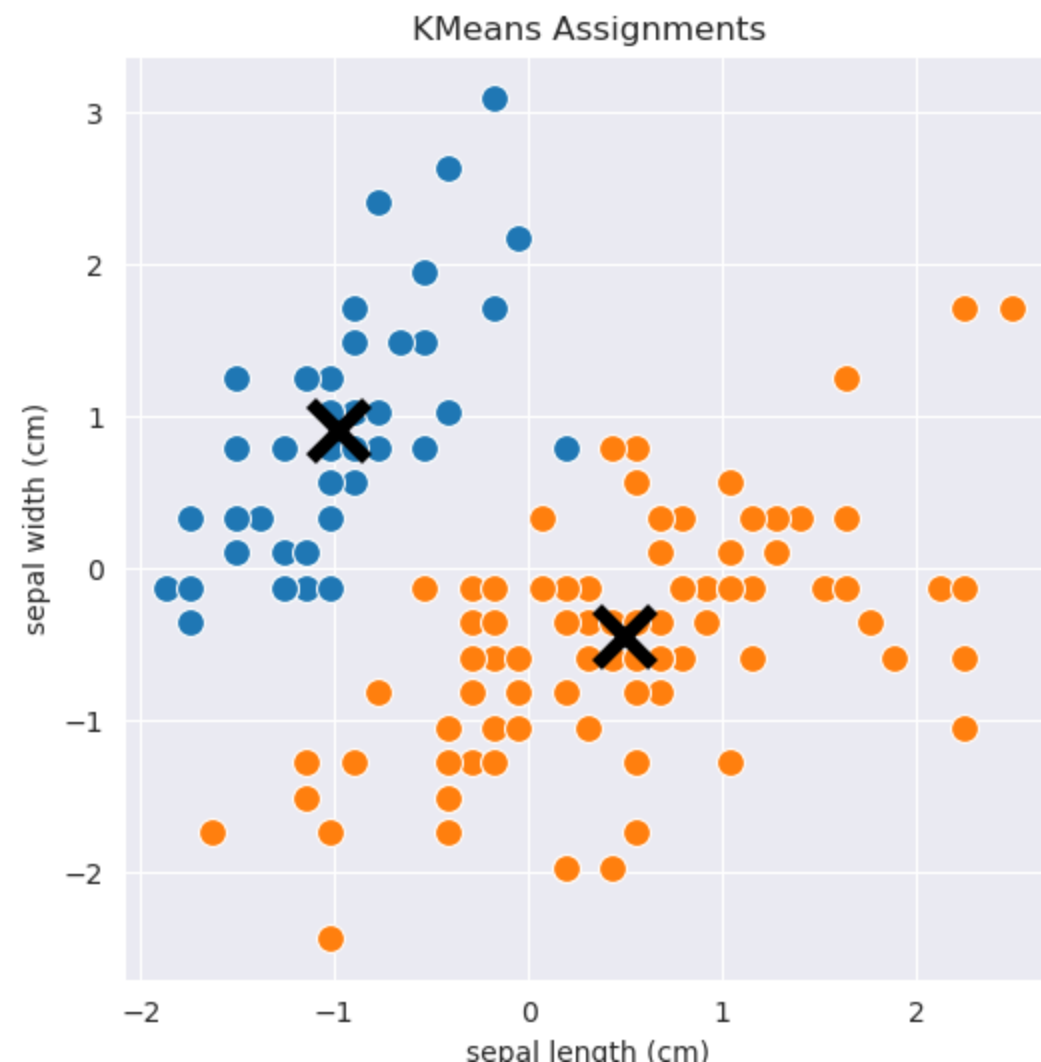
```
In [5]: # cluster centers
km.cluster_centers_.round(2)
```

Out[5]: array([[-0.98, 0.9],
 [0.49, -0.45]])

Plotting clusters and centers

Plotting clusters and centers

```
In [6]: def plot_clusters(X, c=None, km=None, title=None, ax=None, marker_size=100):  
        '''Plot data colored by cluster assignment'''  
        _, ax = plt.subplots(1, 1, figsize=(6, 6)) if ax is None else (None, ax)  
        c = km.fit_predict(X) if c is None else c  
        for i in range(np.max(c)+1):  
            sns.scatterplot(x=X.loc[c == i].iloc[:, 0], y=X.loc[c == i].iloc[:, 1], s=marker_size, ax=ax);  
            if km:  
                ax.plot(km.cluster_centers_[i, 0], km.cluster_centers_[i, 1], marker='x', c='k', ms=20, mew=5)  
        ax.set_title(title)  
  
        plot_clusters(X_iris, km=km, title="KMeans Assignments")
```



K-Means: How good are the clusters?

- One way: **Within Cluster Sum of Squared Distances (SSD)**
- How close is every point to its assigned cluster center?

$$SSD = \sum_{k=1}^K \sum_{x_i \in C_k} ||x_i - \mu_k||_2^2$$

where $||x - \mu||_2 = \sqrt{\sum_{j=1}^d (x_j - \mu_j)^2}$

- If this is high, items in cluster are far from their means.
- If this is low, items in cluster are close to their means.

K-Means: How good are the clusters?

- One way: **Within Cluster Sum of Squared Distances (SSD)**
- How close is every point to its assigned cluster center?

$$SSD = \sum_{k=1}^K \sum_{x_i \in C_k} ||x_i - \mu_k||_2^2$$

where $||x - \mu||_2 = \sqrt{\sum_{j=1}^d (x_j - \mu_j)^2}$

- If this is high, items in cluster are far from their means.
- If this is low, items in cluster are close to their means.

```
In [10]: # SSD stored in KMeans as `.inertia_`  
         round(km.inertia_,2)
```

```
Out[10]: 166.95
```

KMeans in Action

KMeans in Action

```
In [11]: import ipywidgets as widgets  
kmeans_video = widgets.Video.from_file('images/kmeans.mp4', width=750, autoplay=False, controls=True)  
kmeans_video
```

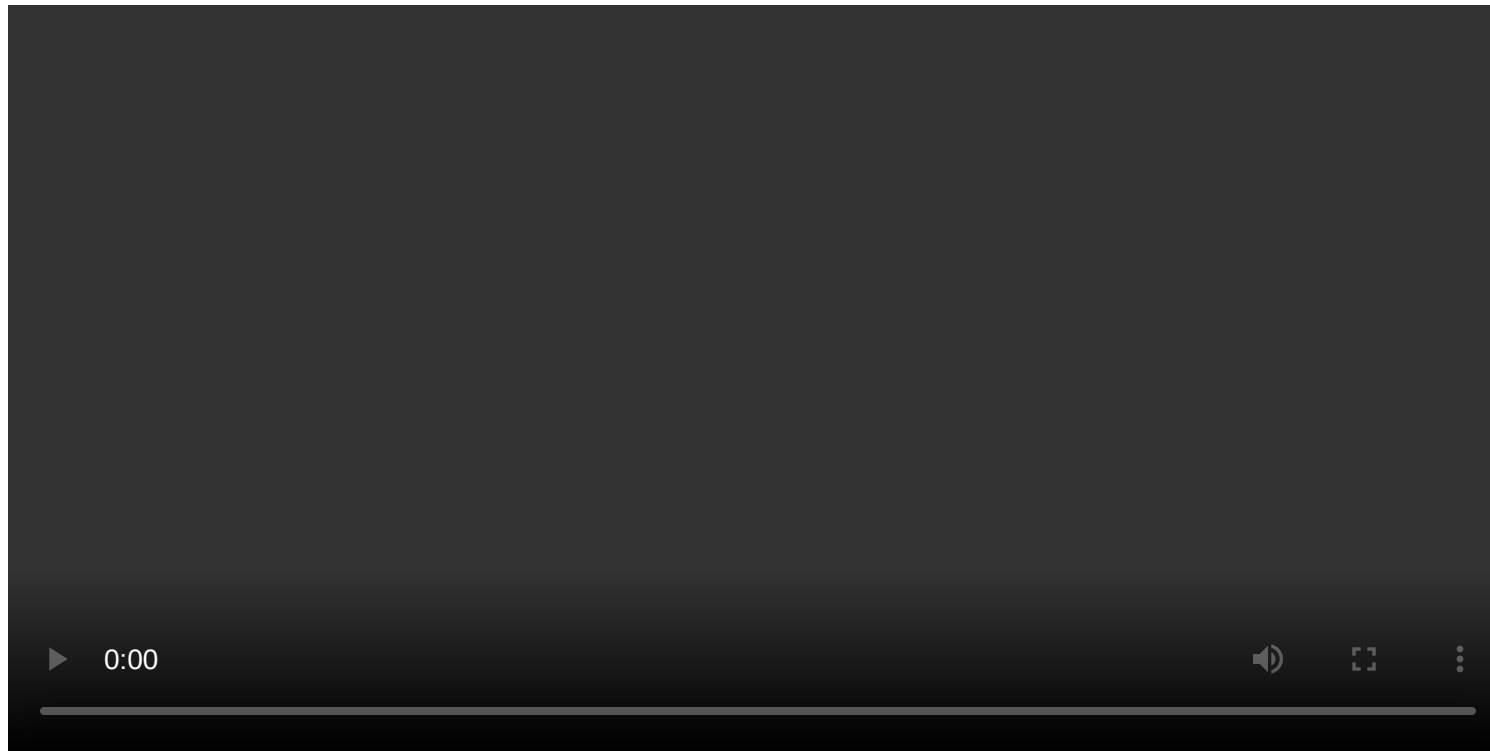
Out[11]:



KMeans in Action

```
In [11]: import ipywidgets as widgets  
kmeans_video = widgets.Video.from_file('images/kmeans.mp4', width=750, autoplay=False, controls=True)  
kmeans_video
```

Out[11]:



From <https://dashee87.github.io/data%20science/general/Clustering-with-Scikit-with-GIFs/>

Things you need to define for KMeans

- number of clusters k or `n_clusters`
- initial locations of means
 - random
 - k-means++ (pick starting points far apart from each other)

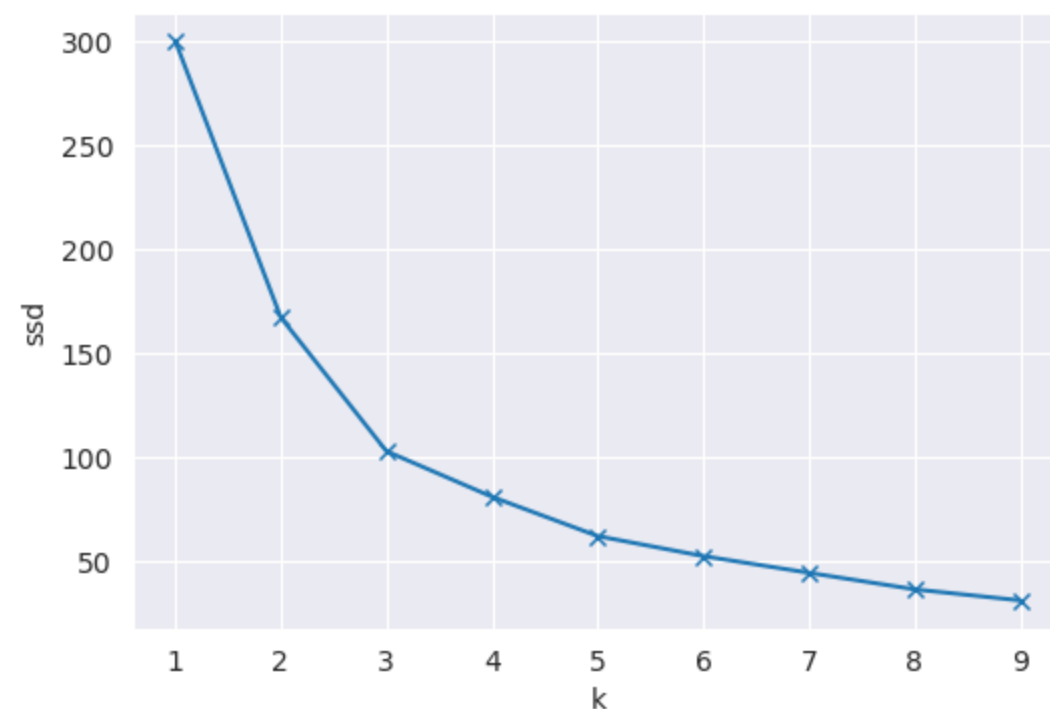
How to choose k or `n_clusters`?

- One way: use "elbow" in SSD or `KMeans.inertia_`
- "elbow" is where SSD ceases to drop rapidly

How to choose k or `n_clusters`?

- One way: use "elbow" in SSD or `KMeans.inertia_`
- "elbow" is where SSD ceases to drop rapidly

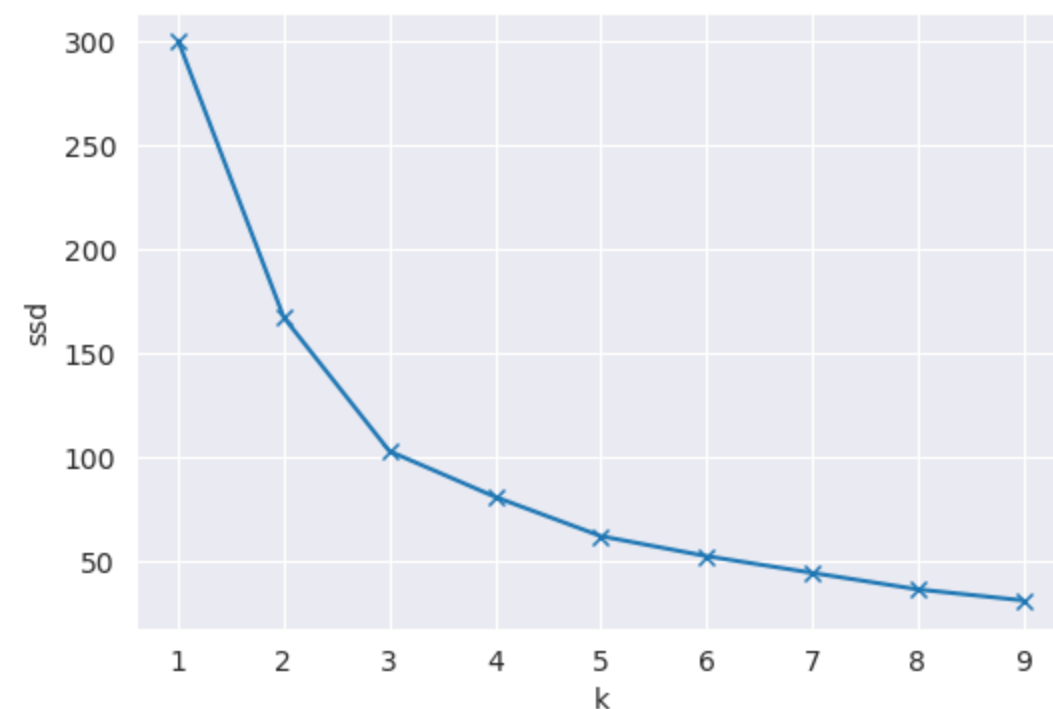
```
In [12]: ssd = []  
for i in range(1,10):  
    ssd.append(KMeans(n_clusters=i).fit(X_iris).inertia_)  
fig,ax=plt.subplots(1,1,figsize=(6,4))  
ax.plot(range(1,10),ssd,marker='x');  
ax.set_xlabel('k');ax.set_ylabel('ssd');
```



How to choose k or `n_clusters`?

- One way: use "elbow" in SSD or `KMeans.inertia_`
- "elbow" is where SSD ceases to drop rapidly

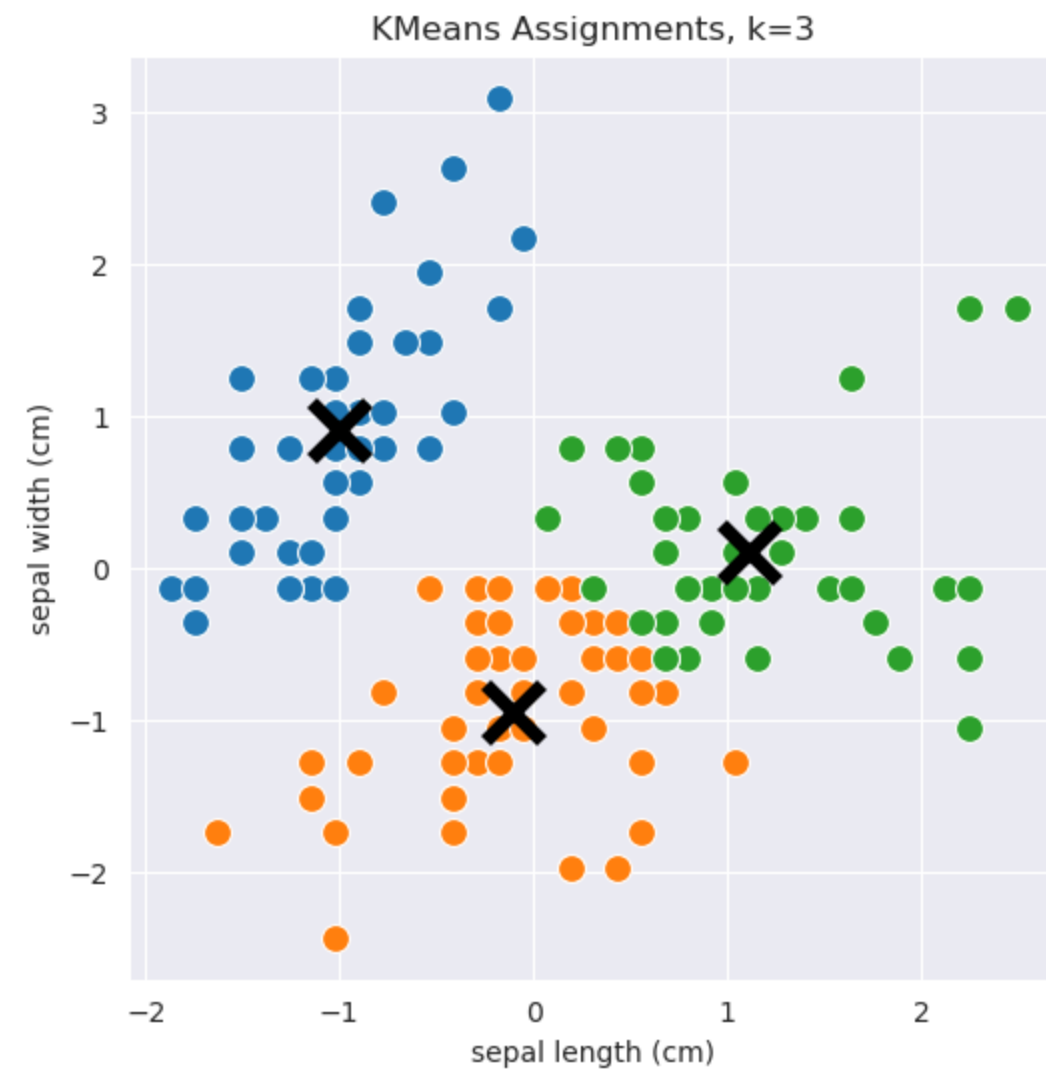
```
In [12]: ssd = []
for i in range(1,10):
    ssd.append(KMeans(n_clusters=i).fit(X_iris).inertia_)
fig,ax=plt.subplots(1,1,figsize=(6,4))
ax.plot(range(1,10),ssd,marker='x');
ax.set_xlabel('k');ax.set_ylabel('ssd');
```



Refitting with $k=3$

Refitting with k=3

```
In [13]: plot_clusters(X_iris, km=KMeans(n_clusters=3, random_state=0), title="KMeans Assignments, k=3")
```



KMeans: Another Example

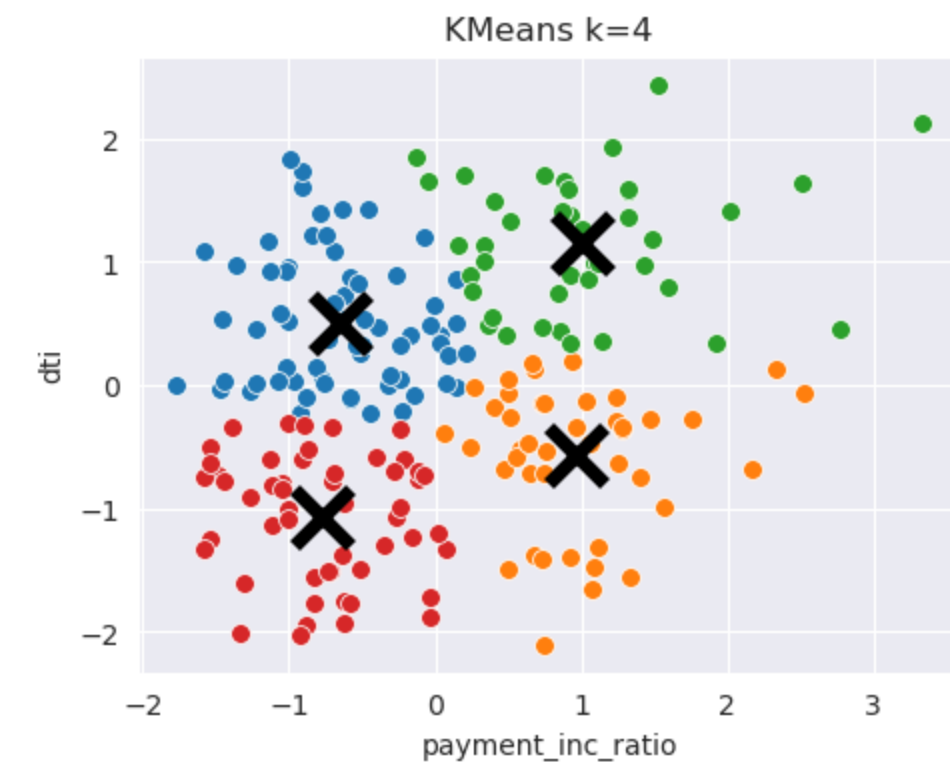
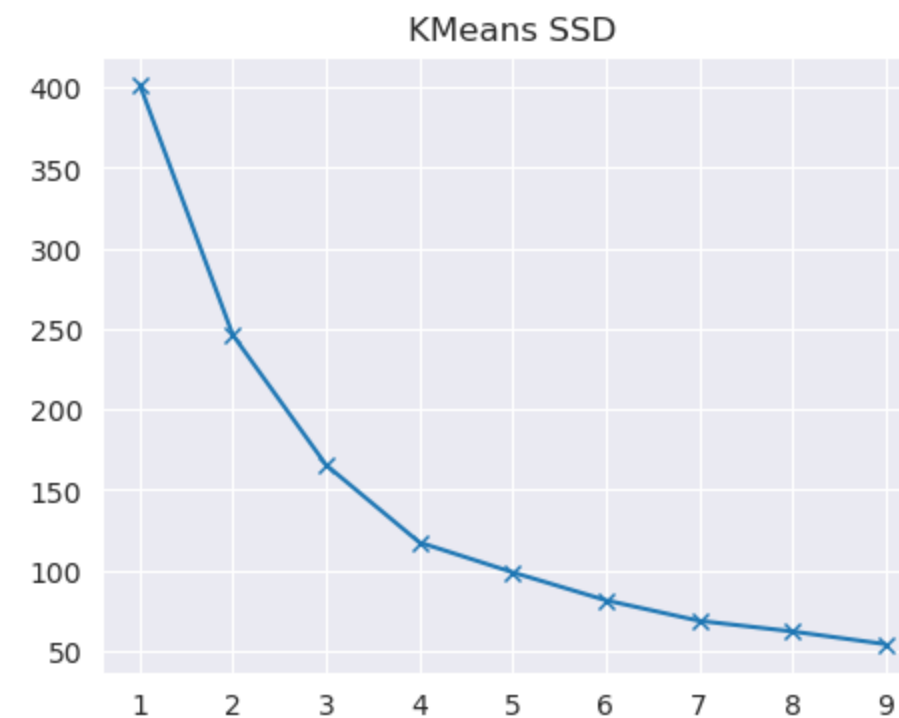
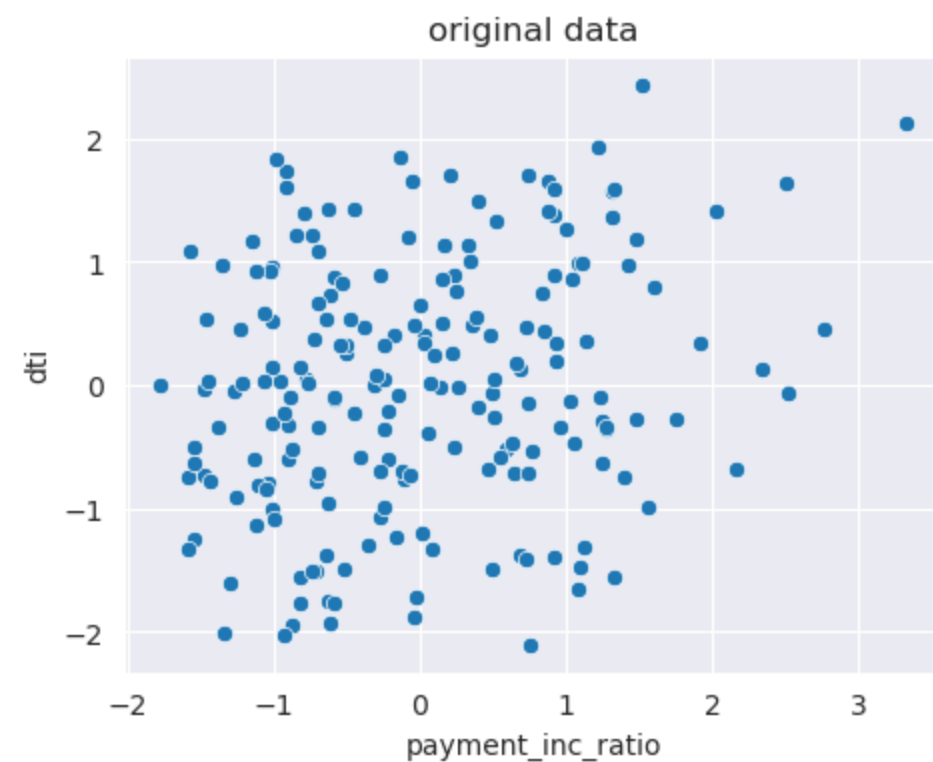
KMeans: Another Example

```
In [15]: # loading and plotting the data
data = pd.read_csv('../data/loan200.csv')[['payment_inc_ratio', 'dti']]
from sklearn.preprocessing import StandardScaler
X_loan = pd.DataFrame(StandardScaler().fit_transform(data), columns=data.columns)

fig, ax = plt.subplots(1, 3, figsize=(18, 4))
sns.scatterplot(x=X_loan.iloc[:, 0], y=X_loan.iloc[:, 1], ax=ax[0]);
ax[0].set_title('original data');

ssd = [KMeans(n_clusters=i).fit(X_loan).inertia_ for i in range(1, 10)]
ax[1].plot(range(1, 10), ssd, marker='x');
ax[1].set_title('KMeans SSD');

plot_clusters(X_loan, km=KMeans(n_clusters=4, random_state=0), title='KMeans k=4', marker_size=50, ax=ax[2])
```



KMeans: Synthetic Example

KMeans: Synthetic Example

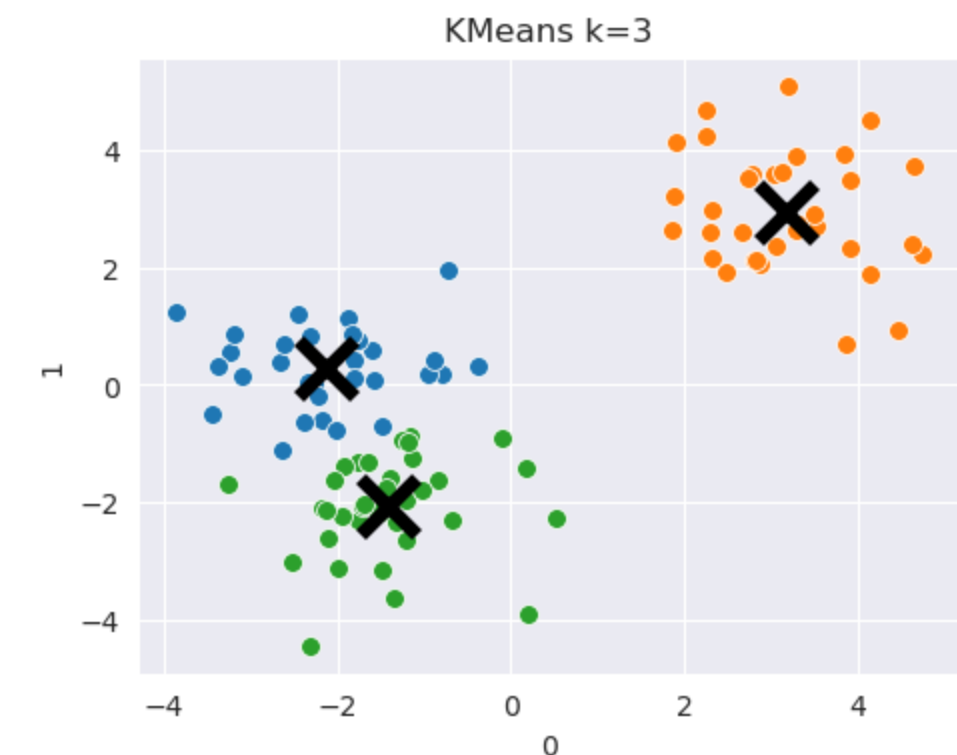
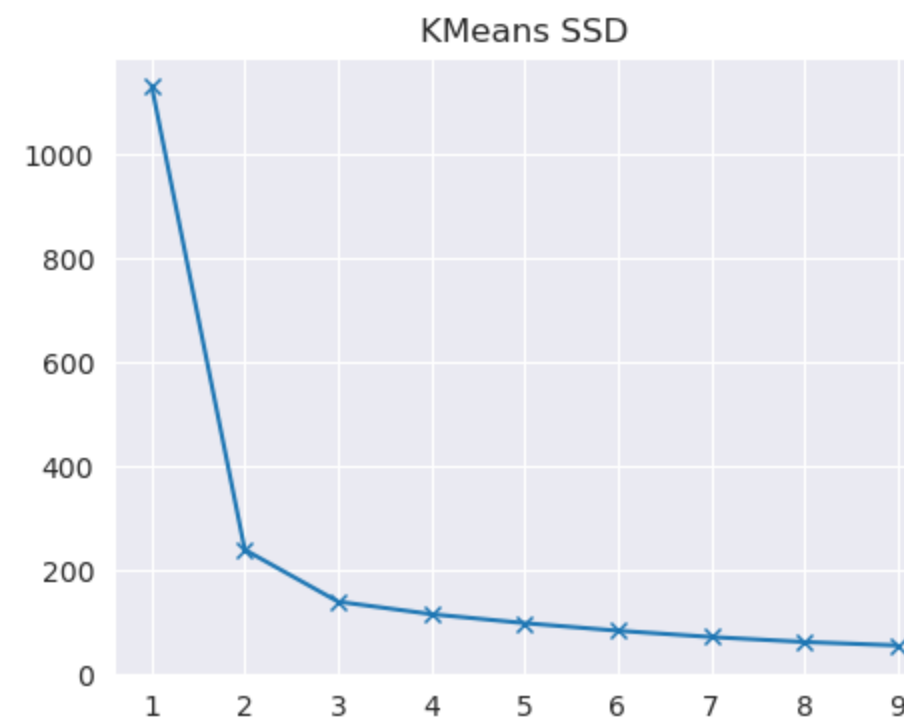
```
In [16]: from sklearn.datasets import make_blobs
X_blobs,y_blobs = make_blobs(centers=[[3,3],(-2,0),(-2,-2)],random_state=1)
X_blobs = pd.DataFrame(X_blobs)

fig,ax = plt.subplots(1,3,figsize=(18,4))

sns.scatterplot(x=X_blobs.iloc[:,0],y=X_blobs.iloc[:,1],ax=ax[0]);
ax[0].set_title('original data');

ssd = [KMeans(n_clusters=i).fit(X_blobs).inertia_ for i in range(1,10)]
ax[1].plot(range(1,10),ssd,marker='x');
ax[1].set_title('KMeans SSD')

plot_clusters(X_blobs,km=KMeans(n_clusters=3, random_state=0),title='KMeans k=3',marker_size=50,ax=ax[2])
```



Hierarchical Agglomerative Clustering (HAC)

- group clusters together from the bottom up
- don't have to specify number of clusters up front
- generates binary tree over data

HAC: How it works

FIRST: every point is it's own cluster

A: Find pair of clusters that are "closest"

B: Merge into single cluster

GOTO A and Repeat till there is a single cluster

HAC in Action

HAC in Action

```
In [17]: import ipywidgets as widgets  
         hac_video = widgets.Video.from_file('images/hac.mp4', width=750, autoplay=False, controls=True)  
         hac_video
```

Out[17]:



HAC in Action

```
In [17]: import ipywidgets as widgets  
         hac_video = widgets.Video.from_file('images/hac.mp4', width=750, autoplay=False, controls=True)  
         hac_video
```

Out[17]:



From <https://dashee87.github.io/data%20science/general/Clustering-with-Scikit-with-GIFs/>

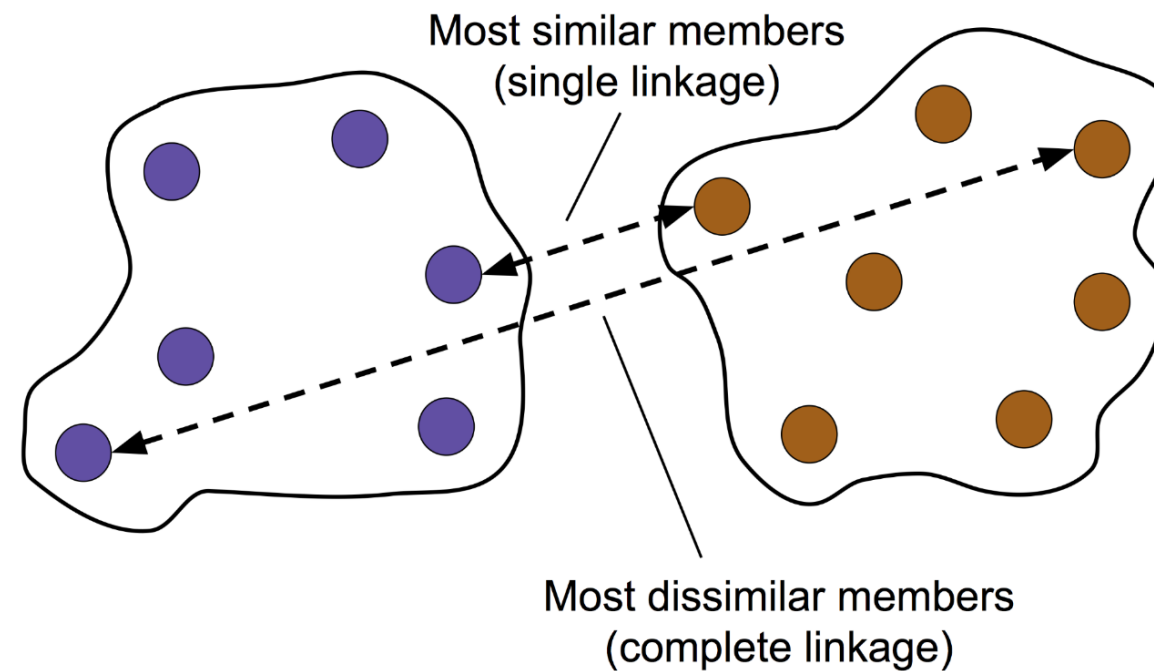
What is "close"?

- Need to define what we mean by "closeness" by choosing
 - distance metric (how to measure distance)
 - linkage criteria (how to compare clusters)

Need to define: Distance Metric

- **Euclidean**: $\sqrt{\sum_{i=1}^n (a_i - b_i)^2}$
 - easy to use analytically, sensitive to outliers
- **Manhattan**: $\sum_{i=1}^n |a_i - b_i|$
 - more difficult to use analytically, robust to outliers
- **Cosine**: $1 - \frac{\sum a_i b_i}{\|a_i\|_2 \|b_i\|_2}$
 - angle between vectors while ignoring their scale
- many more (see <https://numerics.mathdotnet.com/Distance.html>)

Need to define: Linkage



single : shortest distance from item of one cluster to item of the other

complete : greatest distance from item of one cluster to item of the other

average : average distance of items in one cluster to items in the other

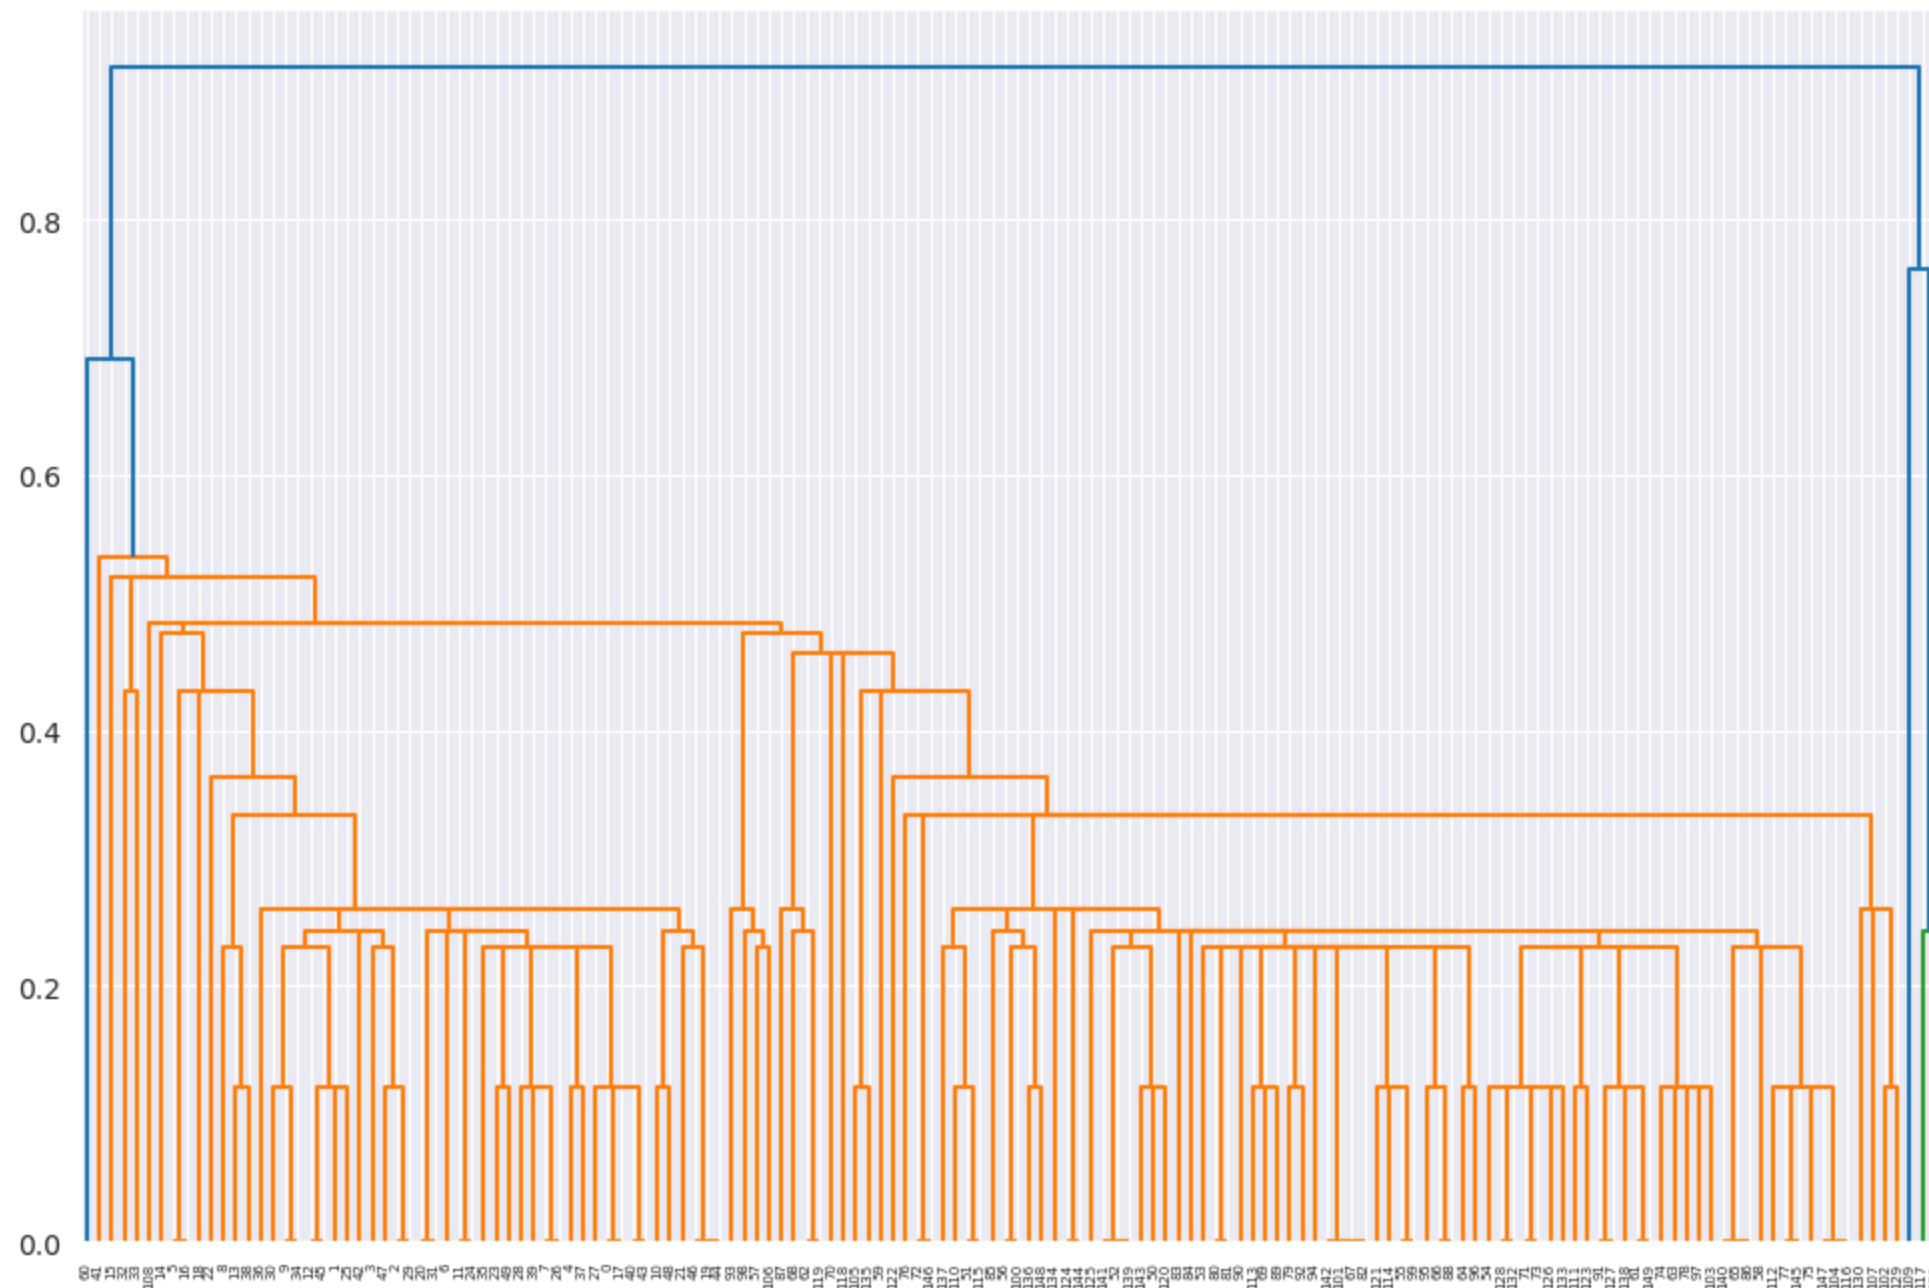
ward : minimize variance of clusters being merged (only euclidean metric)

HAC and Dendrograms: Single Linkage

HAC and Dendrograms: Single Linkage

```
In [18]: # nice helper function for creating a dendrogram
from scipy.cluster import hierarchy

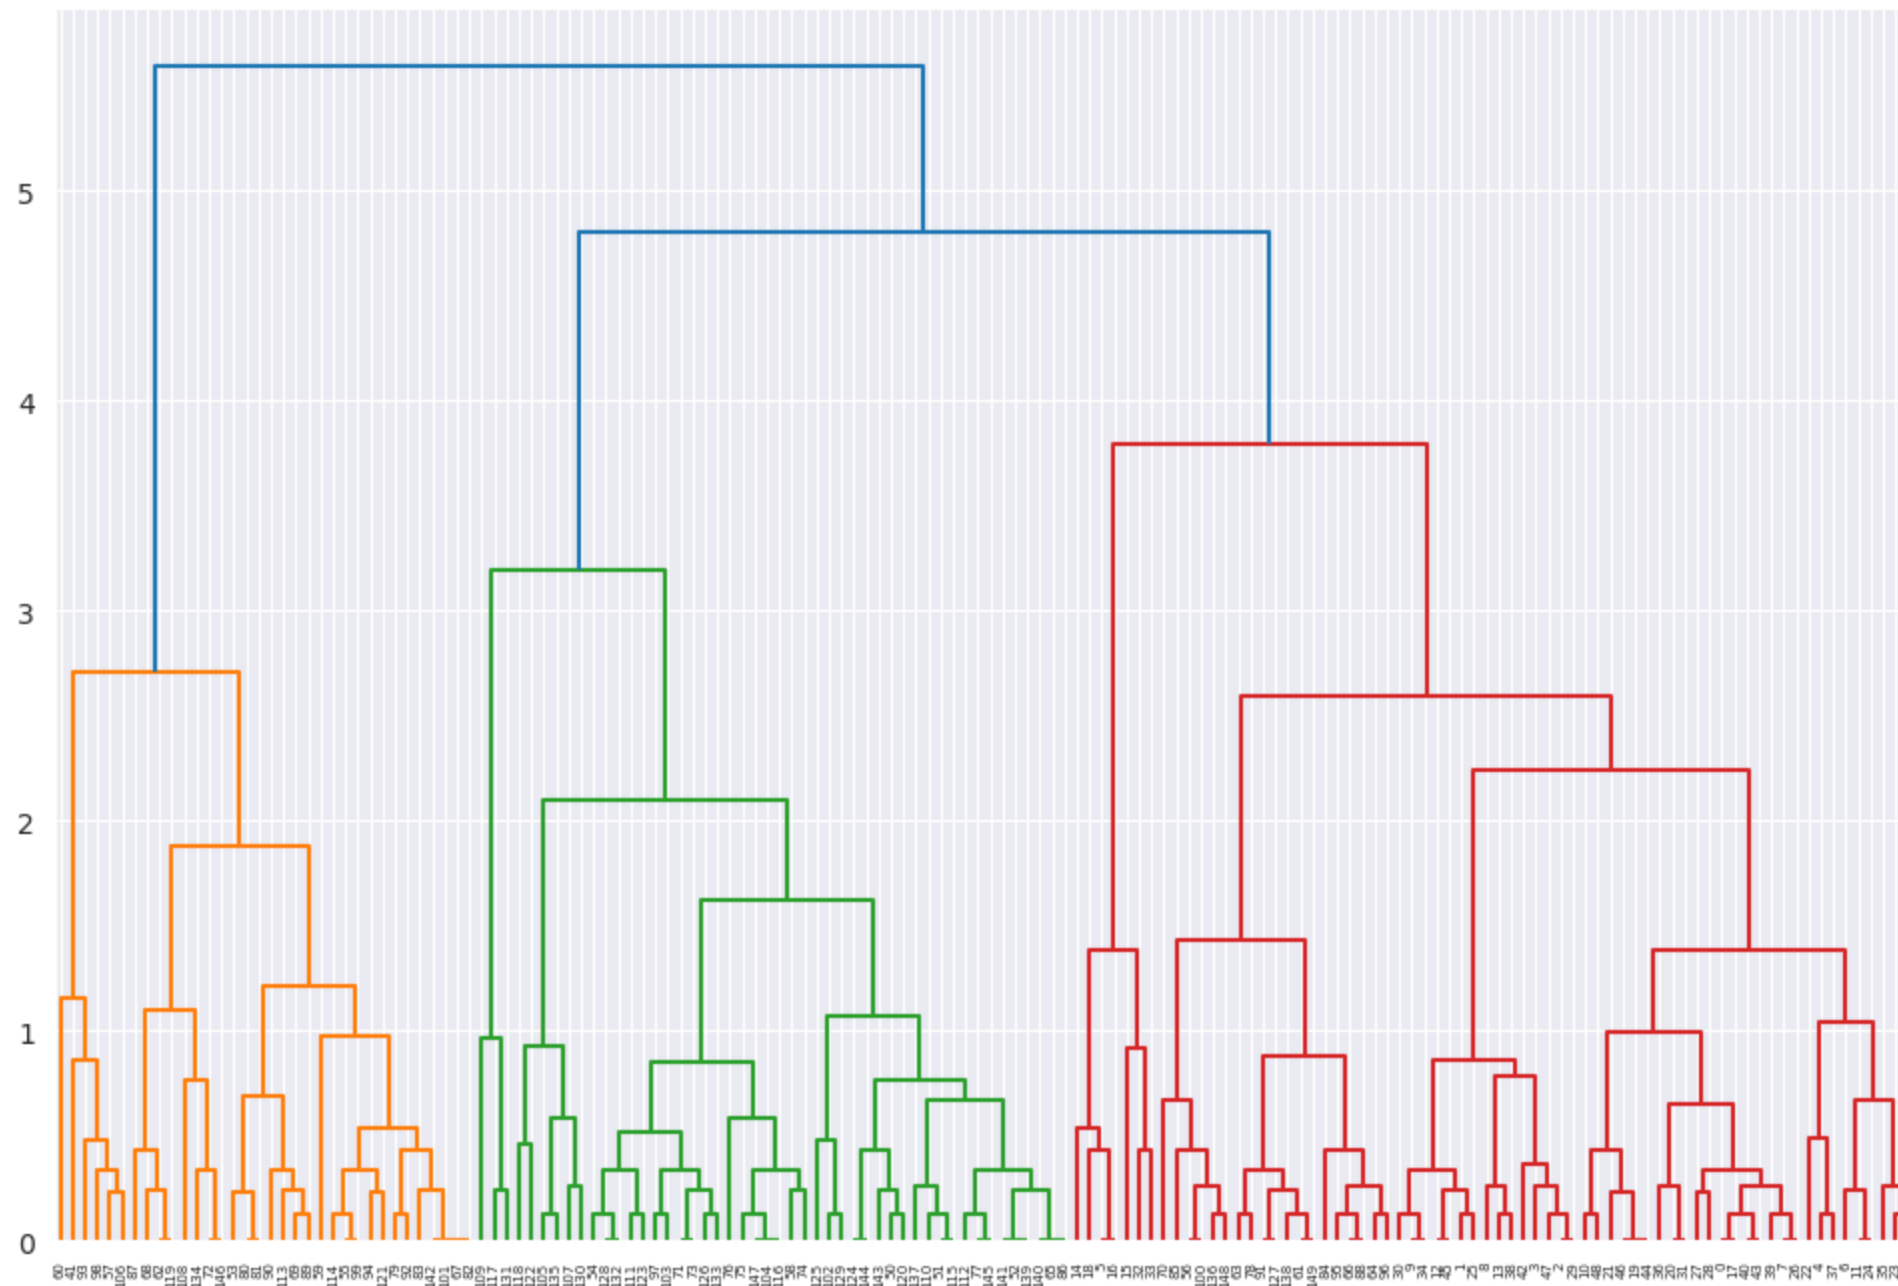
Z = hierarchy.linkage(X_iris, 'single') # metric = 'euclidean' by default
fig = plt.figure(figsize=(12,8)); hierarchy.dendrogram(Z);
```



HAC and Dendrograms: Complete Linkage

HAC and Dendrograms: Complete Linkage

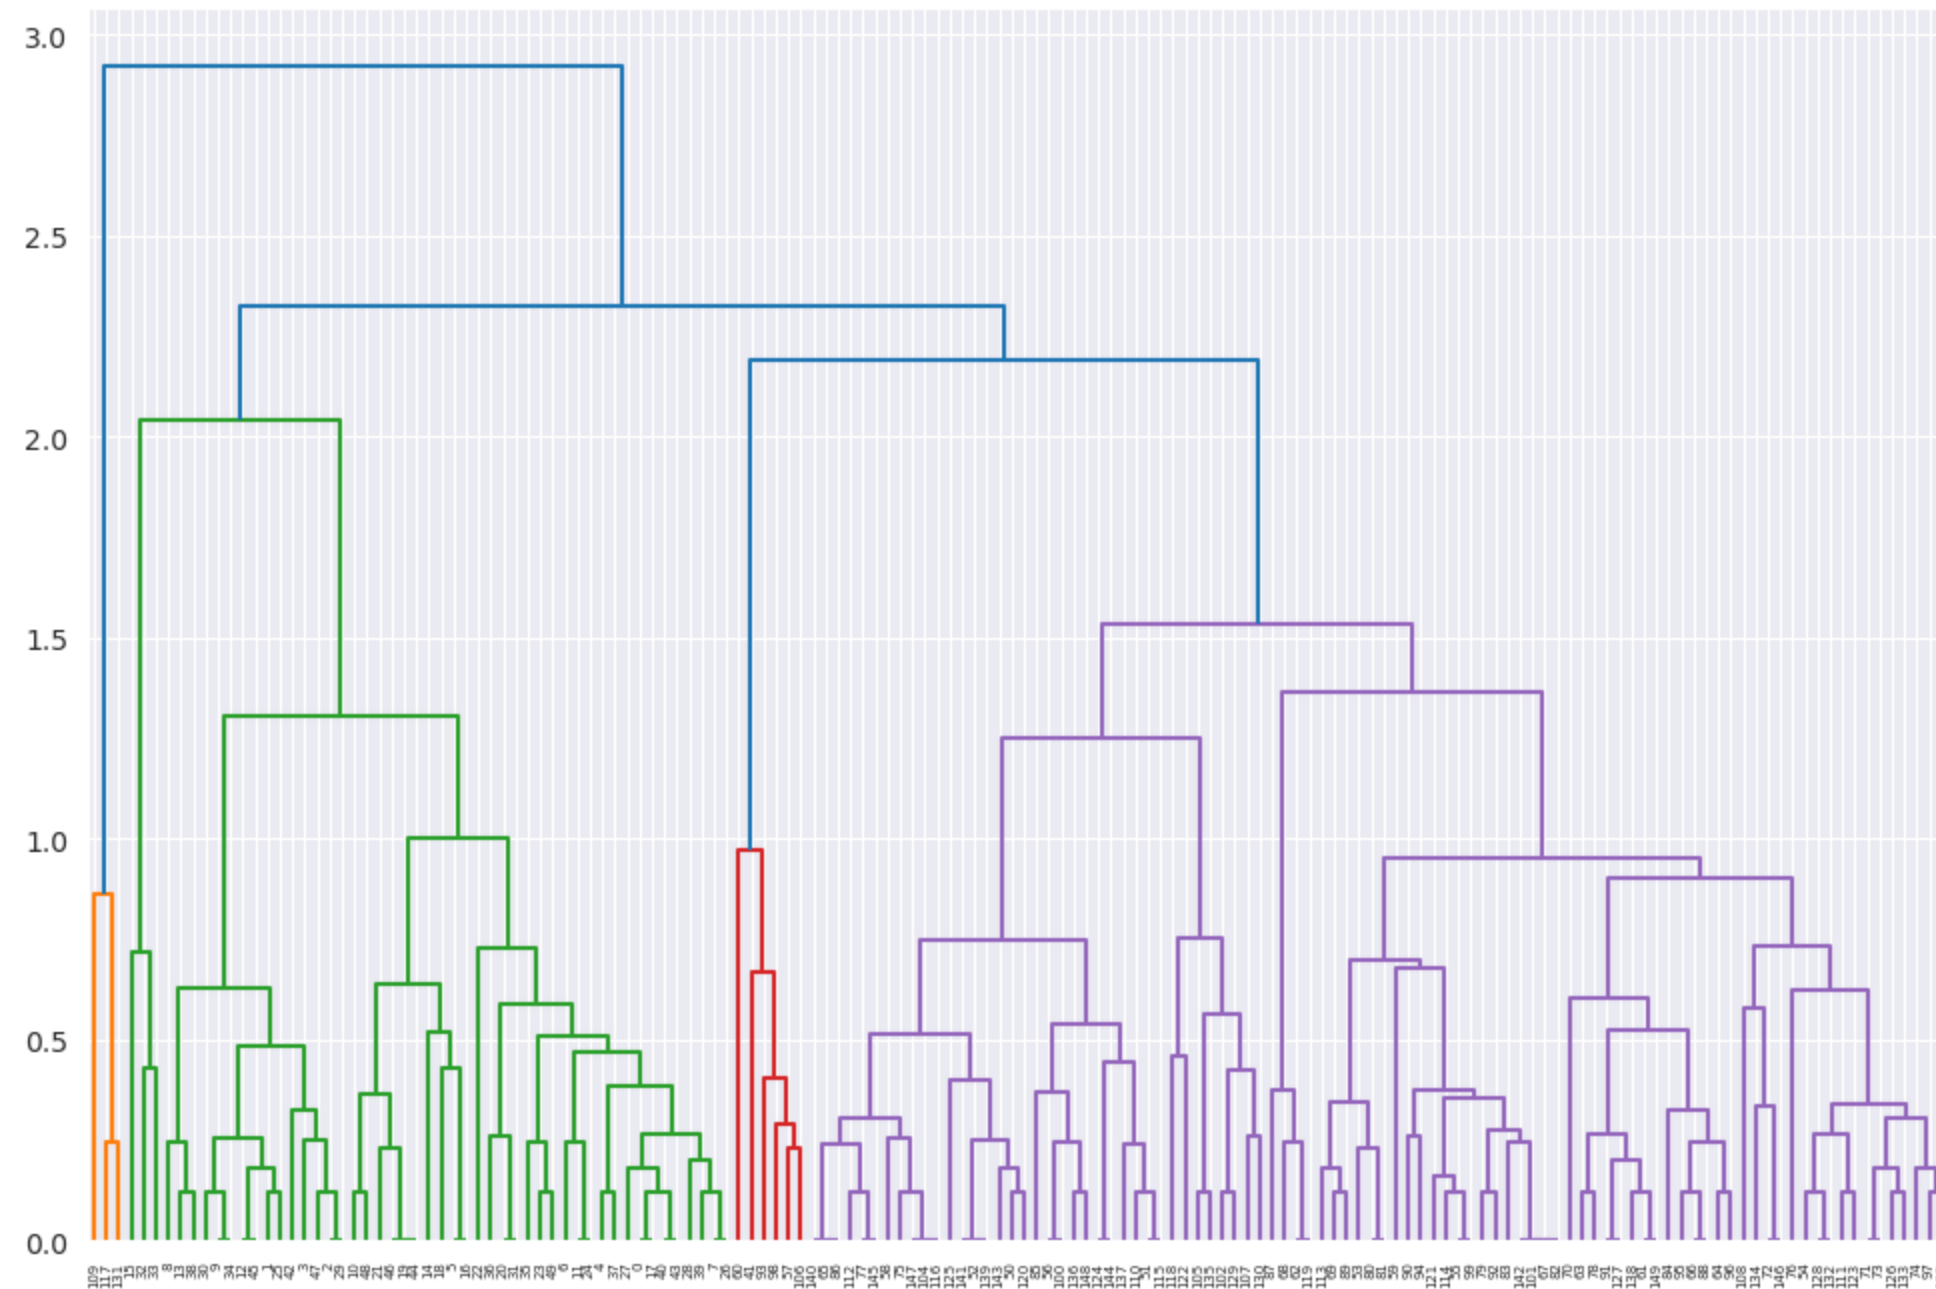
```
In [19]: Z = hierarchy.linkage(X_iris, 'complete')  
fig = plt.figure(figsize=(12,8)); hierarchy.dendrogram(Z);
```



HAC and Dendrograms: Average Linkage

HAC and Dendrograms: Average Linkage

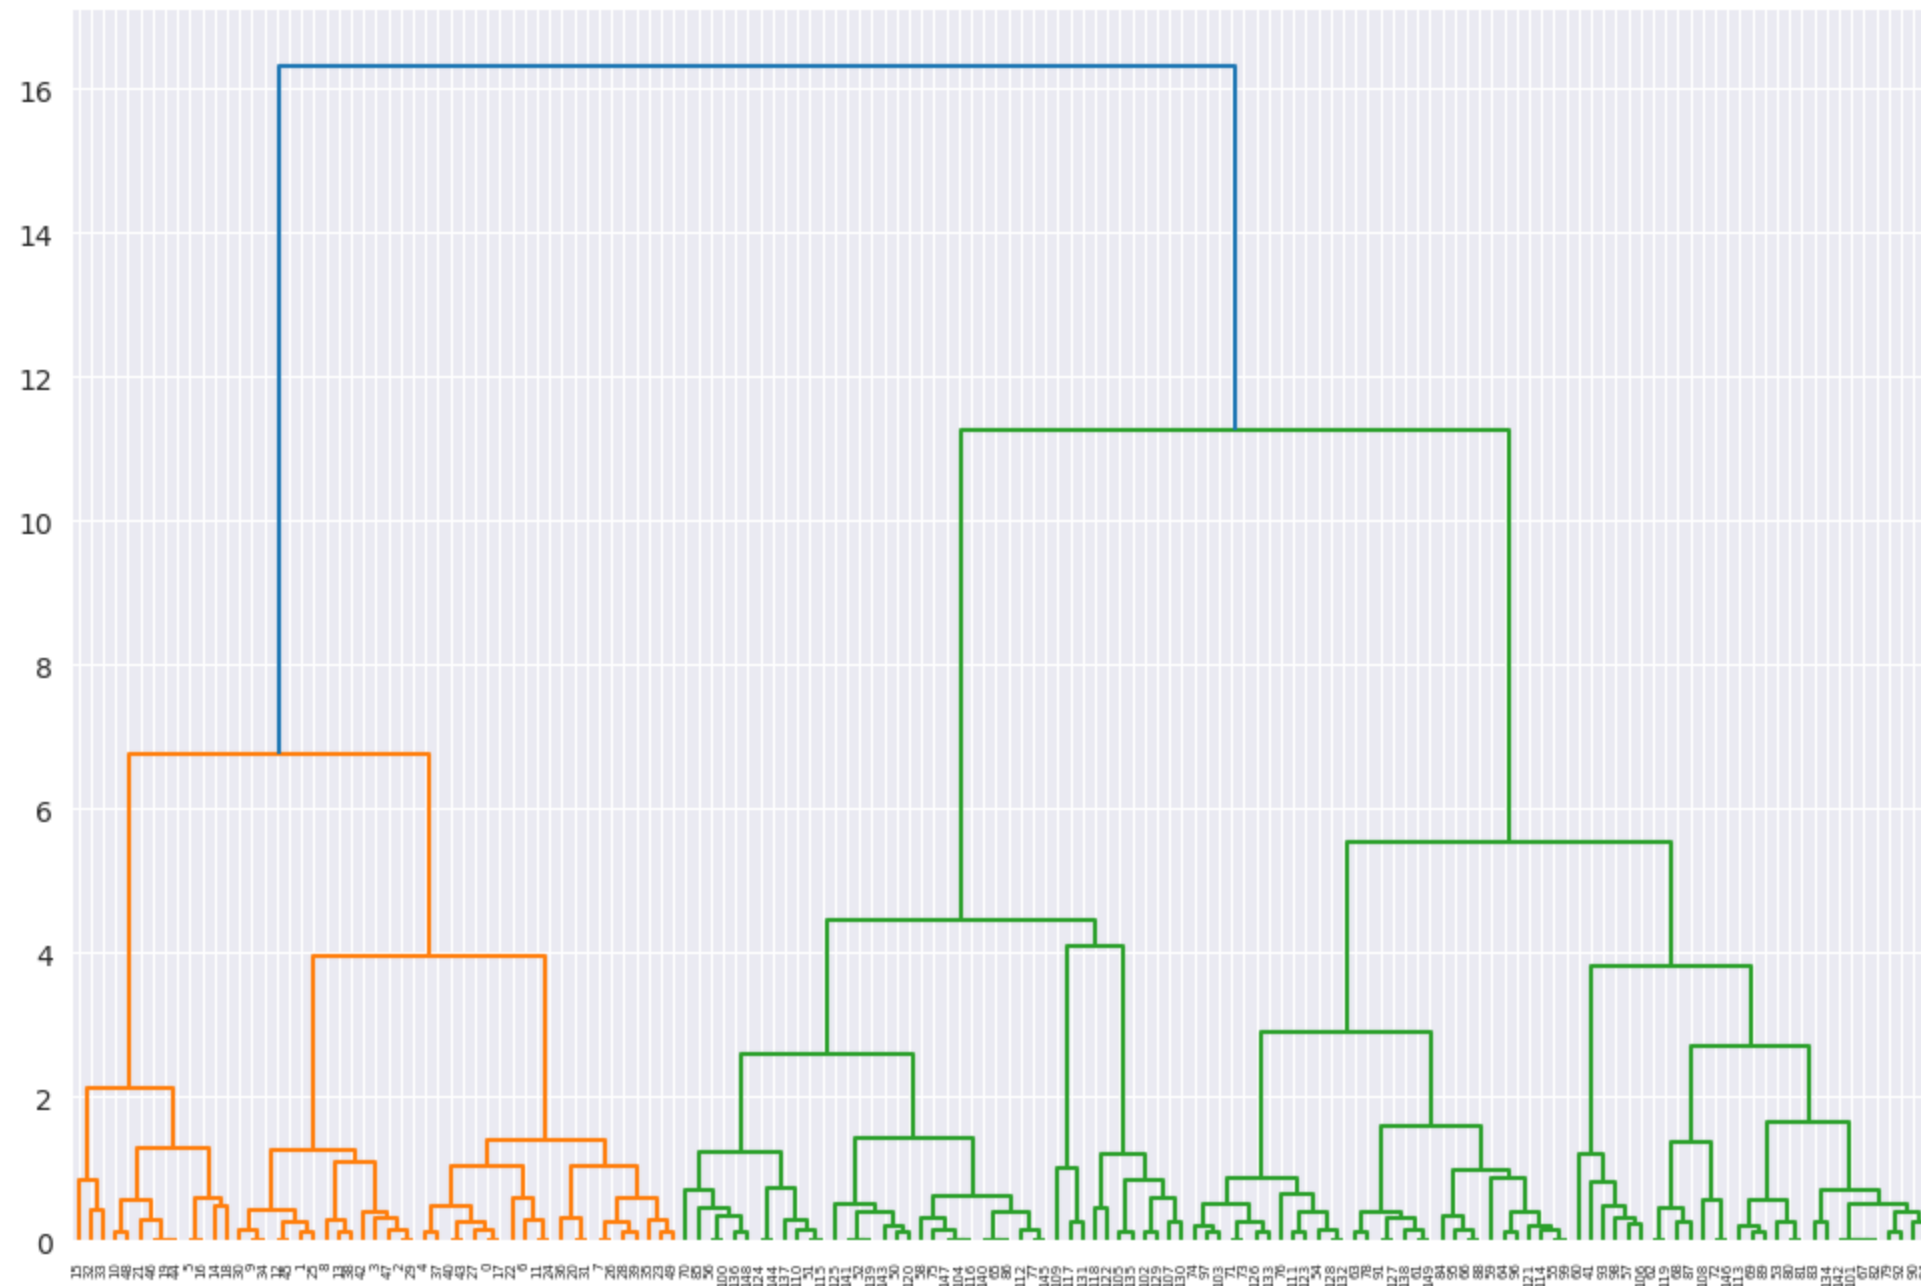
```
In [20]: Z = hierarchy.linkage(X_iris, 'average')  
fig = plt.figure(figsize=(12,8)); hierarchy.dendrogram(Z);
```



HAC and Dendrograms: Ward Linkage

HAC and Dendrograms: Ward Linkage

```
In [21]: Z = hierarchy.linkage(X_iris, 'ward')  
fig = plt.figure(figsize=(12,8)); hierarchy.dendrogram(Z);
```



HAC in sklearn

HAC in sklearn

```
In [22]: from sklearn.cluster import AgglomerativeClustering

         hac = AgglomerativeClustering(linkage='single',      # ward by default
                                       affinity='euclidean',  # default
                                       n_clusters=4)          # 2 by default
         c_single = hac.fit_predict(X_iris)

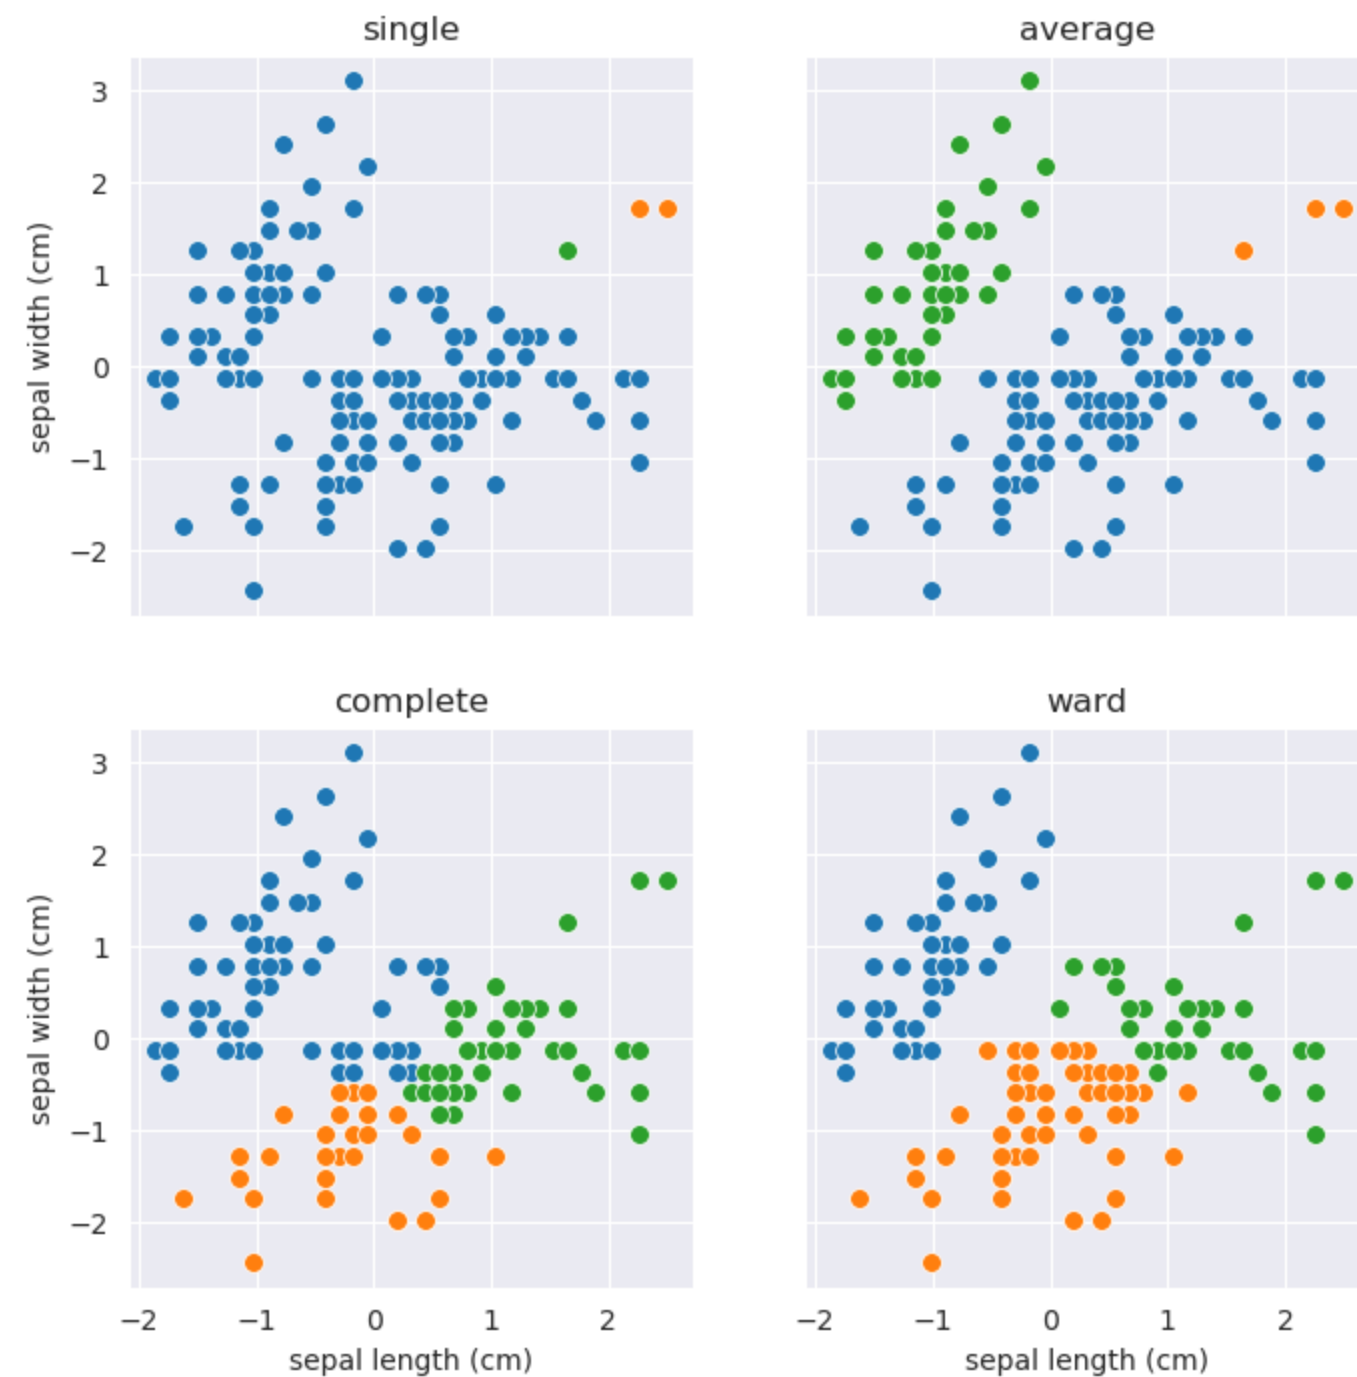
         # generate models and assignments for all linkages
         models,assignments = [],[]
         linkages = ['single', 'average', 'complete', 'ward']
         for linkage in linkages:
             models.append(AgglomerativeClustering(linkage=linkage, affinity='euclidean', n_clusters=3))
             assignments.append(models[-1].fit_predict(X_iris))

         # plot on the next slide
```

HAC in sklearn

HAC in sklearn

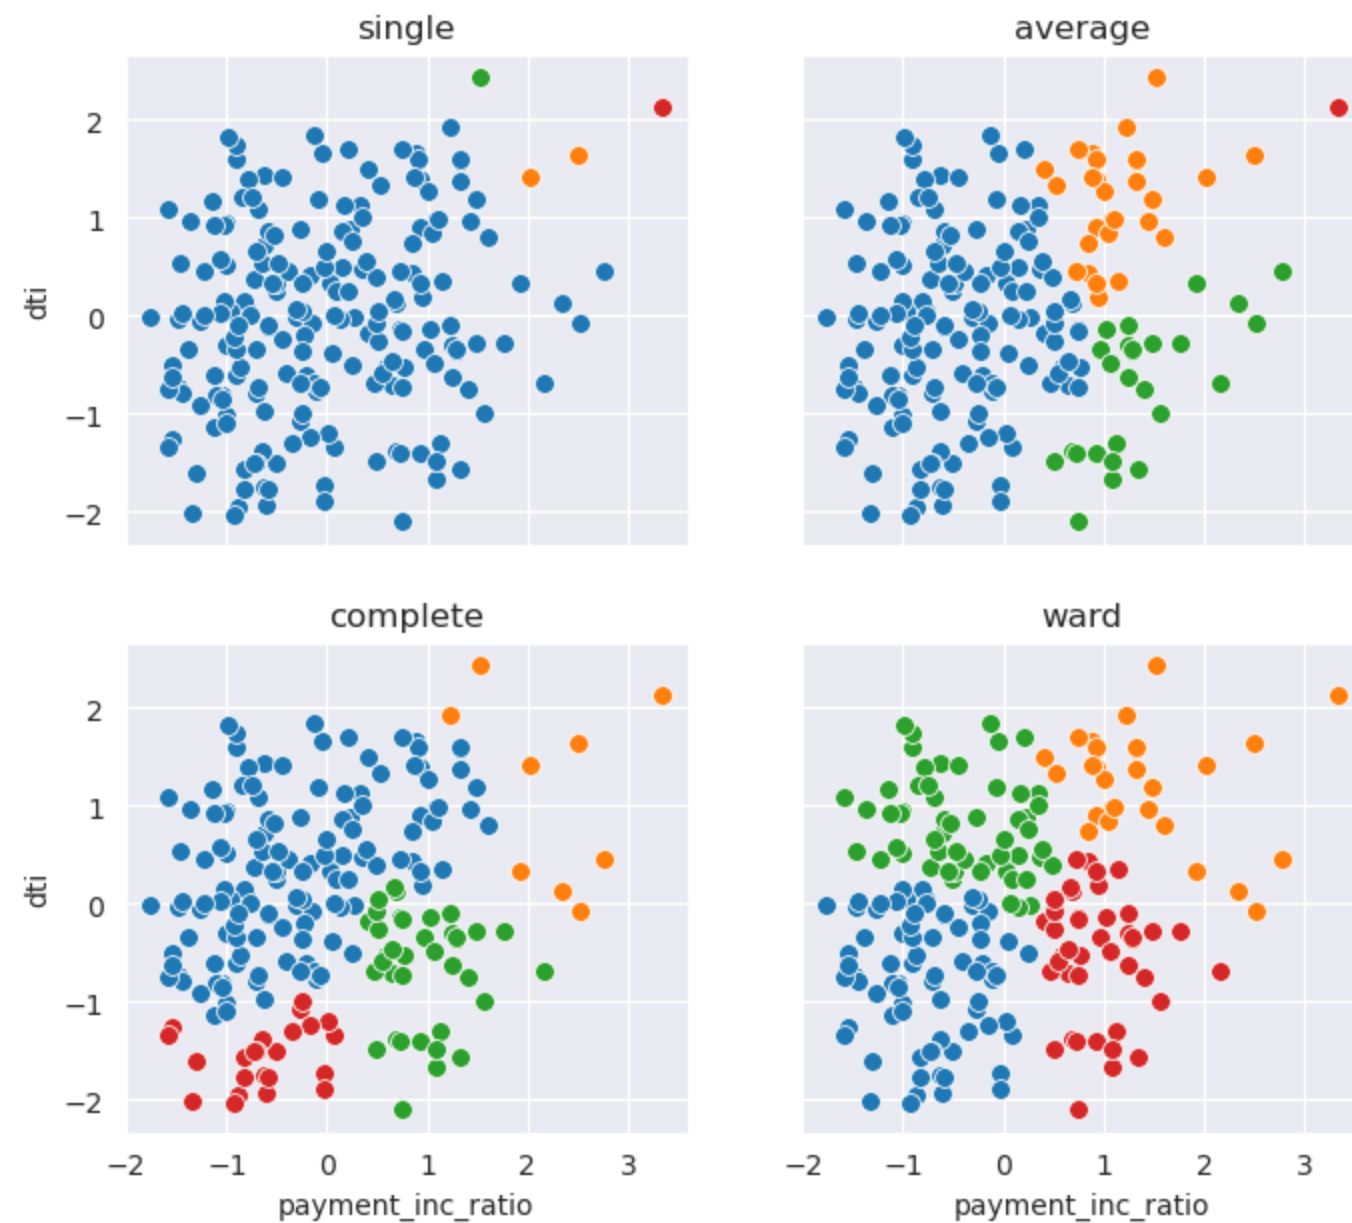
```
In [23]: fig, ax = plt.subplots(2, 2, figsize=(8, 8), sharex=True, sharey=True)
         axs = ax.flatten()
         for i in range(len(linkage)):
             plot_clusters(X_iris, c=assignments[i], title=linkages[i], ax=axs[i], marker_size=50)
```



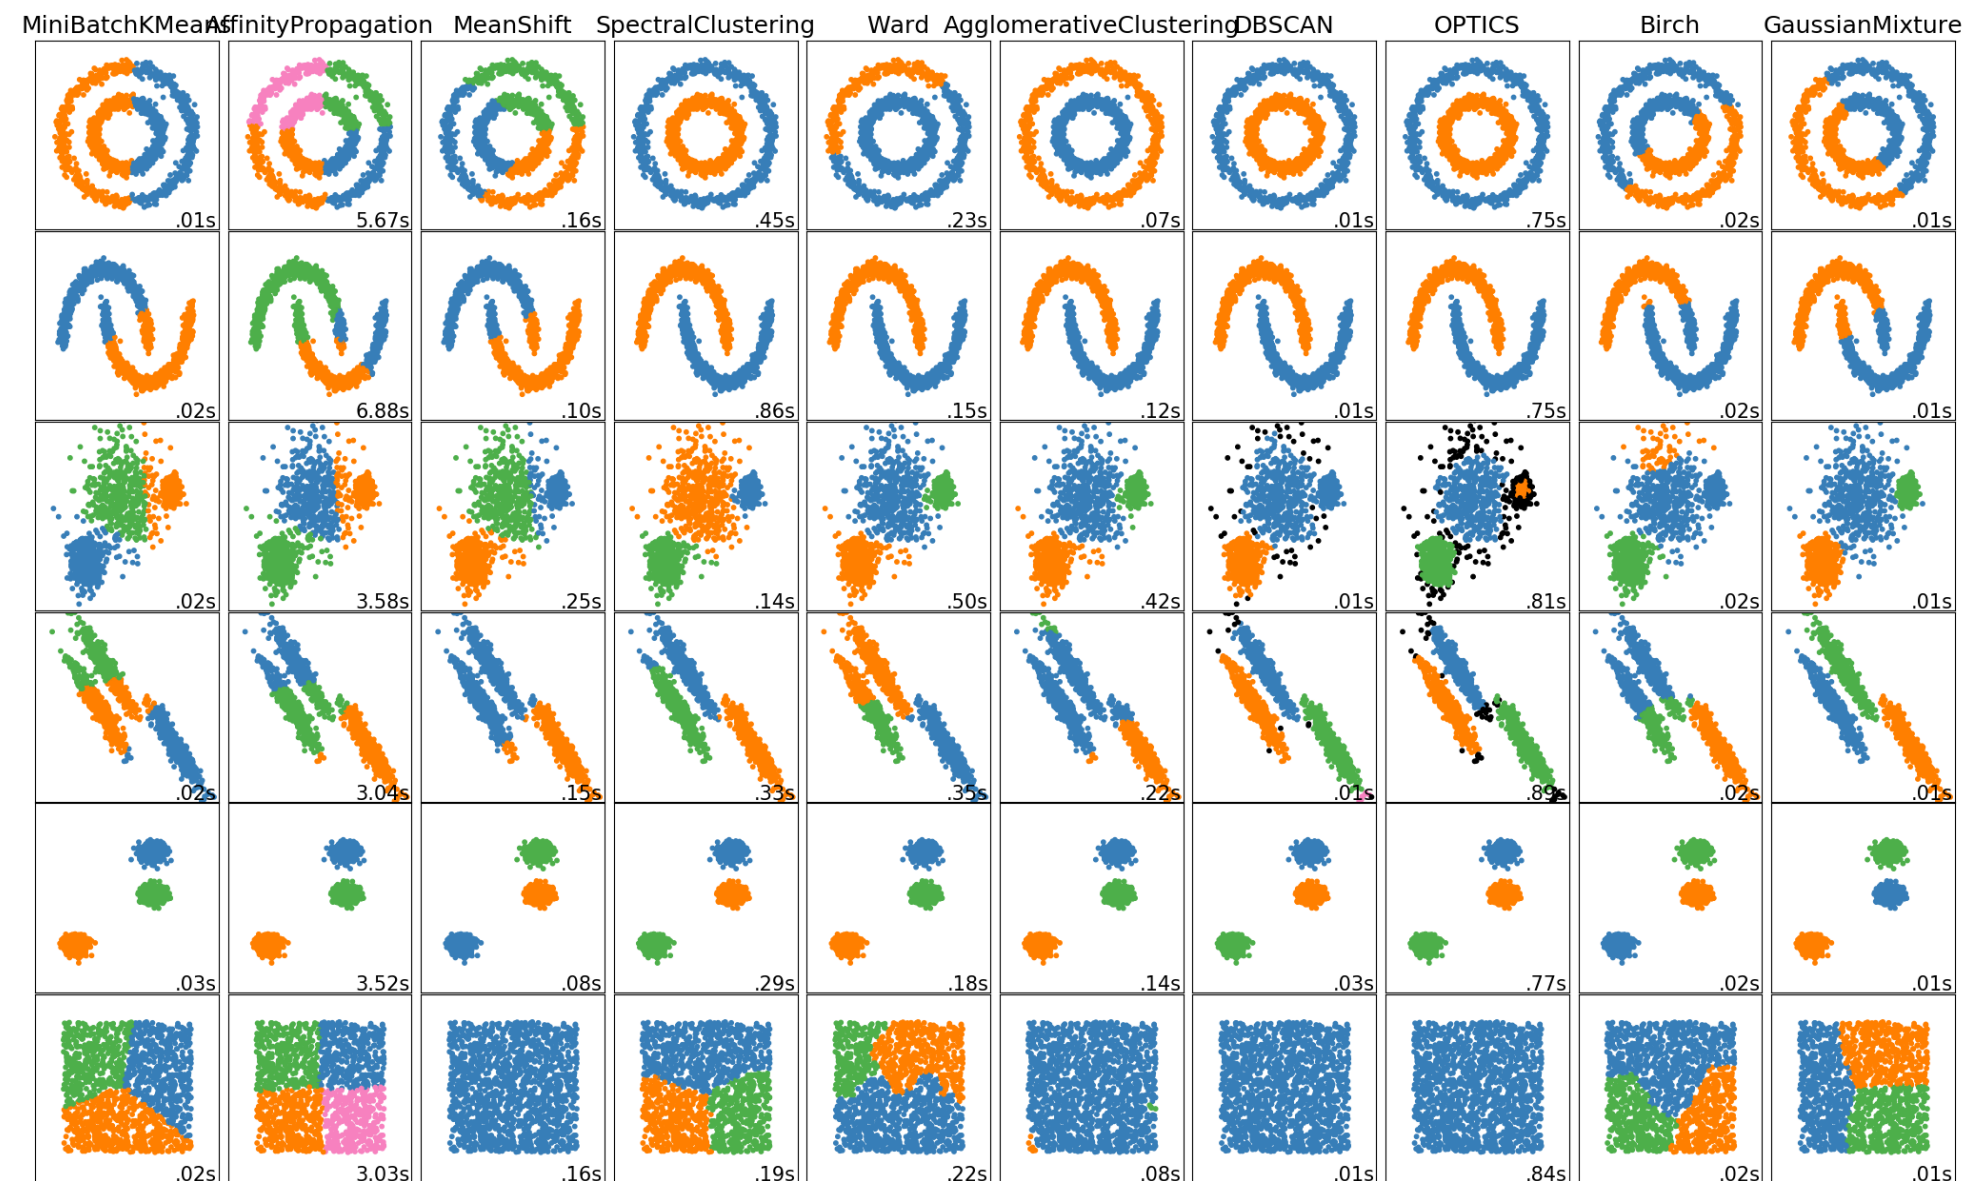
HAC: Another Example

HAC: Another Example

```
In [24]: models,assignments,linkages = [],[],['single','average','complete','ward']
for linkage in linkages:
    models.append(AgglomerativeClustering(linkage=linkage,affinity='euclidean',n_clusters=4))
    assignments.append(models[-1].fit_predict(X_loan))
fig,ax = plt.subplots(2,2,figsize=(8,7),sharex=True,sharey=True)
axs = ax.flatten()
for i in range(len(linkage)):
    plot_clusters(X_loan,assignments[i],title=linkages[i],ax=axs[i],marker_size=50)
```



Clustering: Many Other Methods



From <https://scikit-learn.org/stable/modules/clustering.html>

How to evaluate clustering?

- **Within Cluster Sum of Squared Distances (SSD)**
- If we have labels:
 - **Homogeneity:** each cluster contains only members of a single class
 - **Completeness:** all members of a given class are assigned to the same cluster
 - **V-score:** harmonic mean of Homogeneity and Completeness
- Silhouette plots (see PML)
- many others ([see sklearn](#))

Clustering Review

- k-Means
- Heirarchical Agglomerative Clustering
 - linkages
 - distance metrics
- Evaluating

Questions re Clustering?

Recommendation Engines

- Given a user and a set of items to recommend (or rank):
 - **Content-Based Filtering:** Recommend things **similar to the things I've liked**
 - **Collaborative Filtering:** Recommend things **that people with similar tastes have liked**
 - Hybrid/Ensemble
 - Recommendation as Classification

Example: Housing Data

Example: Housing Data

```
In [25]: df_house = pd.read_csv('../data/house_sales_subset.csv')
df_house = df_house.iloc[:10].loc[:, ['SqFtTotLiving', 'SqFtLot', 'AdjSalePrice']]
X_house_scaled = StandardScaler().fit_transform(df_house)
df_house_scaled = pd.DataFrame(X_house_scaled, columns=['SqFtTotLiving_scaled', 'SqFtLot_scaled', 'AdjSalePrice_scaled'])
df_house_scaled.head().round(2)
```

Out[25]:

| | SqFtTotLiving_scaled | SqFtLot_scaled | AdjSalePrice_scaled |
|---|----------------------|----------------|---------------------|
| 0 | 0.40 | -0.47 | -0.70 |
| 1 | 2.03 | 0.65 | 2.48 |
| 2 | -0.01 | 1.26 | 1.19 |
| 3 | 1.36 | -0.54 | -0.12 |
| 4 | -0.41 | -0.54 | -0.71 |

Content-Based Filtering

- Find other things similar to the things I've liked
- Assume: If I like product A, and product B is like product A, I'll like product B
- Use similarity of items
- Matrix: items x items
- Values: Similarity of items

Calculate Distances

- to maximize similarity \rightarrow minimize distance

Calculate Distances

- to maximize similarity → minimize distance

```
In [26]: # using euclidean distance
from sklearn.metrics.pairwise import euclidean_distances

# calculate all pairwise distances between houses
dists = euclidean_distances(X_house_scaled)

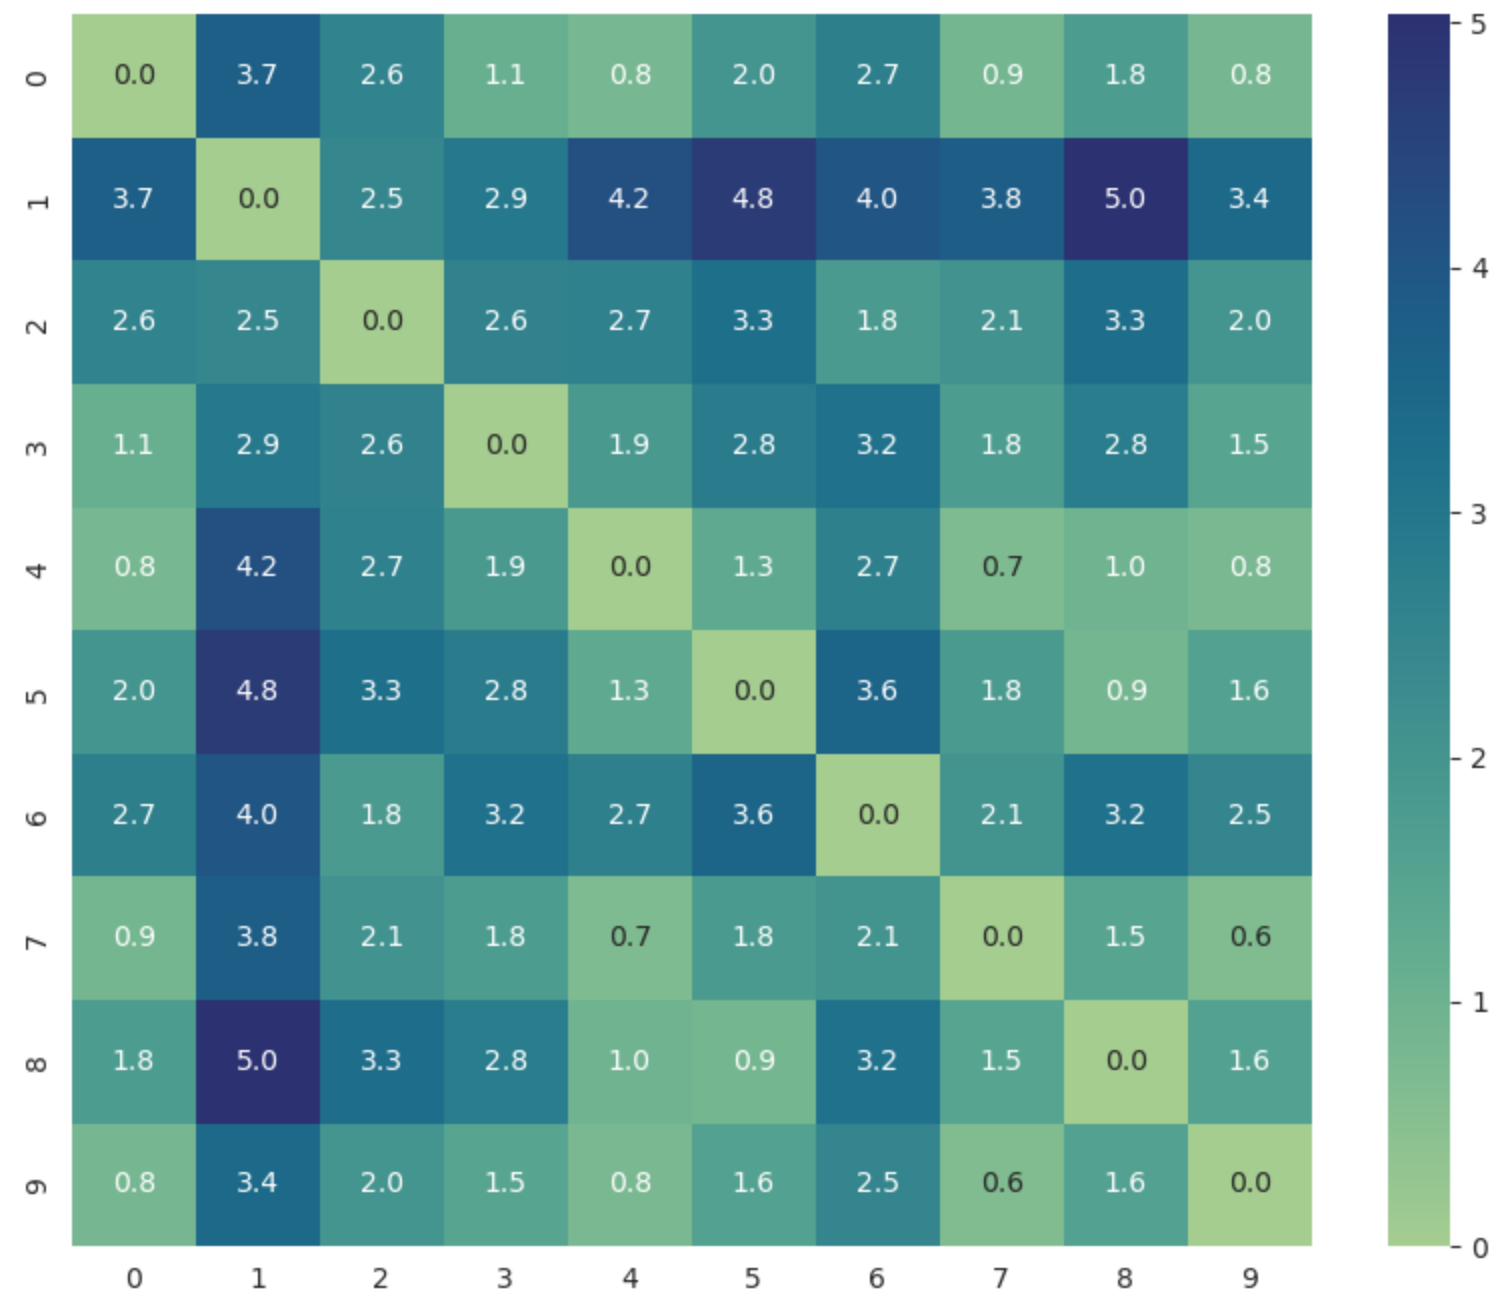
np.round(dists,2)
```

```
Out[26]: array([[0.   , 3.74, 2.59, 1.12, 0.82, 2.01, 2.73, 0.87, 1.76, 0.84],
 [3.74, 0.   , 2.49, 2.94, 4.19, 4.78, 4.01, 3.79, 5.03, 3.44],
 [2.59, 2.49, 0.   , 2.61, 2.65, 3.25, 1.83, 2.07, 3.31, 2.01],
 [1.12, 2.94, 2.61, 0.   , 1.87, 2.83, 3.19, 1.76, 2.8 , 1.47],
 [0.82, 4.19, 2.65, 1.87, 0.   , 1.32, 2.69, 0.68, 0.97, 0.78],
 [2.01, 4.78, 3.25, 2.83, 1.32, 0.   , 3.59, 1.81, 0.87, 1.61],
 [2.73, 4.01, 1.83, 3.19, 2.69, 3.59, 0.   , 2.05, 3.2 , 2.51],
 [0.87, 3.79, 2.07, 1.76, 0.68, 1.81, 2.05, 0.   , 1.5 , 0.64],
 [1.76, 5.03, 3.31, 2.8 , 0.97, 0.87, 3.2 , 1.5 , 0.   , 1.61],
 [0.84, 3.44, 2.01, 1.47, 0.78, 1.61, 2.51, 0.64, 1.61, 0.   ]])
```


Visualizing Distances With a Heatmap

Visualizing Distances With a Heatmap

```
In [27]: fig, ax = plt.subplots(1, 1, figsize=(10, 8))  
sns.heatmap(dists, annot=True, fmt=".1f", cmap='crest');
```



Query For Similarity

- Imagine I like house 5
- What houses are similar to house 5?

Query For Similarity

- Imagine I like house 5
- What houses are similar to house 5?

```
In [28]: query_idx = 5  
df_house.iloc[query_idx]
```

```
Out[28]: SqFtTotLiving    930.0  
SqFtLot                1012.0  
AdjSalePrice           411781.0  
Name: 5, dtype: float64
```

Query For Similarity

- Imagine I like house 5
- What houses are similar to house 5?

```
In [28]: query_idx = 5  
df_house.iloc[query_idx]
```

```
Out[28]: SqFtTotLiving    930.0  
SqFtLot              1012.0  
AdjSalePrice         411781.0  
Name: 5, dtype: float64
```

```
In [29]: # Distances to house 5  
[f'{x:0.1f}' for x in dists[query_idx]]
```

```
Out[29]: ['2.0', '4.8', '3.3', '2.8', '1.3', '0.0', '3.6', '1.8', '0.9', '1.6']
```

Query For Similarity Cont.

Query For Similarity Cont.

```
In [30]: # find indexes of best scores (for distances, want ascending)
         best_idx asc = np.argsort(dists[query_idx])
         best_idx asc
```

```
Out[30]: array([5, 8, 4, 9, 7, 0, 3, 2, 6, 1])
```

Query For Similarity Cont.

```
In [30]: # find indexes of best scores (for distances, want ascending)
best_idx asc = np.argsort(dists[query_idx])
best_idx asc
```


```
Out[30]: array([5, 8, 4, 9, 7, 0, 3, 2, 6, 1])
```

```
In [31]: # the top 10 recommendations with their distances
list(zip(['house ' + str(x) for x in best_idx asc],
        np.round(dists[query_idx][best_idx asc], 2)
        )
    )
```

```
Out[31]: [('house 5', 0.0),
          ('house 8', 0.87),
          ('house 4', 1.32),
          ('house 9', 1.61),
          ('house 7', 1.81),
          ('house 0', 2.01),
          ('house 3', 2.83),
          ('house 2', 3.25),
          ('house 6', 3.59),
          ('house 1', 4.78)]
```


(User Based) Collaborative Filtering

- Recommend things that **people with similar tastes have liked**
- Assume: If both you and I like Movie A, and you like Movie B, I'll like movie B
- Use similarity of user preferences
- Matrix: Users x Items
- Values: Rankings

| | |  |  |  |  |
|---|---|---|---|---|---|
| A |  |  |  |  |  |
| B |  | |  |  |  |
| C |  |  |  |  | |
| D |  |  | |  | |
| E |  |  |  |  |  |

Example: User Interests

Can we recommend topics based on a users existing interests?

Example: User Interests

Can we recommend topics based on a users existing interests?

```
In [32]: # from Data Science from Scratch by Joel Grus
#https://github.com/joelgrus/data-science-from-scratch.git

users_interests = [
    ["Hadoop", "Big Data", "HBase", "Java", "Spark", "Storm", "Cassandra"],
    ["NoSQL", "MongoDB", "Cassandra", "HBase", "Postgres"],
    ["Python", "scikit-learn", "scipy", "numpy", "statsmodels", "pandas"],
    ["R", "Python", "statistics", "regression", "probability"],
    ["machine learning", "regression", "decision trees", "libsvm"],
    ["Python", "R", "Java", "C++", "Haskell", "programming languages"],
    ["statistics", "probability", "mathematics", "theory"],
    ["machine learning", "scikit-learn", "Mahout", "neural networks"],
    ["neural networks", "deep learning", "Big Data", "artificial intelligence"],
    ["Hadoop", "Java", "MapReduce", "Big Data"],
    ["statistics", "R", "statsmodels"],
    ["C++", "deep learning", "artificial intelligence", "probability"],
    ["pandas", "R", "Python"],
    ["databases", "HBase", "Postgres", "MySQL", "MongoDB"],
    ["libsvm", "regression", "support vector machines"]
]
```

Example: User Interests

Can we recommend topics based on a users existing interests?

```
In [32]: # from Data Science from Scratch by Joel Grus
#https://github.com/joelgrus/data-science-from-scratch.git

users_interests = [
    ["Hadoop", "Big Data", "HBase", "Java", "Spark", "Storm", "Cassandra"],
    ["NoSQL", "MongoDB", "Cassandra", "HBase", "Postgres"],
    ["Python", "scikit-learn", "scipy", "numpy", "statsmodels", "pandas"],
    ["R", "Python", "statistics", "regression", "probability"],
    ["machine learning", "regression", "decision trees", "libsvm"],
    ["Python", "R", "Java", "C++", "Haskell", "programming languages"],
    ["statistics", "probability", "mathematics", "theory"],
    ["machine learning", "scikit-learn", "Mahout", "neural networks"],
    ["neural networks", "deep learning", "Big Data", "artificial intelligence"],
    ["Hadoop", "Java", "MapReduce", "Big Data"],
    ["statistics", "R", "statsmodels"],
    ["C++", "deep learning", "artificial intelligence", "probability"],
    ["pandas", "R", "Python"],
    ["databases", "HBase", "Postgres", "MySQL", "MongoDB"],
    ["libsvm", "regression", "support vector machines"]
]
```

```
In [33]: # interests of user0
sorted(users_interests[0])
```

```
Out[33]: ['Big Data', 'Cassandra', 'HBase', 'Hadoop', 'Java', 'Spark', 'Storm']
```

All Unique Interests

All Unique Interests

```
In [34]: # get a sorted list of unique interests (here using set)
unique_interests = sorted({interest
                           for user_interests in users_interests
                           for interest in user_interests})

# the first 20 unique interests
unique_interests[:20]
```

```
Out[34]: ['Big Data',
          'C++',
          'Cassandra',
          'HBase',
          'Hadoop',
          'Haskell',
          'Java',
          'Mahout',
          'MapReduce',
          'MongoDB',
          'MySQL',
          'NoSQL',
          'Postgres',
          'Python',
          'R',
          'Spark',
          'Storm',
          'artificial intelligence',
          'databases',
          'decision trees']
```

Transform User Interest Matrix

Transform User Interest Matrix

```
In [35]: # Transform between lists of strings and fixed length lists of ints
from sklearn.preprocessing import MultiLabelBinarizer

mlb = MultiLabelBinarizer(classes=unique_interests)

# a matrix of "user" rows and "interest" columns
user_interest_matrix = mlb.fit_transform(users_interests)

# The interests for user0
user_interest_matrix[0]
```

```
Out[35]: array([1, 0, 1, 1, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0,
                0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0])
```


Transform User Interest Matrix

```
In [35]: # Transform between lists of strings and fixed length lists of ints
from sklearn.preprocessing import MultiLabelBinarizer

mlb = MultiLabelBinarizer(classes=unique_interests)

# a matrix of "user" rows and "interest" columns
user_interest_matrix = mlb.fit_transform(users_interests)

# The interests for user0
user_interest_matrix[0]
```

```
Out[35]: array([1, 0, 1, 1, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0,
               0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0])
```

```
In [36]: # transforming back from interest matrix to list of interests
mlb.inverse_transform(user_interest_matrix)[0]
```

```
Out[36]: ('Big Data', 'Cassandra', 'HBase', 'Hadoop', 'Java', 'Spark', 'Storm')
```

Calculate Similarity

Calculate Similarity

```
In [37]: from sklearn.metrics.pairwise import cosine_similarity

# using similarity, higher values are better
user_similarities = cosine_similarity(user_interest_matrix)

# what are the similarites for user0 to other users?
user_similarities[0].round(1)
```

```
Out[37]: array([1. , 0.3, 0. , 0. , 0. , 0.2, 0. , 0. , 0.2, 0.6, 0. , 0. , 0. ,
               0.2, 0. ])
```

Calculate Similarity

```
In [37]: from sklearn.metrics.pairwise import cosine_similarity

# using similarity, higher values are better
user_similarities = cosine_similarity(user_interest_matrix)

# what are the similarites for user0 to other users?
user_similarities[0].round(1)
```

```
Out[37]: array([1. , 0.3, 0. , 0. , 0. , 0.2, 0. , 0. , 0.2, 0.6, 0. , 0. , 0. ,
               0.2, 0. ])
```

```
In [38]: # what users does user0 share interests with?
np.where(user_similarities[0])[0]
```

```
Out[38]: array([ 0,  1,  5,  8,  9, 13])
```

Find Similar Users

Find Similar Users

```
In [39]: # return a sorted list of users based on similarity
# skip query user and similarity == 0
def most_similar_users_to(query_idx):
    users_scores = [(idx, np.round(sim, 2))
                     for idx, sim in enumerate(user_similarities[query_idx])
                     if idx != query_idx and sim > 0]
    return sorted(users_scores, key=lambda x: x[1])

pd.DataFrame(most_similar_users_to(0), columns=['user', 'similarity'])
```

Out[39]:

| | user | similarity |
|---|------|------------|
| 0 | 5 | 0.15 |
| 1 | 13 | 0.17 |
| 2 | 8 | 0.19 |
| 3 | 1 | 0.34 |
| 4 | 9 | 0.57 |

Recommend Based On User Similarity

- Want to return items liked by other users sorted by the similarity of those users

Recommend Based On User Similarity

- Want to return items liked by other users sorted by the similarity of those users

```
In [40]: from collections import defaultdict

def user_based_suggestions(user_idx):
    suggestions = defaultdict(float)

    # iterate over interests of similar users
    for other_idx, sim in most_similar_users_to(user_idx): # for each similar user
        for interest in users_interests[other_idx]:        # for each interest of that user
            suggestions[interest] += sim                    # add weight based on the similarity of that user

    # sort suggestions based on weight
    suggestions = sorted(suggestions.items(),
                        key=lambda x:x[1],
                        reverse=True)

    # return only new interests
    return [(suggestion, weight.round(2))
            for suggestion, weight in suggestions
            if suggestion not in users_interests[user_idx]] # weed out existing interests
```


Recommend Based On User Similarity

Recommend Based On User Similarity

```
In [41]: # reminder: original interests  
users_interests[0]
```

```
Out[41]: ['Hadoop', 'Big Data', 'HBase', 'Java', 'Spark', 'Storm', 'Cassandra']
```

Recommend Based On User Similarity

```
In [41]: # reminder: original interests
users_interests[0]
```

```
Out[41]: ['Hadoop', 'Big Data', 'HBase', 'Java', 'Spark', 'Storm', 'Cassandra']
```

```
In [42]: # top 5 new recommended interests
pd.DataFrame(user_based_suggestions(0)[:5], columns=['interest', 'weight'])
```

```
Out[42]:
```

| | interest | weight |
|---|-----------------|--------|
| 0 | MapReduce | 0.57 |
| 1 | Postgres | 0.51 |
| 2 | MongoDB | 0.51 |
| 3 | NoSQL | 0.34 |
| 4 | neural networks | 0.19 |

Issues with Collab. Filtering

- **the cold start problem** : What if it's your first time?
- **sparsity** : How to recommend movies no one's seen?

Recommendation as Classification

- set1 features + set2 features -> label
- generate label based on history
- Examples
 - user + item -> purchased or not
 - candidate + job -> hired or not
 - ...
- Feature Engineering!

Recommendation as Classification: Example

Recommendation as Classification: Example

```
In [43]: person_features = ['Age', 'Country', 'Interest']
book_features = ['Price', 'Language', 'Topic']

features = ['Person_Age',
            'Person_Country',
            'Person_Hobby',
            'Book_Price',
            'Book_Language',
            'Book_Topic',
            'Interest_Topic_Match',
            'Country_Language_Match',
            ]

# target: "Did person purchase book? 1 == yes, 0 == no"

# dataset: Generate all person x book pairs and calculate target
```

Recommendation as Classification: Prediction

Recommendation as Classification: Prediction

```
train classifier on dataset using one of our Classification Models
```

```
then, for a query_person:
```

1. generate all query_person x book pairs
2. calculate $P(y=1|X)$ for all pairs using `.predict_proba()`
3. rank by $P(y=1|X)$
4. return the top N books

Issues with Recommendation as Classification

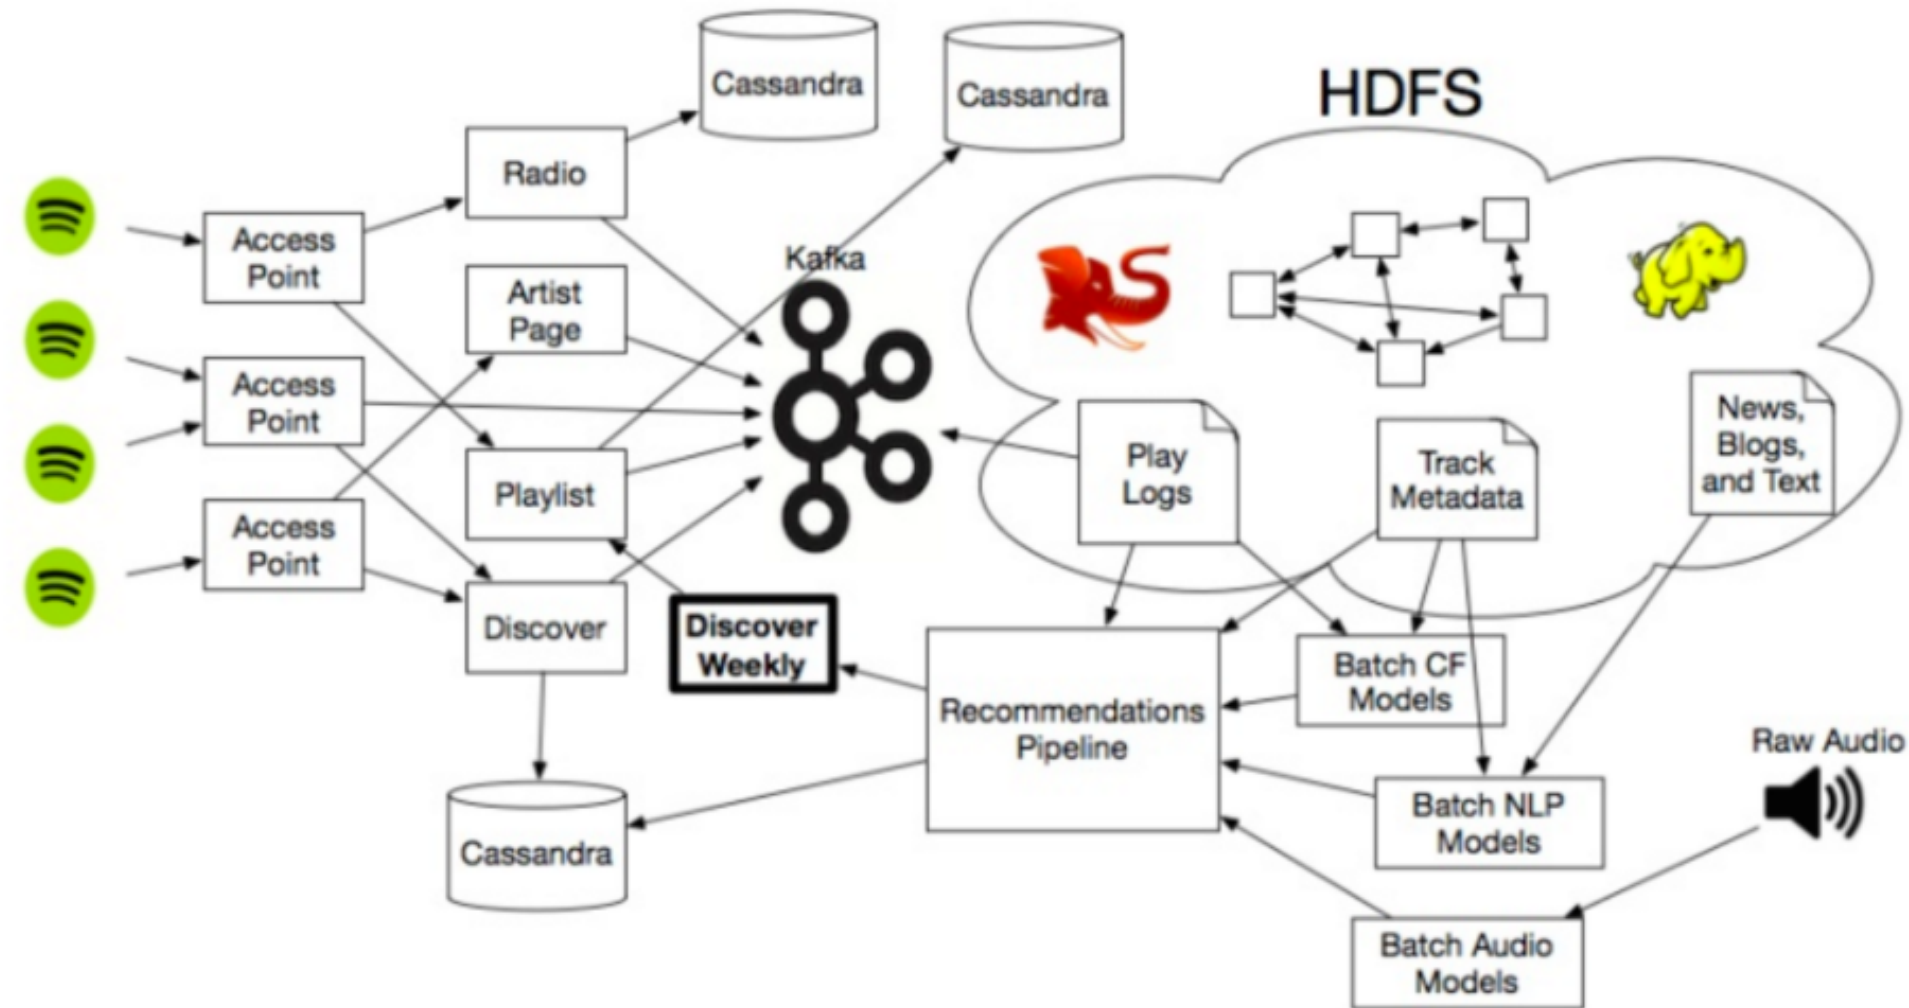
- Imbalanced classes
 - Example: each person bought different 1 of 100 books -> 1 pos to 99 neg
- False Negatives
 - Example: a person+book pair may be a good match even though it wasn't purchased

Evaluating Recommendation Systems

- **Precision At K:** Out of top K, how many were true/good? TP / K
- **Recall At K:** Out of all true/good, how many were in top K? $TP / (TP+FN)$
- Surprise/Novelty?
- Diversity?

Spotify's Recommendation Engine

How Does Spotify Know You So Well?



Recommendation Engines Review

- Content-Based
- User-Based Collaborative Filtering
- Recommendation as Classification
- Issues
- Evaluating

Questions re Recommendation Engines?

Imbalanced Classes

- Imbalanced classes:
 - when there is significantly more of one class than another in a classification task
- common in real world datasets
- Ex: credit card fraud
 - very small number of fraud transactions relative to total transactions

Dealing With Imbalanced Classes

- Stratified Sampling
- Random Undersampling
- Random Oversampling
- Oversample Synthetic Minority Items
 - SMOTE
 - ADASYN
- Other methods

Stratified Sampling

Stratified Sampling

```
In [44]: from sklearn.model_selection import StratifiedKFold

X = np.ones(9)
y = np.array([0, 0, 0, 0, 0, 0, 1, 1, 1])

skf = StratifiedKFold(n_splits=3)
for train_idx, test_idx in skf.split(X, y):
    print(f"indices : {train_idx} {test_idx}")
    print(f"values  : {y[train_idx]} {y[test_idx]}")
    print()
```

```
indices : [2 3 4 5 7 8] [0 1 6]
values  : [0 0 0 0 1 1] [0 0 1]
```

```
indices : [0 1 4 5 6 8] [2 3 7]
values  : [0 0 0 0 1 1] [0 0 1]
```

```
indices : [0 1 2 3 6 7] [4 5 8]
values  : [0 0 0 0 1 1] [0 0 1]
```

Random Sampling

- Randomly Oversample minority class
- Randomly Undersample majority class

Example Dataset

Example Dataset

```
In [45]: from sklearn.datasets import make_classification
from collections import Counter
X_imb, y_imb = make_classification(n_samples=5000, n_features=2, n_informative=2, n_redundant=0, n_repeated=0, n_classes=3,
                                n_clusters_per_class=1, weights=[0.01, 0.05, 0.94], class_sep=0.8, random_state=0)
df_imb = pd.DataFrame(X_imb); df_imb['y'] = y_imb; df_imb.y.value_counts()
```

```
Out[45]: 2    4674
         1     262
         0      64
         Name: y, dtype: int64
```

Example Dataset

```
In [45]: from sklearn.datasets import make_classification
from collections import Counter
X_imb, y_imb = make_classification(n_samples=5000, n_features=2, n_informative=2, n_redundant=0, n_repeated=0, n_classes=3,
                                n_clusters_per_class=1, weights=[0.01, 0.05, 0.94], class_sep=0.8, random_state=0)
df_imb = pd.DataFrame(X_imb); df_imb['y'] = y_imb; df_imb.y.value_counts()
```

```
Out[45]: 2    4674
         1     262
         0      64
         Name: y, dtype: int64
```

```
In [46]: fig,ax=plt.subplots(1,1,figsize=(8,6)); sns.scatterplot(x=0,y=1,hue='y',data=df_imb,palette="colorblind",alpha=.3,s=50);
```



Using imblearn

- `imblearn` is library to created to deal with imbalanced classes
- need to install from `conda-forge` as `imbalanced-learn`
- import from `imblearn`

Random Oversampling of minority class

Random Oversampling of minority class

```
In [48]: from imblearn.over_sampling import RandomOverSampler

ros = RandomOverSampler(random_state=0)
X_ros, y_ros = ros.fit_resample(X_imb, y_imb)
df_ros = pd.DataFrame(X_ros); df_ros['y'] = y_ros; df_ros.y.value_counts()
```

```
Out[48]: 2    4674
         1    4674
         0    4674
         Name: y, dtype: int64
```

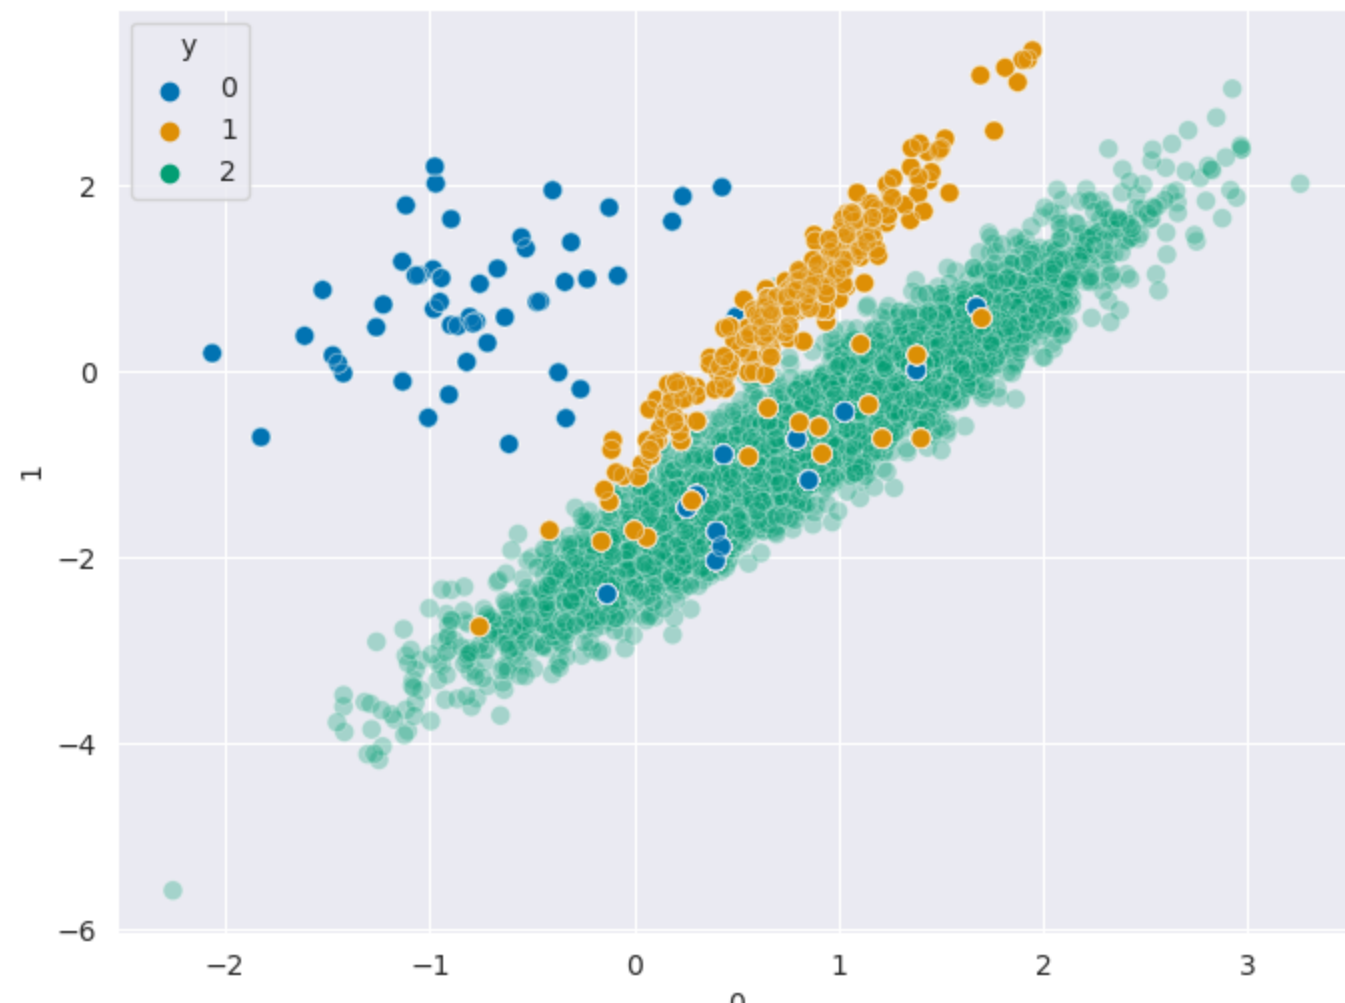
Random Oversampling of minority class

```
In [48]: from imblearn.over_sampling import RandomOverSampler

ros = RandomOverSampler(random_state=0)
X_ros, y_ros = ros.fit_resample(X_imb, y_imb)
df_ros = pd.DataFrame(X_ros); df_ros['y'] = y_ros; df_ros.y.value_counts()
```

```
Out[48]: 2    4674
         1    4674
         0    4674
         Name: y, dtype: int64
```

```
In [49]: fig,ax=plt.subplots(1,1,figsize=(8,6)); sns.scatterplot(x=0,y=1,hue='y',data=df_ros,palette="colorblind",alpha=.3,s=50);
```



Random Undersampling of majority class

Random Undersampling of majority class

```
In [50]: from imblearn.under_sampling import RandomUnderSampler

rus = RandomUnderSampler(random_state=0)
X_rus, y_rus, = rus.fit_resample(X_imb, y_imb)
df_rus = pd.DataFrame(X_rus); df_rus['y'] = y_rus; df_rus.y.value_counts()
```

```
Out[50]: 0    64
         1    64
         2    64
         Name: y, dtype: int64
```

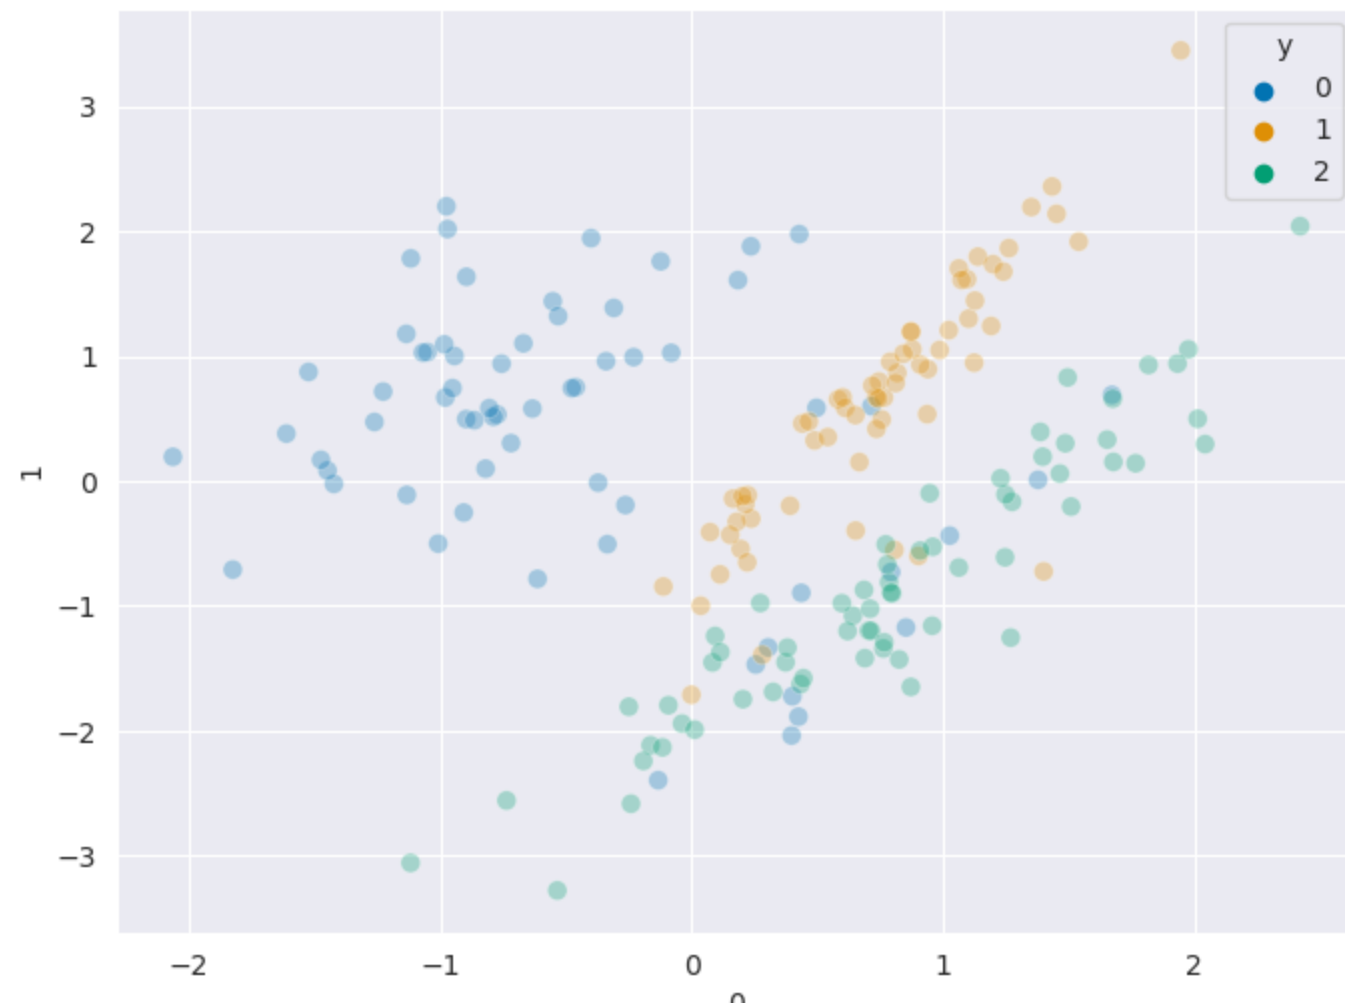
Random Undersampling of majority class

```
In [50]: from imblearn.under_sampling import RandomUnderSampler

rus = RandomUnderSampler(random_state=0)
X_rus, y_rus, = rus.fit_resample(X_imb, y_imb)
df_rus = pd.DataFrame(X_rus); df_rus['y'] = y_rus; df_rus.y.value_counts()
```

```
Out[50]: 0    64
         1    64
         2    64
         Name: y, dtype: int64
```

```
In [51]: fig,ax=plt.subplots(1,1,figsize=(8,6)); sns.scatterplot(x=0,y=1,hue='y',data=df_rus,palette="colorblind",alpha=.3,s=50);
```



Oversample Sythetic Minority Items

- SMOTE: Synthetic Minority Oversampling
- ADASYN: Adaptive Synthetic Minority Oversampling

SMOTE: Synthetic Minority Oversampling

- Create new synthetic points between existing points

SMOTE: Synthetic Minority Oversampling

- Create new synthetic points between existing points

```
In [52]: from imblearn.over_sampling import SMOTE
X_smote, y_smote = SMOTE().fit_resample(X_imb, y_imb)
df_smote = pd.DataFrame(X_smote); df_smote['y'] = y_smote; df_smote.y.value_counts()
```

```
Out[52]: 2    4674
         1    4674
         0    4674
         Name: y, dtype: int64
```


SMOTE: Synthetic Minority Oversampling

- Create new synthetic points between existing points

```
In [52]: from imblearn.over_sampling import SMOTE
X_smote, y_smote = SMOTE().fit_resample(X_imb, y_imb)
df_smote = pd.DataFrame(X_smote); df_smote['y'] = y_smote; df_smote.y.value_counts()
```

```
Out[52]: 2    4674
         1    4674
         0    4674
         Name: y, dtype: int64
```

```
In [53]: fig,ax=plt.subplots(1,1,figsize=(8,6)); sns.scatterplot(x=0,y=1,hue='y',data=df_smote,palette="colorblind",alpha=.3,s=50);
```

ADASYN: Adaptive Synthetic Minority Oversampling

- Create new synthetic points between existing points *where classes overlap*

ADASYN: Adaptive Synthetic Minority Oversampling

- Create new synthetic points between existing points *where classes overlap*

```
In [54]: from imblearn.over_sampling import ADASYN
X_adasyn, y_adasyn = ADASYN().fit_resample(X_imb, y_imb)
df_adasyn = pd.DataFrame(X_adasyn); df_adasyn['y'] = y_adasyn; df_adasyn.y.value_counts()
```

```
Out[54]: 2    4674
         0    4673
         1    4662
         Name: y, dtype: int64
```

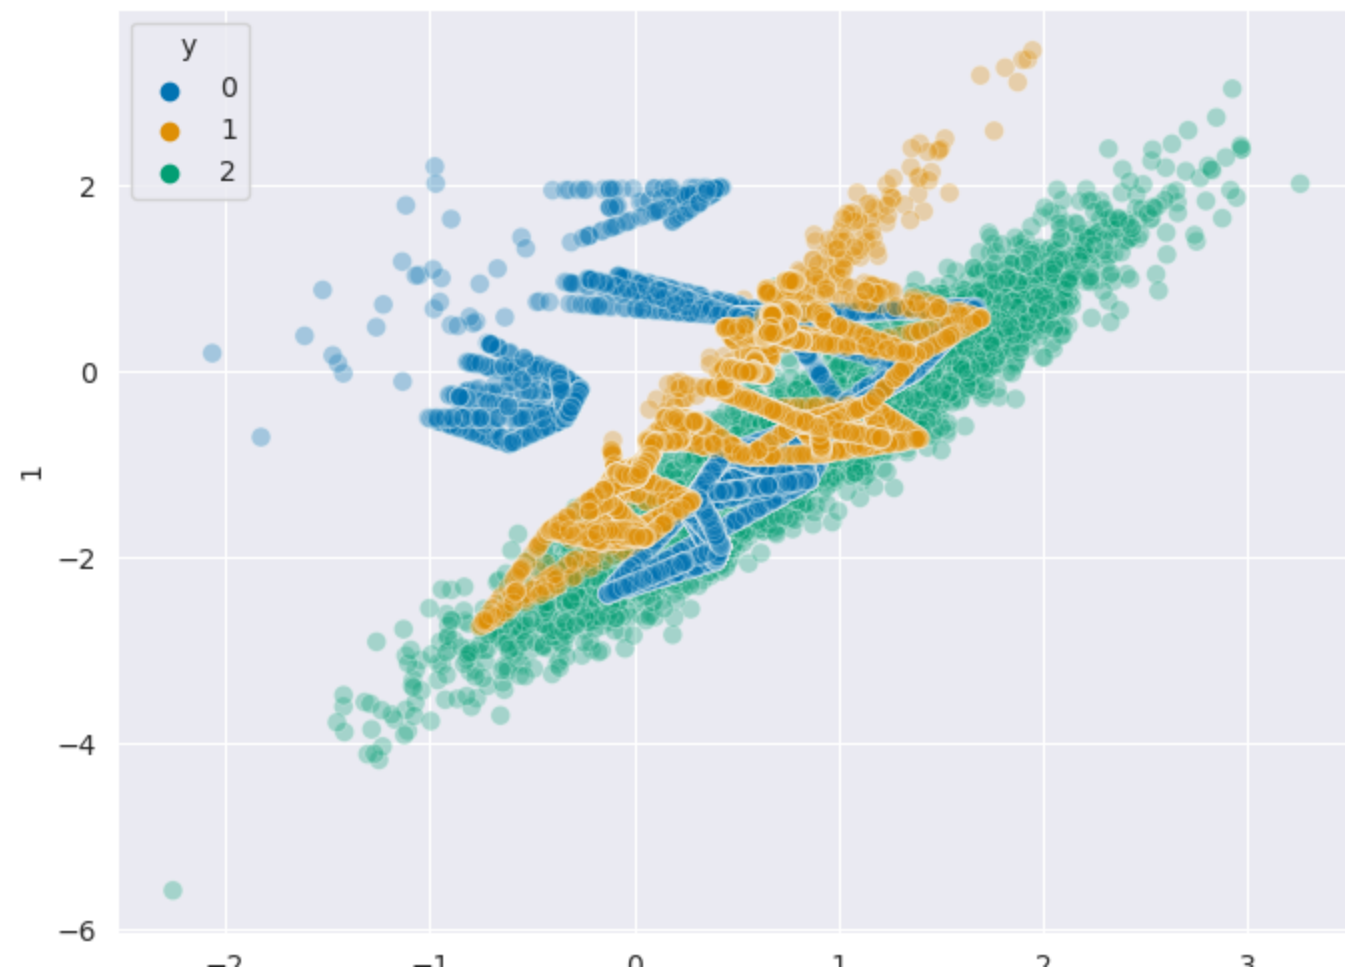
ADASYN: Adaptive Synthetic Minority Oversampling

- Create new synthetic points between existing points *where classes overlap*

```
In [54]: from imblearn.over_sampling import ADASYN
X_adasyn, y_adasyn = ADASYN().fit_resample(X_imb, y_imb)
df_adasyn = pd.DataFrame(X_adasyn); df_adasyn['y'] = y_adasyn; df_adasyn.y.value_counts()
```

```
Out[54]: 2    4674
         0    4673
         1    4662
         Name: y, dtype: int64
```

```
In [55]: fig,ax=plt.subplots(1,1,figsize=(8,6)); sns.scatterplot(x=0,y=1,hue='y',data=df_adasyn,palette="colorblind",alpha=.3,s=50);
```



Other methods for dealing with imbalanced classes

- Adjust class weight (sklearn)
 - Adjust decision threshold (sklearn)
 - Treat as anomaly detection
 - Generate/buy more labels
-
- See https://imbalanced-learn.readthedocs.io/en/stable/auto_examples/over-sampling/plot_comparison_over_sampling.html

Questions re Imbalanced Classes?