

# **Elements Of Data Science - F2022**

## **Week 12: Time Series, Model API with Flask**

**11/30/2020**

# TODOs

- Readings:
  - Recommended: DSFS: Chap 9: Getting Data
  - Recommended: DSFS: Chap 23: Databases and SQL
- HW4, **Due Friday December 2nd 11:59pm ET**
- Quiz 12, **Due Tuesday December 6th 11:59pm ET**
- Final
  - Review sheet in github repo (soon!)
  - Online via Gradescope, open-book, open-note, open-python
  - Released *Wednesday December 7th 11:59pm ET*
  - **Due Friday December 9th 11:59pm ET**
  - Have maximum of 24hrs after starting to finish
  - 30-40 questions (fill in the blank/multiple choice/short answer)
  - Questions asked/answered **privately** via Ed

# Quiz Common Mistakes (points off)

- don't remove instructions from quiz/homework
- `.info()` not `.info`: make sure function/method calls are made with `()`
- Pandas `.sample()` default `n=1`: need to set `n=` or `frac=`
- LinearRegression (regression) vs LogisticRegression (classification)
- using a model "with default settings" means `Model()` or just a subset of parameters set
- Be careful which dataset you're training/evaluating on: `X_train` vs `X_test`
- Make sure all plotting settings get used (eg `hue=`)
- `y_digits[cluster_assignments_km == 9]`
- `sns.scatterplot(x=X_2D[:,0], y=X_2D[:,1], ...)`

# Today

- Time Series Transformations
- Model Prediction API with Flask

## Questions?

# Environment Setup

# Environment Setup

```
In [1]: import numpy
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

sns.set_style('darkgrid')
%matplotlib inline
```

# Time Series

- Data ordered in time
- Applications
  - Financial
  - Economic
  - Scientific
  - etc.

# Time Series Differences

- **Non-i.i.d.** : not independent and identically distributed
- not independent
  - Ex: Stock price
- not-identically distributed
  - Ex: Seasonality
- In other words: Order matters!



# Representing Time in Python

- `datetime` library
- Pandas `Timestamp`

**datetime.date**

# datetime.date

```
In [2]: from datetime import date

        friday = date(2022,11,1) # year,month,day
        friday
```

```
Out[2]: datetime.date(2022, 11, 1)
```

# datetime.date

```
In [2]: from datetime import date

        friday = date(2022,11,1) # year,month,day
        friday
```

Out[2]: datetime.date(2022, 11, 1)

```
In [3]: today = date.today()
        today
```

Out[3]: datetime.date(2022, 11, 30)

# datetime.date

```
In [2]: from datetime import date

        friday = date(2022,11,1) # year,month,day
        friday
```

Out[2]: datetime.date(2022, 11, 1)

```
In [3]: today = date.today()
        today
```

Out[3]: datetime.date(2022, 11, 30)

```
In [4]: today.year
```

Out[4]: 2022

**datetime.time**

# datetime.time

```
In [5]: from datetime import time

        class_start = time(19,10,0) # hour,minute,second,microsecond
        class_start
```

```
Out[5]: datetime.time(19, 10)
```

# datetime.time

```
In [5]: from datetime import time

class_start = time(19,10,0) # hour,minute,second,microsecond
class_start
```

```
Out[5]: datetime.time(19, 10)
```

```
In [6]: class_start.hour
```

```
Out[6]: 19
```



**datetime.datetime**

# datetime.datetime

```
In [7]: from datetime import datetime

        # year, month, day, hour, minute, second, microsecond
        wednesday_afternoon = datetime(2022, 11, 30, 15)
        wednesday_afternoon
```

```
Out[7]: datetime.datetime(2022, 11, 30, 15, 0)
```

# datetime.datetime

```
In [7]: from datetime import datetime

        # year, month, day, hour, minute, second, microsecond
        wednesday_afternoon = datetime(2022, 11, 30, 15)
        wednesday_afternoon
```

```
Out[7]: datetime.datetime(2022, 11, 30, 15, 0)
```

```
In [8]: now = datetime.now()
        now
```

```
Out[8]: datetime.datetime(2022, 11, 30, 18, 4, 6, 48207)
```

**`datetime.timedelta`**

# datetime.timedelta

```
In [9]: diff = datetime(2022, 11, 30, 1) - datetime(2022, 11, 29, 0)
diff
```

```
Out[9]: datetime.timedelta(days=1, seconds=3600)
```

# datetime.timedelta

```
In [9]: diff = datetime(2022,11,30,1) - datetime(2022,11,29,0)
diff
```

```
Out[9]: datetime.timedelta(days=1, seconds=3600)
```

```
In [10]: diff.total_seconds()
```

```
Out[10]: 90000.0
```

# datetime.timedelta

```
In [9]: diff = datetime(2022, 11, 30, 1) - datetime(2022, 11, 29, 0)
diff
```

```
Out[9]: datetime.timedelta(days=1, seconds=3600)
```

```
In [10]: diff.total_seconds()
```

```
Out[10]: 90000.0
```

```
In [11]: from datetime import timedelta

#days, seconds, microseconds, milliseconds, minutes, hours, weeks
one_day = timedelta(1)

date(2022, 11, 30) + 2*one_day
```

```
Out[11]: datetime.date(2022, 12, 2)
```

Printing Datetimes: `strftime()`



# Printing Datetimes: `strftime()`

```
In [12]: now = datetime.now()  
         print(now)
```

```
2022-11-30 18:04:06.077953
```

# Printing Datetimes: `strftime()`

```
In [12]: now = datetime.now()  
         print(now)
```

```
2022-11-30 18:04:06.077953
```

```
In [13]: now.strftime('%a %h %d, %Y %I:%M %p')
```

```
Out[13]: 'Wed Nov 30, 2022 06:04 PM'
```

# Printing Datetimes: `strftime()`

```
In [12]: now = datetime.now()  
         print(now)
```

```
2022-11-30 18:04:06.077953
```

```
In [13]: now.strftime('%a %h %d, %Y %I:%M %p')
```

```
Out[13]: 'Wed Nov 30, 2022 06:04 PM'
```

```
%Y 4-digit year  
%y 2-digit year  
%m 2-digit month  
%d 2-digit day  
%H Hour (24-hour)  
%M 2-digit minute  
%S 2-digit second  
...
```

# Printing Datetimes: `strftime()`

```
In [12]: now = datetime.now()  
         print(now)
```

```
2022-11-30 18:04:06.077953
```

```
In [13]: now.strftime('%a %h %d, %Y %I:%M %p')
```

```
Out[13]: 'Wed Nov 30, 2022 06:04 PM'
```

```
%Y 4-digit year  
%y 2-digit year  
%m 2-digit month  
%d 2-digit day  
%H Hour (24-hour)  
%M 2-digit minute  
%S 2-digit second  
...
```

See [strftime.org](http://strftime.org) and [strfti.me](http://strfti.me)

# Parsing Datetimes: `pandas.to_datetime()`

- `dateutil.parser` available
- pandas has parser built in: `pd.to_datetime()`

# Parsing Datetimes: `pandas.to_datetime()`

- `dateutil.parser` available
- pandas has parser built in: `pd.to_datetime()`

```
In [14]: pd.to_datetime('11/30/2022 7:36pm')
```

```
Out[14]: Timestamp('2022-11-30 19:36:00')
```



# pandas.Timestamp

- like `datetime.datetime`
- can include **timezone** and **frequency** info
- can handle a missing time: `NaT`
- can be used anywhere `datetime` can be used
- an array of Timestamps can be used as an index



# pandas.Timestamp

- like `datetime.datetime`
- can include **timezone** and **frequency** info
- can handle a missing time: `NaT`
- can be used anywhere `datetime` can be used
- an array of Timestamps can be used as an index

```
In [16]: pd.Timestamp(2022, 11, 30, 19)
```

```
Out[16]: Timestamp('2022-11-30 19:00:00')
```

# pandas.Timestamp

- like `datetime.datetime`
- can include **timezone** and **frequency** info
- can handle a missing time: `NaT`
- can be used anywhere `datetime` can be used
- an array of Timestamps can be used as an index

```
In [16]: pd.Timestamp(2022, 11, 30, 19)
```

```
Out[16]: Timestamp('2022-11-30 19:00:00')
```

```
In [17]: pd.Timestamp('20221130 7:00pm EST')
```

```
Out[17]: Timestamp('2022-11-30 19:00:00-0500', tz='tzlocal()')
```

```
In [18]: pd.Timestamp('20221130 7:00pm', tz='US/Pacific')
```

```
Out[18]: Timestamp('2022-11-30 19:00:00-0800', tz='US/Pacific')
```

# pandas.Timestamp

- like `datetime.datetime`
- can include **timezone** and **frequency** info
- can handle a missing time: `NaT`
- can be used anywhere `datetime` can be used
- an array of Timestamps can be used as an index

```
In [16]: pd.Timestamp(2022, 11, 30, 19)
```

```
Out[16]: Timestamp('2022-11-30 19:00:00')
```

```
In [17]: pd.Timestamp('20221130 7:00pm EST')
```

```
Out[17]: Timestamp('2022-11-30 19:00:00-0500', tz='tzlocal()')
```

```
In [18]: pd.Timestamp('20221130 7:00pm', tz='US/Pacific')
```

```
Out[18]: Timestamp('2022-11-30 19:00:00-0800', tz='US/Pacific')
```

```
In [19]: dt_index[0]
```

```
Out[19]: Timestamp('2020-11-26 00:00:00')
```

# Accessing Datetime Components with `.dt`

# Accessing Datetime Components with `.dt`

```
In [20]: df_taxi = pd.read_csv('../data/yellowcab_tripdata_2017-01_subset10000rows.csv',  
                                parse_dates=['tpep_pickup_datetime']).head(3)  
df_taxi.tpep_pickup_datetime
```

```
Out[20]: 0    2017-01-10 18:37:59  
        1    2017-01-05 15:14:52  
        2    2017-01-11 14:47:52  
        Name: tpep_pickup_datetime, dtype: datetime64[ns]
```

# Accessing Datetime Components with `.dt`

```
In [20]: df_taxi = pd.read_csv('../data/yellowcab_tripdata_2017-01_subset10000rows.csv',  
                                parse_dates=['tpep_pickup_datetime']).head(3)  
df_taxi.tpep_pickup_datetime
```

```
Out[20]: 0    2017-01-10 18:37:59  
        1    2017-01-05 15:14:52  
        2    2017-01-11 14:47:52  
        Name: tpep_pickup_datetime, dtype: datetime64[ns]
```

```
In [21]: df_taxi.tpep_pickup_datetime.dt.day
```

```
Out[21]: 0     10  
        1      5  
        2     11  
        Name: tpep_pickup_datetime, dtype: int64
```

# Accessing Datetime Components with `.dt`

```
In [20]: df_taxi = pd.read_csv('../data/yellowcab_tripdata_2017-01_subset10000rows.csv',  
                                parse_dates=['tpep_pickup_datetime']).head(3)  
df_taxi.tpep_pickup_datetime
```

```
Out[20]: 0    2017-01-10 18:37:59  
        1    2017-01-05 15:14:52  
        2    2017-01-11 14:47:52  
        Name: tpep_pickup_datetime, dtype: datetime64[ns]
```

```
In [21]: df_taxi.tpep_pickup_datetime.dt.day
```

```
Out[21]: 0     10  
        1      5  
        2     11  
        Name: tpep_pickup_datetime, dtype: int64
```

```
In [22]: df_taxi.tpep_pickup_datetime.dt.day_of_week # Monday=0 ... Sunday=6
```

```
Out[22]: 0      1  
        1      3  
        2      2  
        Name: tpep_pickup_datetime, dtype: int64
```

# Accessing Datetime Components with `.dt`

```
In [20]: df_taxi = pd.read_csv('../data/yellowcab_tripdata_2017-01_subset10000rows.csv',  
                                parse_dates=['tpep_pickup_datetime']).head(3)  
df_taxi.tpep_pickup_datetime
```

```
Out[20]: 0    2017-01-10 18:37:59  
         1    2017-01-05 15:14:52  
         2    2017-01-11 14:47:52  
         Name: tpep_pickup_datetime, dtype: datetime64[ns]
```

```
In [21]: df_taxi.tpep_pickup_datetime.dt.day
```

```
Out[21]: 0     10  
         1      5  
         2     11  
         Name: tpep_pickup_datetime, dtype: int64
```

```
In [22]: df_taxi.tpep_pickup_datetime.dt.day_of_week # Monday=0 ... Sunday=6
```

```
Out[22]: 0      1  
         1      3  
         2      2  
         Name: tpep_pickup_datetime, dtype: int64
```

```
In [23]: df_taxi.tpep_pickup_datetime.dt.hour
```

```
Out[23]: 0     18  
         1     15  
         2     14  
         Name: tpep_pickup_datetime, dtype: int64
```



# DateIndex Indexing/Selecting/Slicing

# DateIndex Indexing/Selecting/Slicing

```
In [24]: s = pd.Series(['Dec 1 2021', 'Jan 2 2022', 'Feb 3 2022'],  
                      index=pd.to_datetime(['Dec 1 2021', 'Jan 2 2022', 'Feb 3 2022']))  
s
```

```
Out[24]: 2021-12-01    Dec 1 2021  
         2022-01-02    Jan 2 2022  
         2022-02-03    Feb 3 2022  
         dtype: object
```

# DateIndex Indexing/Selecting/Slicing

```
In [24]: s = pd.Series(['Dec 1 2021', 'Jan 2 2022', 'Feb 3 2022'],  
                      index=pd.to_datetime(['Dec 1 2021', 'Jan 2 2022', 'Feb 3 2022']))  
s
```

```
Out[24]: 2021-12-01    Dec 1 2021  
         2022-01-02    Jan 2 2022  
         2022-02-03    Feb 3 2022  
         dtype: object
```

```
In [25]: # can index normally using iloc  
s.iloc[0:2]
```

```
Out[25]: 2021-12-01    Dec 1 2021  
         2022-01-02    Jan 2 2022  
         dtype: object
```

# DateIndex Indexing/Selecting/Slicing Cont.

# DateIndex Indexing/Selecting/Slicing Cont.

```
In [26]: # only rows from the year 2022  
s.loc['2022']
```

```
Out[26]: 2022-01-02    Jan 2 2022  
         2022-02-03    Feb 3 2022  
         dtype: object
```

# DateIndex Indexing/Selecting/Slicing Cont.

```
In [26]: # only rows from the year 2022  
s.loc['2022']
```

```
Out[26]: 2022-01-02    Jan 2 2022  
         2022-02-03    Feb 3 2022  
         dtype: object
```

```
In [27]: # only rows from January 2022  
s.loc['2022-01']
```

```
Out[27]: 2022-01-02    Jan 2 2022  
         dtype: object
```

# DateIndex Indexing/Selecting/Slicing Cont.

```
In [26]: # only rows from the year 2022  
s.loc['2022']
```

```
Out[26]: 2022-01-02    Jan 2 2022  
         2022-02-03    Feb 3 2022  
         dtype: object
```

```
In [27]: # only rows from January 2022  
s.loc['2022-01']
```

```
Out[27]: 2022-01-02    Jan 2 2022  
         dtype: object
```

```
In [28]: # only rows between Jan 1st 2021 and Jan 2nd 2022, inclusive  
s.loc['01/01/2021':'01/02/2022']
```

```
Out[28]: 2021-12-01    Dec 1 2021  
         2022-01-02    Jan 2 2022  
         dtype: object
```

# Datetimes in DataFrames



# Datetimes in DataFrames

```
In [29]: df = pd.DataFrame([['12/1/2021', 101, 'A'],  
                           ['1/1/2022', 102, 'B']], columns=['col1', 'col2', 'col3'])  
df.col1 = pd.to_datetime(df.col1)  
df.set_index('col1', drop=True, inplace=True)  
df
```

Out[29]:

	col2	col3
col1		
2021-12-01	101	A
2022-01-01	102	B

# Datetimes in DataFrames

```
In [29]: df = pd.DataFrame([[ '12/1/2021', 101, 'A'],  
                           [ '1/1/2022', 102, 'B']], columns=[ 'col1', 'col2', 'col3'])  
df.col1 = pd.to_datetime(df.col1)  
df.set_index('col1', drop=True, inplace=True)  
df
```

```
Out[29]:
```

	col2	col3
col1		
2021-12-01	101	A
2022-01-01	102	B

```
In [30]: # only return rows from 2022  
df.loc['2022']
```

```
Out[30]:
```

	col2	col3
col1		
2022-01-01	102	B

# Timestamp Index: Setting Frequency

# Timestamp Index: Setting Frequency

```
In [31]: s = pd.Series(['Nov 1 2022', 'Nov 3 2022'], index=pd.to_datetime(['Nov 1 2022', 'Nov 3 2022']))  
s
```

```
Out[31]: 2022-11-01    Nov 1 2022  
         2022-11-03    Nov 3 2022  
         dtype: object
```

# Timestamp Index: Setting Frequency

```
In [31]: s = pd.Series(['Nov 1 2022', 'Nov 3 2022'], index=pd.to_datetime(['Nov 1 2022', 'Nov 3 2022']))  
s
```

```
Out[31]: 2022-11-01    Nov 1 2022  
         2022-11-03    Nov 3 2022  
         dtype: object
```

```
In [32]: # Use resample() and asfreq() to set frequency  
s.resample('D').asfreq()
```

```
Out[32]: 2022-11-01    Nov 1 2022  
         2022-11-02         NaN  
         2022-11-03    Nov 3 2022  
         Freq: D, dtype: object
```

# Timestamp Index: Setting Frequency

```
In [31]: s = pd.Series(['Nov 1 2022', 'Nov 3 2022'], index=pd.to_datetime(['Nov 1 2022', 'Nov 3 2022']))  
s
```

```
Out[31]: 2022-11-01    Nov 1 2022  
         2022-11-03    Nov 3 2022  
         dtype: object
```

```
In [32]: # Use resample() and asfreq() to set frequency  
s.resample('D').asfreq()
```

```
Out[32]: 2022-11-01    Nov 1 2022  
         2022-11-02         NaN  
         2022-11-03    Nov 3 2022  
         Freq: D, dtype: object
```

```
In [33]: pd.to_datetime(['Nov 1 2022', 'Nov 3 2022'])
```

```
Out[33]: DatetimeIndex(['2022-11-01', '2022-11-03'], dtype='datetime64[ns]', freq=None)
```

# Timestamp Index: Setting Frequency

```
In [31]: s = pd.Series(['Nov 1 2022', 'Nov 3 2022'], index=pd.to_datetime(['Nov 1 2022', 'Nov 3 2022']))  
s
```

```
Out[31]: 2022-11-01    Nov 1 2022  
         2022-11-03    Nov 3 2022  
         dtype: object
```

```
In [32]: # Use resample() and asfreq() to set frequency  
s.resample('D').asfreq()
```

```
Out[32]: 2022-11-01    Nov 1 2022  
         2022-11-02         NaN  
         2022-11-03    Nov 3 2022  
         Freq: D, dtype: object
```

```
In [33]: pd.to_datetime(['Nov 1 2022', 'Nov 3 2022'])
```

```
Out[33]: DatetimeIndex(['2022-11-01', '2022-11-03'], dtype='datetime64[ns]', freq=None)
```

```
In [34]: # Use date_range with freq to get a range of dates of a certain frequency  
pd.date_range(start='Nov 1 2022', end='Nov 3 2022', freq='D')
```

```
Out[34]: DatetimeIndex(['2022-11-01', '2022-11-02', '2022-11-03'], dtype='datetime64[ns]', freq='D')
```

## Sample of Available Frequencies

B	business day frequency
D	calendar day frequency
W	weekly frequency
M	month end frequency
BM	business month end frequency
...	
Q	quarter end frequency
BQ	business quarter end frequency
...	
Y	year end frequency
BY	business year end frequency
...	
BH	business hour frequency
H	hourly frequency
T,min	minutely frequency
S	secondly frequency
L,ms	milliseconds
U,us	microseconds
N	nanoseconds



## Sample of Available Frequencies

B	business day frequency
D	calendar day frequency
W	weekly frequency
M	month end frequency
BM	business month end frequency
...	
Q	quarter end frequency
BQ	business quarter end frequency
...	
Y	year end frequency
BY	business year end frequency
...	
BH	business hour frequency
H	hourly frequency
T,min	minutely frequency
S	secondly frequency
L,ms	milliseconds
U,us	microseconds
N	nanoseconds

# Timezones

- Handled by `pytz` library

# Timezones

- Handled by `pytz` library

```
In [35]: import pytz

[x for x in pytz.common_timezones if x.startswith('U')]
```

```
Out[35]: ['US/Alaska',
          'US/Arizona',
          'US/Central',
          'US/Eastern',
          'US/Hawaii',
          'US/Mountain',
          'US/Pacific',
          'UTC']
```

# Timezones

- Handled by `pytz` library

```
In [35]: import pytz

[x for x in pytz.common_timezones if x.startswith('U')]
```

```
Out[35]: ['US/Alaska',
          'US/Arizona',
          'US/Central',
          'US/Eastern',
          'US/Hawaii',
          'US/Mountain',
          'US/Pacific',
          'UTC']
```

UTC: coordinated universal time (EST is 5 hours behind, -5:00)

# Timezones Cont.

# Timezones Cont.

```
In [36]: ts = pd.date_range('11/2/2022 9:30am', periods=2, freq='D')
         ts
```

```
Out[36]: DatetimeIndex(['2022-11-02 09:30:00', '2022-11-03 09:30:00'], dtype='datetime64[ns]', freq='D')
```

# Timezones Cont.

```
In [36]: ts = pd.date_range('11/2/2022 9:30am', periods=2, freq='D')
ts
```

```
Out[36]: DatetimeIndex(['2022-11-02 09:30:00', '2022-11-03 09:30:00'], dtype='datetime64[ns]', freq='D')
```

```
In [37]: # Set timezone using .localize()
ts_utc = ts.tz_localize('US/Eastern')
ts_utc
```

```
Out[37]: DatetimeIndex(['2022-11-02 09:30:00-04:00', '2022-11-03 09:30:00-04:00'], dtype='datetime64[ns, US/Eastern]', freq=None)
```

# Timezones Cont.

```
In [36]: ts = pd.date_range('11/2/2022 9:30am', periods=2, freq='D')
ts
```

```
Out[36]: DatetimeIndex(['2022-11-02 09:30:00', '2022-11-03 09:30:00'], dtype='datetime64[ns]', freq='D')
```

```
In [37]: # Set timezone using .localize()
ts_utc = ts.tz_localize('US/Eastern')
ts_utc
```

```
Out[37]: DatetimeIndex(['2022-11-02 09:30:00-04:00', '2022-11-03 09:30:00-04:00'], dtype='datetime64[ns, US/Eastern]', freq=None)
```

```
In [38]: # Change timezones using .tz_convert()
ts_utc.tz_convert('UTC')
```

```
Out[38]: DatetimeIndex(['2022-11-02 13:30:00+00:00', '2022-11-03 13:30:00+00:00'], dtype='datetime64[ns, UTC]', freq=None)
```



# Timezones Cont.

```
In [36]: ts = pd.date_range('11/2/2022 9:30am', periods=2, freq='D')
ts
```

```
Out[36]: DatetimeIndex(['2022-11-02 09:30:00', '2022-11-03 09:30:00'], dtype='datetime64[ns]', freq='D')
```

```
In [37]: # Set timezone using .localize()
ts_utc = ts.tz_localize('US/Eastern')
ts_utc
```

```
Out[37]: DatetimeIndex(['2022-11-02 09:30:00-04:00', '2022-11-03 09:30:00-04:00'], dtype='datetime64[ns, US/Eastern]', freq=None)
```

```
In [38]: # Change timezones using .tz_convert()
ts_utc.tz_convert('UTC')
```

```
Out[38]: DatetimeIndex(['2022-11-02 13:30:00+00:00', '2022-11-03 13:30:00+00:00'], dtype='datetime64[ns, UTC]', freq=None)
```

```
In [39]: # Can also initialize with timezone set
ts = pd.date_range('11/2/2022 9:30am', periods=2, freq='D', tz='US/Eastern')
ts
```

```
Out[39]: DatetimeIndex(['2022-11-02 09:30:00-04:00', '2022-11-03 09:30:00-04:00'], dtype='datetime64[ns, US/Eastern]', freq='D')
```

# Time Series in Python so far:

- `datetime .date .time .datetime .timedelta`
- format with `.strftime()`
- parse time with `pd.to_datetime()`
- `pandas Timestamp Timedelta DatetimeIndex`
- Indexing with `DatetimeIndex`
- Frequencies
- Timezones

Next: Operations on Time Series data

- Shifting
- Resampling
- Moving Windows

# Shifting/Lagging

- Moving data backward or forward in time (lagging/leading)
- Ex: calculate percent change

# Shifting/Lagging

- Moving data backward or forward in time (lagging/leading)
- Ex: calculate percent change

```
In [40]: ts = pd.Series([1,2,8],  
                        index=pd.date_range('1/1/2022', periods=3, freq='M')) # Month End frequency (MS: Month Start)  
ts
```

```
Out[40]: 2022-01-31    1  
         2022-02-28    2  
         2022-03-31    8  
         Freq: M, dtype: int64
```

# Shifting/Lagging

- Moving data backward or forward in time (lagging/leading)
- Ex: calculate percent change

```
In [40]: ts = pd.Series([1,2,8],  
                        index=pd.date_range('1/1/2022',periods=3,freq='M')) # Month End frequency (MS: Month Start)  
ts
```

```
Out[40]: 2022-01-31    1  
         2022-02-28    2  
         2022-03-31    8  
         Freq: M, dtype: int64
```

```
In [41]: ts.shift(1) # last month's value
```

```
Out[41]: 2022-01-31    NaN  
         2022-02-28    1.0  
         2022-03-31    2.0  
         Freq: M, dtype: float64
```

# Shifting

- percent change, use one of :
  - $(\text{new\_value} - \text{old\_value}) / \text{old\_value}$
  - $(\text{new\_value} / \text{old\_value}) - 1$

# Shifting

- percent change, use one of :
  - $(\text{new\_value} - \text{old\_value}) / \text{old\_value}$
  - $(\text{new\_value} / \text{old\_value}) - 1$

```
In [42]: # multiply by 100 to turn into a percent  
((ts / ts.shift(1)) - 1) * 100
```

```
Out[42]: 2022-01-31      NaN  
2022-02-28      100.0  
2022-03-31      300.0  
Freq: M, dtype: float64
```

# Example Dataset: Twitter Stock



# Example Dataset: Twitter Stock

```
In [43]: # from pandas_datareader import data
# df_twtr = data.DataReader('TWTR', start='2015', end='11/27/2022', data_source='yahoo')
# df_twtr.to_csv('../data/twtr_20150102-20221127.csv')
df_twtr = pd.read_csv('../data/twtr_20150102-20221127.csv', parse_dates=['Date'], index_col='Date')
df_twtr.head(3).round(2)
```

Out[43]:

	High	Low	Open	Close	Volume	Adj Close
Date						
2015-01-02	36.74	35.54	36.23	36.56	12062461.0	36.56
2015-01-05	37.11	35.64	36.26	36.38	15062744.0	36.38
2015-01-06	39.45	36.04	36.27	38.76	33050812.0	38.76

# Example Dataset: Twitter Stock

```
In [43]: # from pandas_datareader import data
# df_twtr = data.DataReader('TWTR', start='2015', end='11/27/2022', data_source='yahoo')
# df_twtr.to_csv('../data/twtr_20150102-20221127.csv')
df_twtr = pd.read_csv('../data/twtr_20150102-20221127.csv', parse_dates=['Date'], index_col='Date')
df_twtr.head(3).round(2)
```

```
Out[43]:
```

	High	Low	Open	Close	Volume	Adj Close
Date						
2015-01-02	36.74	35.54	36.23	36.56	12062461.0	36.56
2015-01-05	37.11	35.64	36.26	36.38	15062744.0	36.38
2015-01-06	39.45	36.04	36.27	38.76	33050812.0	38.76

```
In [44]: df_twtr.info() # Adj Close factors in corporate actions, such as stock splits, dividends, and rights offerings
```

```
<class 'pandas.core.frame.DataFrame'>
DatetimeIndex: 1970 entries, 2015-01-02 to 2022-10-27
Data columns (total 6 columns):
#   Column      Non-Null Count  Dtype
---  -
0   High        1970 non-null   float64
1   Low         1970 non-null   float64
2   Open        1970 non-null   float64
3   Close       1970 non-null   float64
4   Volume      1970 non-null   float64
5   Adj Close   1970 non-null   float64
dtypes: float64(6)
memory usage: 107.7 KB
```

# Example Dataset: Twitter Stock

# Example Dataset: Twitter Stock

```
In [45]: fig,ax = plt.subplots(1,1,figsize=(24,10))
df_twtr[['Close']].plot(ax=ax,marker='x');
ax.set_ylabel('stock price');
```



# Shifting Example: Percent Change Twitter Close

# Shifting Example: Percent Change Twitter Close

```
In [46]: ((df_twtr.Close / df_twtr.Close.shift(1)) - 1).tail(3).round(3) # # (today / yesterday) - 1
```


```
Out[46]: Date  
2022-10-25    0.024  
2022-10-26    0.011  
2022-10-27    0.007  
Name: Close, dtype: float64
```

# Shifting Example: Percent Change Twitter Close

```
In [46]: ((df_twtr.Close / df_twtr.Close.shift(1)) - 1).tail(3).round(3) # # (today / yesterday) - 1
```

```
Out[46]: Date
2022-10-25    0.024
2022-10-26    0.011
2022-10-27    0.007
Name: Close, dtype: float64
```

```
In [47]: # plot percent change of close in 2022
fig,ax = plt.subplots(1,1,figsize=(24,8))
close_2020 = df_twtr.loc['2020','Close']
((close_2020 / close_2020.shift(1)) - 1).plot(marker='x',ax=ax,zorder=2);
ax.axhline(ls=':',c='k',zorder=1)
ax.set_ylabel('percent change');
```



# Resampling

- Convert from one frequency to another
- **Downsampling**
  - from higher to lower (day to month)
  - need to aggregate
- **Upsampling**
  - from lower to higher (month to day)
  - need to fill missing
- Can also be used to set frequency from None



# Resampling: Initialize Frequency

# Resampling: Initialize Frequency

```
In [48]: df_twtr.index
```

```
Out[48]: DatetimeIndex(['2015-01-02', '2015-01-05', '2015-01-06', '2015-01-07',  
                        '2015-01-08', '2015-01-09', '2015-01-12', '2015-01-13',  
                        '2015-01-14', '2015-01-15',  
                        ...,  
                        '2022-10-14', '2022-10-17', '2022-10-18', '2022-10-19',  
                        '2022-10-20', '2022-10-21', '2022-10-24', '2022-10-25',  
                        '2022-10-26', '2022-10-27'],  
                        dtype='datetime64[ns]', name='Date', length=1970, freq=None)
```

# Resampling: Initialize Frequency

```
In [48]: df_twtr.index
```

```
Out[48]: DatetimeIndex(['2015-01-02', '2015-01-05', '2015-01-06', '2015-01-07',  
                        '2015-01-08', '2015-01-09', '2015-01-12', '2015-01-13',  
                        '2015-01-14', '2015-01-15',  
                        ...,  
                        '2022-10-14', '2022-10-17', '2022-10-18', '2022-10-19',  
                        '2022-10-20', '2022-10-21', '2022-10-24', '2022-10-25',  
                        '2022-10-26', '2022-10-27'],  
                      dtype='datetime64[ns]', name='Date', length=1970, freq=None)
```

```
In [49]: df_twtr_B = df_twtr.resample('B').asfreq() # set frequency to business day  
df_twtr_B.index
```

```
Out[49]: DatetimeIndex(['2015-01-02', '2015-01-05', '2015-01-06', '2015-01-07',  
                        '2015-01-08', '2015-01-09', '2015-01-12', '2015-01-13',  
                        '2015-01-14', '2015-01-15',  
                        ...,  
                        '2022-10-14', '2022-10-17', '2022-10-18', '2022-10-19',  
                        '2022-10-20', '2022-10-21', '2022-10-24', '2022-10-25',  
                        '2022-10-26', '2022-10-27'],  
                      dtype='datetime64[ns]', name='Date', length=2040, freq='B')
```

# Resampling: Downsampling

- Go from higher/shorter to lower/longer
- Need to aggregate (like groupby)
- Example: Downsampling from business day to business quarter

# Resampling: Downsampling

- Go from higher/shorter to lower/longer
- Need to aggregate (like groupby)
- Example: Downsampling from business day to business quarter

```
In [50]: df_twtr_BQ = df_twtr_B.resample('BQ')  
df_twtr_BQ
```

```
Out[50]: <pandas.core.resample.DatetimeIndexResampler object at 0x7fed6a232560>
```

# Resampling: Downsampling

- Go from higher/shorter to lower/longer
- Need to aggregate (like groupby)
- Example: Downsampling from business day to business quarter

```
In [50]: df_twtr_BQ = df_twtr_B.resample('BQ')  
df_twtr_BQ
```

```
Out[50]: <pandas.core.resample.DatetimeIndexResampler object at 0x7fed6a232560>
```

```
In [51]: print(df_twtr_BQ)
```

```
DatetimeIndexResampler [freq=<BusinessQuarterEnd: startingMonth=12>, axis=0, closed=right, label=right, convention=start, origin=start_day]
```

# Resampling: Downsampling

- Go from higher/shorter to lower/longer
- Need to aggregate (like groupby)
- Example: Downsampling from business day to business quarter

```
In [50]: df_twtr_BQ = df_twtr_B.resample('BQ')
df_twtr_BQ
```

```
Out[50]: <pandas.core.resample.DatetimeIndexResampler object at 0x7fed6a232560>
```

```
In [51]: print(df_twtr_BQ)
```

```
DatetimeIndexResampler [freq=<BusinessQuarterEnd: startingMonth=12>, axis=0, closed=right, label=right, convention=start, origin=start_day]
```

```
In [52]: df_twtr_BQ.mean().head(3).round(2)
```

```
Out[52]:
```

	High	Low	Open	Close	Volume	Adj Close
Date						
2015-03-31	45.10	43.55	44.23	44.34	20840997.51	44.34
2015-06-30	41.63	40.38	41.17	40.87	22287099.56	40.87
2015-09-30	30.64	29.42	30.05	30.00	20065038.11	30.00

# Resampling: Downsampling



# Resampling: Downsampling

```
In [53]: fig,ax = plt.subplots(1,1,figsize=(24,8))
df_twtr_B.Close.plot(style='-', label='by B',ax=ax)
df_twtr_BQ.Close.mean().plot(style='--',marker='x',label='by BQ',ax=ax)
plt.legend(loc='upper right');
```



# Resampling: Upsampling

- Go from lower/longer to higher/shorter
- Need to decide how to handle missing values
- Example: Upsample from business day to hour

# Resampling: Upsampling

- Go from lower/longer to higher/shorter
- Need to decide how to handle missing values
- Example: Upsample from business day to hour

```
In [54]: df_twtr_B.index[:3]
```

```
Out[54]: DatetimeIndex(['2015-01-02', '2015-01-05', '2015-01-06'], dtype='datetime64[ns]', name='Date', freq='B')
```

# Resampling: Upsampling

- Go from lower/longer to higher/shorter
- Need to decide how to handle missing values
- Example: Upsample from business day to hour

```
In [54]: df_twtr_B.index[:3]
```

```
Out[54]: DatetimeIndex(['2015-01-02', '2015-01-05', '2015-01-06'], dtype='datetime64[ns]', name='Date', freq='B')
```

```
In [55]: df_twtr_B.Close.resample('H').asfreq().iloc[0:3].round(2)
```

```
Out[55]: Date
2015-01-02 00:00:00    36.56
2015-01-02 01:00:00     NaN
2015-01-02 02:00:00     NaN
Freq: H, Name: Close, dtype: float64
```

# Resampling: Upsampling

- Go from lower/longer to higher/shorter
- Need to decide how to handle missing values
- Example: Upsample from business day to hour

```
In [54]: df_twtr_B.index[:3]
```

```
Out[54]: DatetimeIndex(['2015-01-02', '2015-01-05', '2015-01-06'], dtype='datetime64[ns]', name='Date', freq='B')
```

```
In [55]: df_twtr_B.Close.resample('H').asfreq().iloc[0:3].round(2)
```

```
Out[55]: Date
2015-01-02 00:00:00    36.56
2015-01-02 01:00:00     NaN
2015-01-02 02:00:00     NaN
Freq: H, Name: Close, dtype: float64
```

```
In [56]: df_twtr_B.Close.resample('H').asfreq().iloc[70:73].round(2)
```

```
Out[56]: Date
2015-01-04 22:00:00     NaN
2015-01-04 23:00:00     NaN
2015-01-05 00:00:00    36.38
Freq: H, Name: Close, dtype: float64
```

# Resampling: Upsampling

- `ffill()` : Forward Fill

# Resampling: Upsampling

- `ffill()` : Forward Fill

```
In [57]: df_twtr_B.Close.resample('H').ffill().head(3).round(2)
```

```
Out[57]: Date
2015-01-02 00:00:00    36.56
2015-01-02 01:00:00    36.56
2015-01-02 02:00:00    36.56
Freq: H, Name: Close, dtype: float64
```

# Resampling: Upsampling

- `ffill()` : Forward Fill

```
In [57]: df_twtr_B.Close.resample('H').ffill().head(3).round(2)
```

```
Out[57]: Date
2015-01-02 00:00:00    36.56
2015-01-02 01:00:00    36.56
2015-01-02 02:00:00    36.56
Freq: H, Name: Close, dtype: float64
```

- `bfill()` : Backward Fill



# Resampling: Upsampling

- `ffill()` : Forward Fill

```
In [57]: df_twtr_B.Close.resample('H').ffill().head(3).round(2)
```

```
Out[57]: Date
2015-01-02 00:00:00    36.56
2015-01-02 01:00:00    36.56
2015-01-02 02:00:00    36.56
Freq: H, Name: Close, dtype: float64
```

- `bfill()` : Backward Fill

```
In [58]: df_twtr_B.Close.resample('H').bfill().head(3).round(3)
```

```
Out[58]: Date
2015-01-02 00:00:00    36.56
2015-01-02 01:00:00    36.38
2015-01-02 02:00:00    36.38
Freq: H, Name: Close, dtype: float64
```

# Moving/Rolling Windows

- Apply function on a fixed window moving across time
- Method of smoothing out the data
- **center** : place values at center of window

# Moving/Rolling Windows

- Apply function on a fixed window moving across time
- Method of smoothing out the data
- **center** : place values at center of window

```
In [59]: df_twtr_B.Close['2020-11-02':'2020-11-06'].round(2)
```

```
Out[59]: Date
2020-11-02    39.47
2020-11-03    41.73
2020-11-04    42.76
2020-11-05    43.71
2020-11-06    43.12
Freq: B, Name: Close, dtype: float64
```

# Moving/Rolling Windows

- Apply function on a fixed window moving across time
- Method of smoothing out the data
- **center** : place values at center of window

```
In [59]: df_twtr_B.Close['2020-11-02':'2020-11-06'].round(2)
```

```
Out[59]: Date
2020-11-02    39.47
2020-11-03    41.73
2020-11-04    42.76
2020-11-05    43.71
2020-11-06    43.12
Freq: B, Name: Close, dtype: float64
```

```
In [60]: rolling_3 = df_twtr_B.Close['2020-11-02':'2020-11-06'].rolling(3, center=True)
rolling_3
```

```
Out[60]: Rolling [window=3,center=True,axis=0,method=single]
```

# Moving/Rolling Windows

- Apply function on a fixed window moving across time
- Method of smoothing out the data
- **center** : place values at center of window

```
In [59]: df_twtr_B.Close['2020-11-02':'2020-11-06'].round(2)
```

```
Out[59]: Date
2020-11-02    39.47
2020-11-03    41.73
2020-11-04    42.76
2020-11-05    43.71
2020-11-06    43.12
Freq: B, Name: Close, dtype: float64
```

```
In [60]: rolling_3 = df_twtr_B.Close['2020-11-02':'2020-11-06'].rolling(3, center=True)
rolling_3
```

```
Out[60]: Rolling [window=3,center=True,axis=0,method=single]
```

```
In [61]: rolling_3.mean()['2020-11-02':'2020-11-06'].round(2)
```


```
Out[61]: Date
2020-11-02    NaN
2020-11-03    41.32
2020-11-04    42.73
2020-11-05    43.20
2020-11-06    NaN
Freq: B, Name: Close, dtype: float64
```

# Moving Windows

# Moving Windows

```
In [62]: sns.set_style("whitegrid")
rolling = df_twtr_B.Close.rolling(5, center=True)

fig, ax = plt.subplots(1, 1, figsize=(24, 8));
df_twtr_B.loc['2020'].Close.plot(style='-', alpha=0.3, label='business day');
rolling.mean().loc['2020'].plot(style='--', label='5 day rolling window mean');
(rolling.mean().loc['2020'] + 2*rolling.std().loc['2020']).plot(style=':', c='g', label='_nolegend_');
(rolling.mean().loc['2020'] - 2*rolling.std().loc['2020']).plot(style=':', c='g', label='_nolegend_');
ax.legend();
```



# Example: Bike Travel (From PDSH Chapter 3.11)

- Bicycle traffic over Fremont Bridge in Seattle in 2012
- Data gathered using: `!curl -o ../data/FremontBridge.csv https://data.seattle.gov/api/views/65db-xm6k/rows.csv?accessType=DOWNLOAD`



# Example: Bike Travel (From PDSH Chapter 3.11)

- Bicycle traffic over Fremont Bridge in Seattle in 2012
- Data gathered using: `!curl -o ../data/FremontBridge.csv https://data.seattle.gov/api/views/65db-xm6k/rows.csv?accessType=DOWNLOAD`

```
In [63]: df_bike_counts = pd.read_csv('../data/FremontBridge_2012-2015.csv', parse_dates=['Date'], index_col='Date')
df_bike_counts.columns = ['Total', 'East', 'West']
df_bike_counts.info()
```

```
<class 'pandas.core.frame.DataFrame'>
DatetimeIndex: 28440 entries, 2012-10-03 00:00:00 to 2015-12-31 23:00:00
Data columns (total 3 columns):
#   Column  Non-Null Count  Dtype
---  -
0   Total    28433 non-null     float64
1   East     28433 non-null     float64
2   West     28433 non-null     float64
dtypes: float64(3)
memory usage: 888.8 KB
```

# Example: Bike Travel (From PDSH Chapter 3.11)

- Bicycle traffic over Fremont Bridge in Seattle in 2012
- Data gathered using: `!curl -o ../data/FremontBridge.csv https://data.seattle.gov/api/views/65db-xm6k/rows.csv?accessType=DOWNLOAD`

```
In [63]: df_bike_counts = pd.read_csv('../data/FremontBridge_2012-2015.csv', parse_dates=['Date'], index_col='Date')
df_bike_counts.columns = ['Total', 'East', 'West']
df_bike_counts.info()
```

```
<class 'pandas.core.frame.DataFrame'>
DatetimeIndex: 28440 entries, 2012-10-03 00:00:00 to 2015-12-31 23:00:00
Data columns (total 3 columns):
#   Column  Non-Null Count  Dtype
---  -
0   Total    28433 non-null     float64
1   East     28433 non-null     float64
2   West     28433 non-null     float64
dtypes: float64(3)
memory usage: 888.8 KB
```

```
In [64]: df_bike_counts.head(3)
```

Out[64]:

	Total	East	West
Date			
2012-10-03 00:00:00	13.0	4.0	9.0
2012-10-03 01:00:00	10.0	4.0	6.0
2012-10-03 02:00:00	2.0	1.0	1.0

# Example: Fill Missing Values

# Example: Fill Missing Values

```
In [65]: f'proportion missing: {sum(df_bike_counts.Total.isna()) / len(df_bike_counts):0.5f}'
```

```
Out[65]: 'proportion missing: 0.00025'
```

# Example: Fill Missing Values

```
In [65]: f'proportion missing: {sum(df_bike_counts.Total.isna()) / len(df_bike_counts):0.5f}'
```

```
Out[65]: 'proportion missing: 0.00025'
```


```
In [66]: df_bike_counts = df_bike_counts.fillna(method='ffill')
df_bike_counts.info()
```

```
<class 'pandas.core.frame.DataFrame'>
DatetimeIndex: 28440 entries, 2012-10-03 00:00:00 to 2015-12-31 23:00:00
Data columns (total 3 columns):
#   Column  Non-Null Count  Dtype
---  -
0   Total   28440 non-null   float64
1   East    28440 non-null   float64
2   West    28440 non-null   float64
dtypes: float64(3)
memory usage: 888.8 KB
```

**Plot data from 2015**

# Plot data from 2015

```
In [67]: fig,ax = plt.subplots(1,1,figsize=(24,8))
df_bike_counts.loc['2015'].plot(style=['-', '--', ':'],ax=ax)
plt.ylabel('Hourly Bicycle Count');
```




**Downsample to weekly sum to smooth things out**



# Downsample to weekly sum to smooth things out


```
In [68]: weekly = df_bike_counts.resample('W').sum()  
weekly.plot(style=[':', '--', '-'], figsize=(24,8))  
plt.ylabel('Weekly bicycle count');
```



**Resample at daily for a more granular view and apply a rolling window of 30 days**

# Resample at daily for a more granular view and apply a rolling window of 30 days

```
In [69]: daily = df_bike_counts.resample('D').sum()  
daily.rolling(30,center=True).mean().plot(style=[':', '--', '-'], figsize=(24,8))  
plt.ylabel('mean daily count');
```



# From Datetime to Time

# From Datetime to Time

```
In [71]: #If we want to only look at time of day  
df_bike_counts.index.time
```

```
Out[71]: array([datetime.time(0, 0), datetime.time(1, 0), datetime.time(2, 0), ...,  
               datetime.time(21, 0), datetime.time(22, 0), datetime.time(23, 0)],  
              dtype=object)
```

# From Datetime to Time

```
In [71]: #If we want to only look at time of day  
df_bike_counts.index.time
```

```
Out[71]: array([datetime.time(0, 0), datetime.time(1, 0), datetime.time(2, 0), ...,  
               datetime.time(21, 0), datetime.time(22, 0), datetime.time(23, 0)],  
              dtype=object)
```


```
In [72]: #Get mean data by time (hourly)  
by_time = df_bike_counts.groupby(df_bike_counts.index.time).mean().round(2)  
display(by_time.head())
```

	Total	East	West
00:00:00	13.34	5.94	7.40
01:00:00	7.15	3.34	3.81
02:00:00	4.97	2.61	2.36
03:00:00	3.43	1.90	1.52
04:00:00	6.13	3.53	2.59

**Plot by hour of the day**

# Plot by hour of the day

```
In [73]: hourly_ticks = 60 * 60 * 4 * np.arange(6) # sec * min * every4hours
by_time.plot(xticks=hourly_ticks, style=[':', '--', '-'], figsize=(24,8));
plt.ylabel('mean hourly count');
```






**Can also look at average by day of week**

# Can also look at average by day of week

```
In [74]: # note that for dayofweek: 0 == Mon, 1 == Tues, ..., 6 == 'Sun'
by_weekday = df_bike_counts.groupby(df_bike_counts.index.dayofweek).mean()
by_weekday = by_weekday.set_index(pd.Index(['Mon', 'Tues', 'Wed', 'Thurs', 'Fri', 'Sat', 'Sun']))

fig, ax = plt.subplots(1, 1, figsize=(24, 8))
by_weekday.plot(style=[':', '--', '-'], ax=ax);
ax.set_xlabel('Day of Week'); ax.set_ylabel('mean daily count');
```


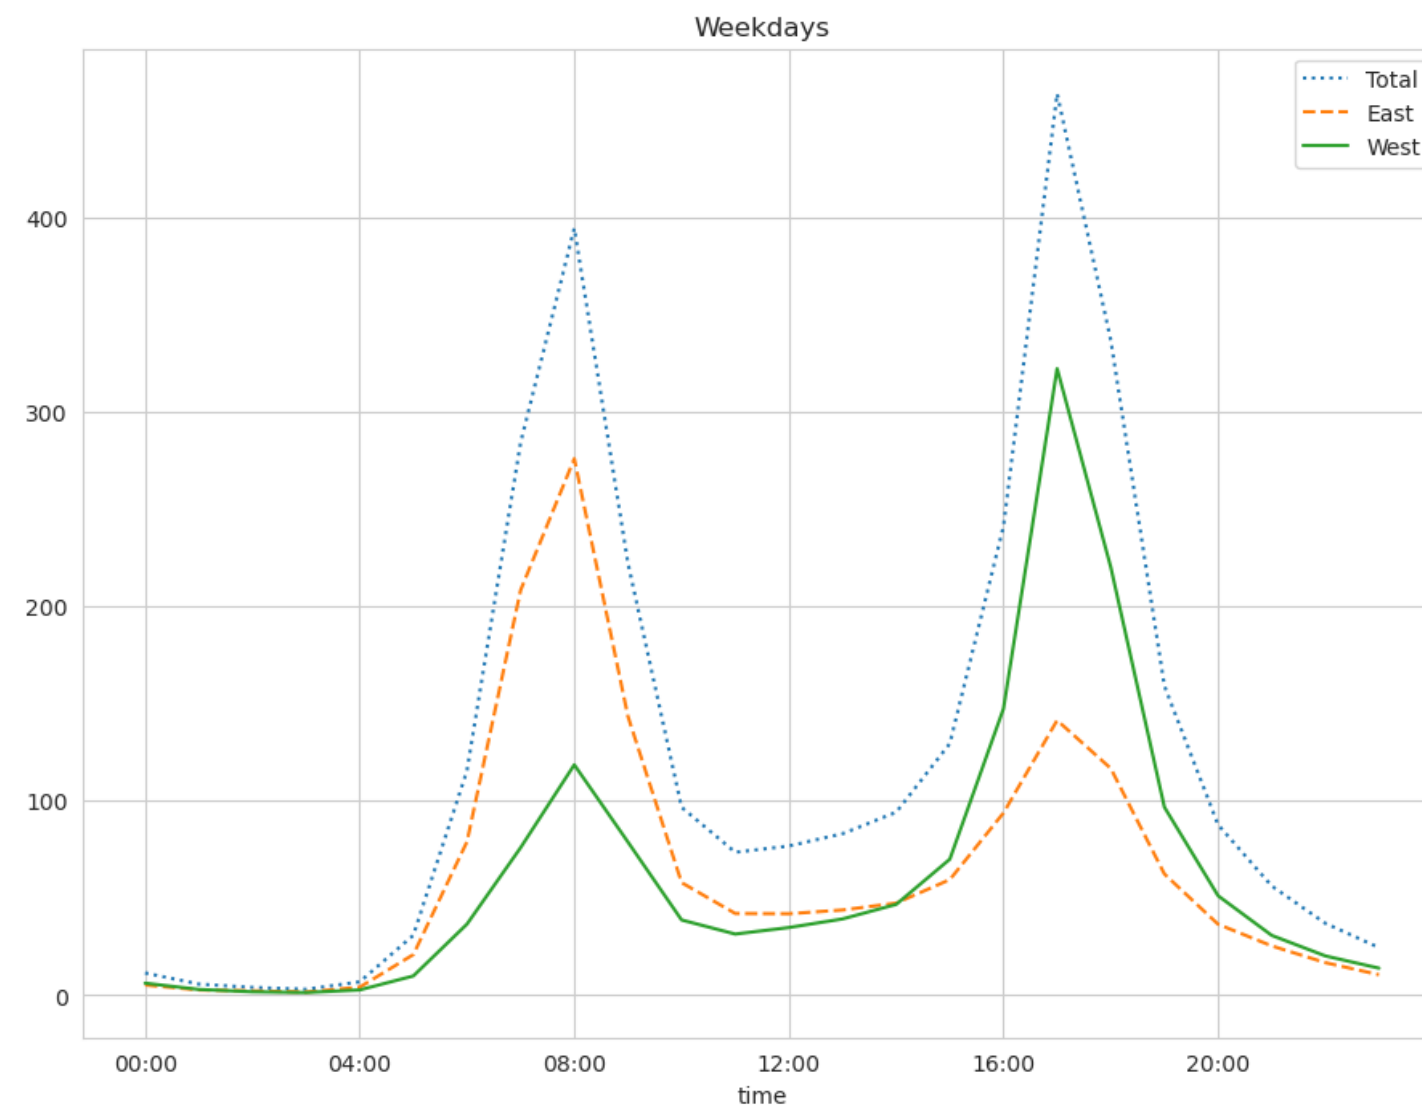


**Separate out weekdays and weekends**

# Separate out weekdays and weekends

```
In [75]: # create a weekend mask
weekend = np.where(df_bike_counts.index.weekday < 5, 'Weekday', 'Weekend')

# get hourly mean values split by weekday, weekend
by_time = df_bike_counts.groupby([weekend, df_bike_counts.index.time]).mean()
fig, ax = plt.subplots(1, 2, figsize=(24, 8))
by_time.loc['Weekday'].plot(ax=ax[0], title='Weekdays', xticks=hourly_ticks, style=[':', '--', '-'])
by_time.loc['Weekend'].plot(ax=ax[1], title='Weekends', xticks=hourly_ticks, style=[':', '--', '-']);
```



**Can we predict daily Total bike traffic?**

# Can we predict daily Total bike traffic?

```
In [76]: df_bike_counts = pd.read_csv('../data/FremontBridge_2012-2015.csv', index_col='Date', parse_dates=True)
df_bike_weather = pd.read_csv('../data/BicycleWeather.csv', index_col='DATE', parse_dates=True)

df_bike = (
    df_bike_counts.loc[:, ['Fremont Bridge Total']] # keep Total as target
    .rename({'Fremont Bridge Total': 'Total'}, axis=1) # rename target column
    .resample('D').sum() # downsample to daily totals
)
print(df_bike.head(3))
```

Date	Total
2012-10-03	3521.0
2012-10-04	3475.0
2012-10-05	3148.0

# Can we predict daily Total bike traffic?

```
In [76]: df_bike_counts = pd.read_csv('../data/FremontBridge_2012-2015.csv', index_col='Date', parse_dates=True)
df_bike_weather = pd.read_csv('../data/BicycleWeather.csv', index_col='DATE', parse_dates=True)

df_bike = (
    df_bike_counts.loc[:, ['Fremont Bridge Total']] # keep Total as target
    .rename({'Fremont Bridge Total': 'Total'}, axis=1) # rename target column
    .resample('D').sum() # downsample to daily totals
)
print(df_bike.head(3))
```

Date	Total
2012-10-03	3521.0
2012-10-04	3475.0
2012-10-05	3148.0

## On to Feature Engineering...

**Add 'day of week'**



# Add 'day of week'

```
In [77]: day_names_map = dict(enumerate(['Mon', 'Tue', 'Wed', 'Thu', 'Fri', 'Sat', 'Sun']))
print(f"{day_names_map = :}")
df_bike['DayOfWeek'] = df_bike.index.dayofweek.map(day_names_map)
df_bike.head(3)
```

day\_names\_map = {0: 'Mon', 1: 'Tue', 2: 'Wed', 3: 'Thu', 4: 'Fri', 5: 'Sat', 6: 'Sun'}

Out[77]:

	Total	DayOfWeek
Date		
2012-10-03	3521.0	Wed
2012-10-04	3475.0	Thu
2012-10-05	3148.0	Fri

**Add 'is it a holiday' dummy feature**

# Add 'is it a holiday' dummy feature

```
In [78]: from pandas.tseries.holiday import USFederalHolidayCalendar
cal = USFederalHolidayCalendar()
holidays = cal.holidays('2012', '2016')

df_bike = df_bike.join(pd.Series(1, index=holidays, name='IsHoliday'))
df_bike['IsHoliday'].fillna(0, inplace=True)
print(df_bike.head(3))
```

	Total	DayOfWeek	IsHoliday
Date			
2012-10-03	3521.0	Wed	0.0
2012-10-04	3475.0	Thu	0.0
2012-10-05	3148.0	Fri	0.0

**Add number of hours of daylight**


# Add number of hours of daylight

```
In [79]: from datetime import datetime

def hours_of_daylight(date, axis=23.44, latitude=47.61):
    """Compute the hours of daylight for the given date"""
    days = (date - datetime(2000, 12, 21)).days # days till winter solstice
    m = (1. - np.tan(np.radians(latitude))
          * np.tan(np.radians(axis) * np.cos(days * 2 * np.pi / 365.25)))
    return 24. * np.degrees(np.arccos(1 - np.clip(m, 0, 2))) / 180.

df_bike['HoursDaylight'] = list(map(hours_of_daylight, df_bike.index));

ax = df_bike[['HoursDaylight']].plot(figsize=(18,4));
ax.set_ylim(8, 16);
```



**Add weather information (Q: can we predict this for future dates?)**

# Add weather information (Q: can we predict this for future dates?)

```
In [80]: # temperatures are in 1/10 deg C; convert to C
df_bike_weather['TMIN'] /= 10
df_bike_weather['TMAX'] /= 10
df_bike_weather['TempC'] = 0.5 * (df_bike_weather['TMIN'] + df_bike_weather['TMAX'])

# precip is in 1/10 mm; convert to inches
df_bike_weather['PRCP'] /= 254
df_bike_weather['IsDryDay'] = (df_bike_weather['PRCP'] == 0).astype(int)

df_bike = df_bike.join(df_bike_weather[['PRCP', 'TempC', 'IsDryDay']],how='inner')
df_bike.head(3).round(2)
```

Out[80]:

	Total	DayOfWeek	IsHoliday	HoursDaylight	PRCP	TempC	IsDryDay
2012-10-03	3521.0	Wed	0.0	11.28	0.0	13.35	1
2012-10-04	3475.0	Thu	0.0	11.22	0.0	13.60	1
2012-10-05	3148.0	Fri	0.0	11.16	0.0	15.30	1

**Add time of year**



# Add time of year

```
In [81]: df_bike['TimeOfYear'] = (df_bike.index - df_bike.index[0]).days / 365.0 # Days since the beginning of the year
df_bike.head(3)
```

Out[81]:

	Total	DayOfWeek	IsHoliday	HoursDaylight	PRCP	TempC	IsDryDay	TimeOfYear
2012-10-03	3521.0	Wed	0.0	11.277359	0.0	13.35	1	0.000000
2012-10-04	3475.0	Thu	0.0	11.219142	0.0	13.60	1	0.002740
2012-10-05	3148.0	Fri	0.0	11.161038	0.0	15.30	1	0.005479

**Generate and evaluate a model**

# Generate and evaluate a model

```
In [82]: from sklearn.ensemble import GradientBoostingRegressor
from sklearn.dummy import DummyRegressor
from sklearn.metrics import mean_absolute_error

# drop any rows with missing data
df_bike.dropna(axis=0, how='any', inplace=True)

X_bike = pd.get_dummies(df_bike.loc[:,df_bike.columns != 'Total'])
display(X_bike.head(1).round(2))
y_bike = df_bike.Total

X_bike_train = X_bike.loc['2012':'2014']
y_bike_train = y_bike.loc['2012':'2014']
X_bike_test = X_bike.loc['2015']
y_bike_test = y_bike.loc['2015']

dummy_bike = DummyRegressor().fit(X_bike_train,y_bike_train)
gb_bike = GradientBoostingRegressor().fit(X_bike_train,y_bike_train)
print(f'dummy training mae      : {mean_absolute_error(y_bike_train,dummy_bike.predict(X_bike_train)).round(2)}')
print(f'one-back training mae  : {mean_absolute_error(y_bike_train,y_bike_train.shift(1).fillna(0)).round(2)}')
print(f'gb training set mae     : {mean_absolute_error(y_bike_train,gb_bike.predict(X_bike_train)).round(2)}')
print(f'gb test set R^2         : {mean_absolute_error(y_bike_test,gb_bike.predict(X_bike_test)).round(2)}')
```


	IsHoliday	HoursDaylight	PRCP	TempC	IsDryDay	TimeOfYear	DayOfWeek_Fri	DayOfWeek_Mon	DayOfWeek_Sat	DayOfWeek_Sun	DayOfWeek_Thu	DayOfWeek_Tue	DayC
2012-10-03	0.0	11.28	0.0	13.35	1	0.0	0	0	0	0	0	0	1

```
dummy training mae      : 1019.45
one-back training mae   : 710.39
gb training set mae     : 213.37
gb test set R^2         : 308.24
```

# Plot predictions vs observed

# Plot predictions vs observed

```
In [83]: df_bike['predicted'] = gb_bike.predict(X_bike)
df_bike[['Total', 'predicted']].rolling(30, center=True, win_type='gaussian').mean(std=7).plot(alpha=0.5, figsize=(24, 8))
plt.tight_layout()
```



# Time Series Operations Review

- Shifting
- Resampling
  - Downsampling
  - Upsampling
- Moving/Rolling Windows
- for more info, including time-series cross-validation:
  - [sklearn: Time-related feature engineering](#)
  - PML Chapter 13 - Modeling Sequential Data Using Recurrent Neural Network (with Tensorflow)
- for more models:
  - [skforecast](#)
  - [statsmodels](#)

# Questions re Time Series Transformations?

# Delivering Data With Flask



# Delivering Data With Flask

- Flask : lightweight web server
- can be used to create a small API to:
  - return transformed data
  - return predictions
  - return datasets
  - ...

# Aside: Python script basics

# Aside: Python script basics

In [84]: `!cat ../src/sample_script.py`

```
# import necessary libraries and function
from datetime import datetime

# python as usual
def current_time():
    return datetime.now()

# will run as script or on import
run_or_imported_at = current_time()
print(f"this was run or imported at {run_or_imported_at}")
print(f"__name__ = :s")

if __name__ == "__main__":
    # will only run if this is a script
    # won't be run if imported
    print("running as a script")
```

# Aside: Python script basics

In [84]: `!cat ../src/sample_script.py`

```
# import necessary libraries and function
from datetime import datetime

# python as usual
def current_time():
    return datetime.now()

# will run as script or on import
run_or_imported_at = current_time()
print(f"this was run or imported at {run_or_imported_at}")
print(f"__name__ = :s}")

if __name__ == "__main__":
    # will only run if this is a script
    # won't be run if imported
    print("running as a script")
```

In [85]: `import sys`  
`sys.path.append('../src/') # need to tell python where to look for this file`

`import sample_script`

```
this was run or imported at 2022-11-30 18:04:12.290787
__name__ = sample_script
```

# Aside: Python script basics

In [84]: `!cat ../src/sample_script.py`

```
# import necessary libraries and function
from datetime import datetime

# python as usual
def current_time():
    return datetime.now()

# will run as script or on import
run_or_imported_at = current_time()
print(f"this was run or imported at {run_or_imported_at}")
print(f"__name__ = :s}")

if __name__ == "__main__":
    # will only run if this is a script
    # won't be run if imported
    print("running as a script")
```

In [85]: `import sys`  
`sys.path.append('../src/') # need to tell python where to look for this file`

`import sample_script`

```
this was run or imported at 2022-11-30 18:04:12.290787
__name__ = sample_script
```

In [86]: `%run ../src/sample_script.py`

```
this was run or imported at 2022-11-30 18:04:12.299807
__name__ = __main__
running as a script
```

<Figure size 640x480 with 0 Axes>

# Aside: Function Decorators

- act like wrappers around functions
- decorators are prefixed by the "@" symbol
- placed above the function to be wrapped

# Aside: Function Decorators

- act like wrappers around functions
- decorators are prefixed by the "@" symbol
- placed above the function to be wrapped

```
In [87]: def my_decorator(func):
        def wrapper():
            print("Happens before the function is called.")
            func()
            print("Happens after the function is called.")
        return wrapper

        @my_decorator
        def say_hello():
            print("Hello")

        say_hello()
```

```
Happens before the function is called.
Hello
Happens after the function is called.
```

# Creating APIs: Flask



# Creating APIs: Flask

```
In [88]: !cat ../src/hello_flask.py

from flask import Flask, escape, request

app = Flask(__name__)

@app.route('/')
def hello():
    name = request.args.get("name", "World")
    return f'Hello, {escape(name)}!\n'

if __name__ == '__main__':
    app.run()
```

# Creating APIs: Flask

```
In [88]: !cat ../src/hello_flask.py

from flask import Flask, escape, request

app = Flask(__name__)

@app.route('/')
def hello():
    name = request.args.get("name", "World")
    return f'Hello, {escape(name)}!\n'

if __name__ == '__main__':
    app.run()
```

1. at command line, run: `$ python hello_flask.py`
2. in another terminal, ipython (or notebook), run:

# Creating APIs: Flask

```
In [88]: !cat ../src/hello_flask.py

from flask import Flask, escape, request

app = Flask(__name__)

@app.route('/')
def hello():
    name = request.args.get("name", "World")
    return f'Hello, {escape(name)}!\n'

if __name__ == '__main__':
    app.run()
```

1. at command line, run: `$ python hello_flask.py`

2. in another terminal, ipython (or notebook), run:

```
# need to change "Cell Type" to code to run in notebook
import requests
r = requests.get('http://127.0.0.1:5000/?name=Bryan')
print(r.text)
```

# Creating APIs: Flask with Multiple Routes

# Creating APIs: Flask with Multiple Routes

In [89]: !cat ../src/die\_flask.py

```
import numpy as np
from flask import Flask, request, jsonify

app = Flask(__name__)

@app.route("/")
def help():
    return "Give the number of sides the die should have.\n"

@app.route("/<int:sides>")
def roll_die(sides):
    return str(np.random.randint(1,sides+1))

@app.route("/json/<int:sides>")
def roll_die_json(sides):
    return jsonify({'sides': sides, 'roll': np.random.randint(1,sides+1)})

if __name__ == '__main__':
    app.run()
```

# GET vs POST

- **GET**: pass information in the url

```
127.0.0.1:5000/?firstname=Bryan&lastname=Gibson
```

- **POST**: pass information as additional http request (often JSON)

```
127.0.0.1:5000/  
{ 'firstname': 'Bryan', 'lastname': 'Gibson' }
```

## Example Model: Titanic

# Example Model: Titanic

```
In [90]: from sklearn.model_selection import train_test_split
titanic_url = ('https://raw.githubusercontent.com/amueller/scipy-2017-sklearn/091d371/notebooks/datasets/titanic3.csv')
df_titanic = pd.read_csv(titanic_url)[['age', 'fare', 'embarked', 'sex', 'pclass', 'survived']]
print(df_titanic.info())
X_titanic = df_titanic.drop('survived', axis=1)
y_titanic = df_titanic['survived']
X_titanic_train, X_titanic_test, y_titanic_train, y_titanic_test = train_test_split(X_titanic, y_titanic,
                                                                                      test_size=0.2,
                                                                                      stratify=y_titanic,
                                                                                      random_state=42)

display(X_titanic_train.head(3))
print(f"y_train prop positive: {y_titanic_train.mean().round(2)}")
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1309 entries, 0 to 1308
Data columns (total 6 columns):
#   Column      Non-Null Count  Dtype
---  -
0   age         1046 non-null   float64
1   fare        1308 non-null   float64
2   embarked    1307 non-null   object
3   sex         1309 non-null   object
4   pclass      1309 non-null   int64
5   survived    1309 non-null   int64
dtypes: float64(2), int64(2), object(2)
memory usage: 61.5+ KB
None
```

	age	fare	embarked	sex	pclass
999	NaN	7.7500	Q	female	3
392	24.0	27.7208	C	female	2
628	11.0	31.2750	S	female	3

y\_train prop positive: 0.38



## Example Model: Titanic

# Example Model: Titanic

```
In [91]: from sklearn.pipeline import Pipeline
from sklearn.compose import ColumnTransformer
from sklearn.impute import SimpleImputer
from sklearn.preprocessing import OneHotEncoder, StandardScaler
from sklearn.model_selection import GridSearchCV
from sklearn.linear_model import LogisticRegression

numeric_features = ['age', 'fare']
numeric_transformer = Pipeline(steps=[('imputer', SimpleImputer(strategy='median')),
                                      ('scaler', StandardScaler())])

categorical_features = ['embarked', 'sex', 'pclass']
categorical_transformer = Pipeline(steps=[('imputer', SimpleImputer(strategy='constant', fill_value='missing')),
                                          ('onehot', OneHotEncoder(handle_unknown='ignore'))])

preprocessor = ColumnTransformer(transformers=[('num', numeric_transformer, numeric_features),
                                             ('cat', categorical_transformer, categorical_features)])

pipe = Pipeline(steps=[('preprocessor', preprocessor),
                       ('classifier', LogisticRegression(solver='lbfgs', random_state=42))])

param_grid = {
    'preprocessor__num__imputer__strategy': ['mean', 'median'],
    'classifier__C': [0.1, 1.0, 10, 100],
}

gs_pipeline = GridSearchCV(pipe, param_grid, cv=3)
gs_pipeline.fit(X_titanic_train, y_titanic_train)
print("best test set score from grid search: {:.3f}".format(gs_pipeline.score(X_titanic_test, y_titanic_test)))
print("best parameter settings: {}".format(gs_pipeline.best_params_))
```

best test set score from grid search: 0.802

best parameter settings: {'classifier\_\_C': 1.0, 'preprocessor\_\_num\_\_imputer\_\_strategy': 'median'}

# Creating APIs: Flask

- Export trained models (and other data structures) using `pickle`

# Creating APIs: Flask

- Export trained models (and other data structures) using `pickle`

```
In [92]: import pickle as pkl  
  
# write/dump to disk  
with open('../data/titanic_pipeline_clf.pkl', 'wb') as f:  
    pkl.dump(gs_pipeline, f)
```



# Creating APIs: Deliver Predictions Using Flask

# Creating APIs: Deliver Predictions Using Flask

In [94]: !cat ../src/titanic\_clf.py

```
from flask import Flask, escape, request, jsonify
import pickle as pkl
import pandas as pd

# need to train and pickle classifier first
with open('../data/titanic_pipeline_clf.pkl','rb') as f:
    clf = pkl.load(f)

app = Flask(__name__)

@app.route('/',methods=['POST'])
def predict():
    prediction = None
    query = pd.DataFrame(request.form,index=[0])
    print(query,flush=True)
    if query is not None:
        prediction = clf.predict(query)
    if prediction:
        return jsonify([str(x) for x in prediction])
    else:
        return 'no predictions made'

if __name__ == '__main__':
    app.run()
```

# Creating APIs: Deliver Predictions Using Flask Cont.



# Creating APIs: Deliver Predictions Using Flask Cont.

```
In [95]: query_label = df_titanic.iloc[0].loc['survived']
```

# Creating APIs: Deliver Predictions Using Flask Cont.

```
In [95]: query_label = df_titanic.iloc[0].loc['survived']
```

```
In [96]: query = df_titanic.iloc[0, :-1].to_dict()  
query
```

```
Out[96]: {'age': 29.0, 'fare': 211.3375, 'embarked': 'S', 'sex': 'female', 'pclass': 1}
```

# Creating APIs: Deliver Predictions Using Flask Cont.

```
In [95]: query_label = df_titanic.iloc[0].loc['survived']
```

```
In [96]: query = df_titanic.iloc[0, :-1].to_dict()  
query
```

```
Out[96]: {'age': 29.0, 'fare': 211.3375, 'embarked': 'S', 'sex': 'female', 'pclass': 1}
```

```
In [97]: query_label
```

```
Out[97]: 1
```

# Creating APIs: Deliver Predictions Using Flask Cont.

```
In [95]: query_label = df_titanic.iloc[0].loc['survived']
```

```
In [96]: query = df_titanic.iloc[0,:-1].to_dict()  
query
```

```
Out[96]: {'age': 29.0, 'fare': 211.3375, 'embarked': 'S', 'sex': 'female', 'pclass': 1}
```

```
In [97]: query_label
```

```
Out[97]: 1
```

```
In [98]: # Start script from command line by first activating the eods-f22 environment and running: python titanic_clf.py  
# Then uncomment the following:  
  
#import requests  
#print(requests.post('http://127.0.0.1:5000/', data=query).text)
```

# Questions re Model API via Flask?