

DIABETES PREDICTION PROJECT

AGENDA

- ❑ Concept Study
- ❑ Data Preparation
- ❑ EDA
- ❑ Model Building
- ❑ Communicate Results
- ❑ Operationalize

Concept Study

In this project, we will be predicting that whether the patient has diabetes or not on the basis of the features(criteria) we will provide to our machine learning model.

Criteria

Pregnancies

Glucose

Blood Pressure

Skin Thickness

Insulin

BMI

Diabetes Pedigree Function

Age

Connect to App

information
Submission

Prediction

Concept Study

❑ Importing libraries

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder, StandardScaler
from sklearn.metrics import confusion_matrix
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.svm import SVC
from sklearn import metrics
```

Data Preparation

❑ Read the dataset

```
df = pd.read_csv('diabetes.csv')
```

❑ Display the 5 first rows

```
df.head()
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
0	6	148	72	35	0	33.6	0.627	50	1
1	1	85	66	29	0	26.6	0.351	31	0
2	8	183	64	0	0	23.3	0.672	32	1
3	1	89	66	23	94	28.1	0.167	21	0
4	0	137	40	35	168	43.1	2.288	33	1

Data Preparation

❑ Display the dimension of the dataset

```
df.shape
```

```
(768, 9)
```

❑ Information about the dataset

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 768 entries, 0 to 767
```

```
Data columns (total 9 columns):
```

#	Column	Non-Null Count	Dtype
0	Pregnancies	768 non-null	int64
1	Glucose	768 non-null	int64
2	BloodPressure	768 non-null	int64
3	SkinThickness	768 non-null	int64
4	Insulin	768 non-null	int64
5	BMI	768 non-null	float64
6	DiabetesPedigreeFunction	768 non-null	float64
7	Age	768 non-null	int64
8	Outcome	768 non-null	int64

```
dtypes: float64(2), int64(7)
```

```
memory usage: 54.1 KB
```

Data Preparation

- ❑ Describe the dataset to know more about the it

```
df.describe().T
```

	count	mean	std	min	25%	50%	75%	max
Pregnancies	768.0	3.845052	3.369578	0.000	1.00000	3.00000	6.00000	17.00
Glucose	768.0	120.894531	31.972618	0.000	99.00000	117.0000	140.25000	199.00
BloodPressure	768.0	69.105469	19.355807	0.000	62.00000	72.00000	80.00000	122.00
SkinThickness	768.0	20.536458	15.952218	0.000	0.00000	23.00000	32.00000	99.00
Insulin	768.0	79.799479	115.244002	0.000	0.00000	30.50000	127.25000	846.00
BMI	768.0	31.992578	7.884160	0.000	27.30000	32.00000	36.60000	67.10
DiabetesPedigreeFunction	768.0	0.471876	0.331329	0.078	0.24375	0.37250	0.62625	2.42
Age	768.0	33.240885	11.760232	21.000	24.00000	29.00000	41.00000	81.00
Outcome	768.0	0.348958	0.476951	0.000	0.00000	0.00000	1.00000	1.00

Data Preparation

- ❑ Check the number of missing values our dataset has

```
df.isnull().sum()
```

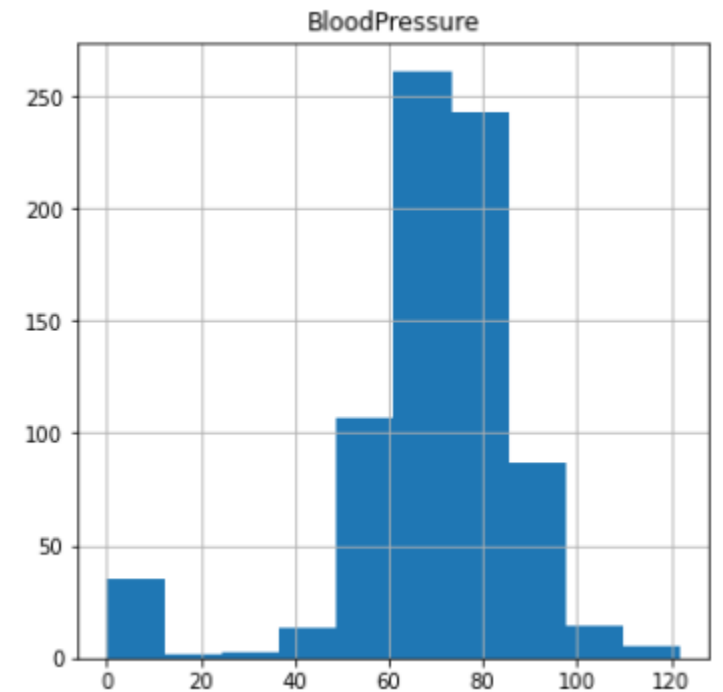

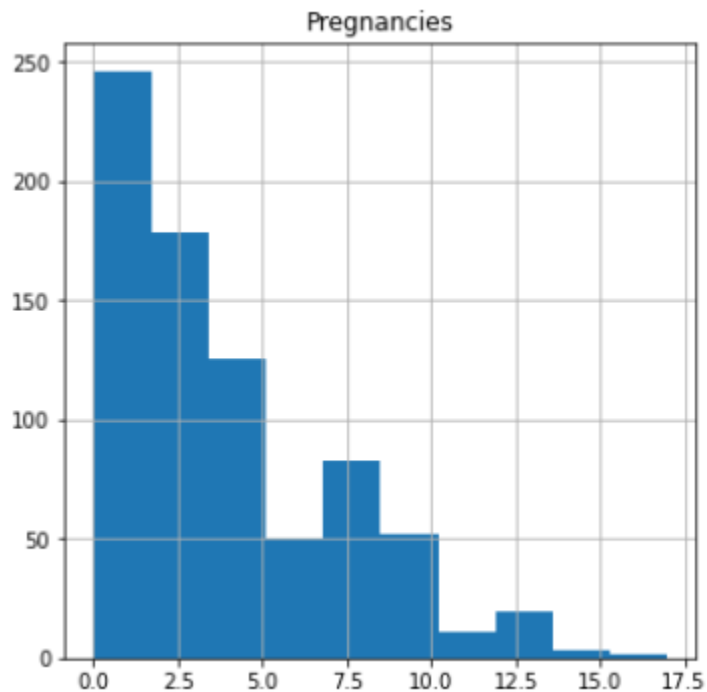
```
Pregnancies      0
Glucose           0
BloodPressure     0
SkinThickness     0
Insulin           0
BMI               0
DiabetesPedigreeFunction  0
Age              0
Outcome          0
dtype: int64
```

We now get is that there are no missing values but that is actually not a true story as in this particular dataset all the missing values were given the outlier(0) as a value which is not good for the authenticity of the dataset. So we will handle the outliers (0).

Data Preparation

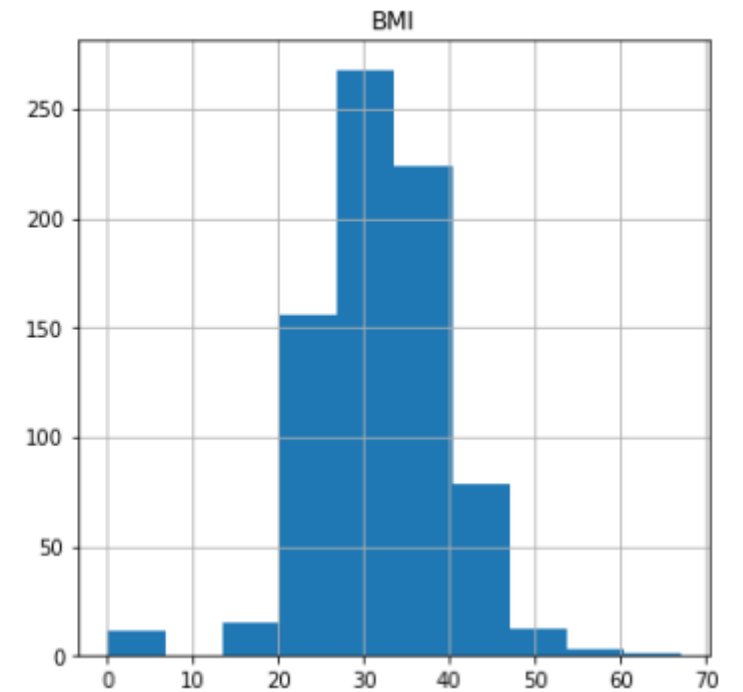
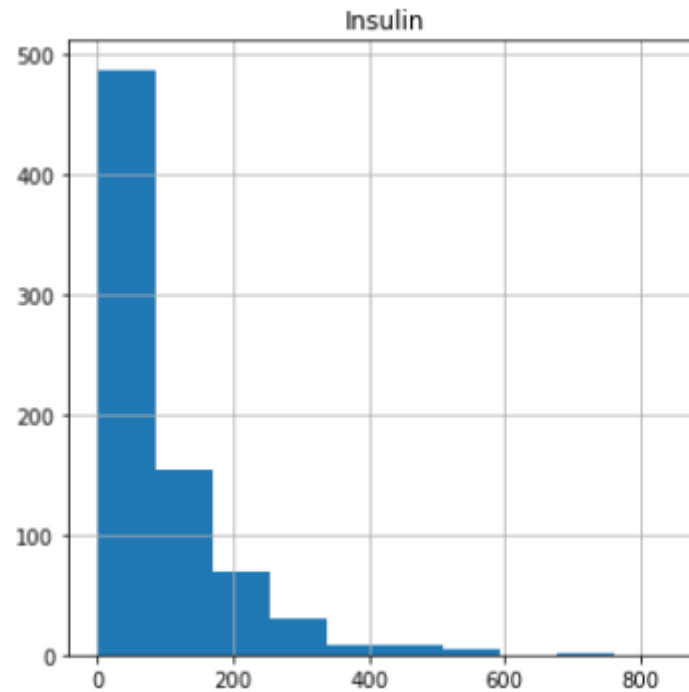
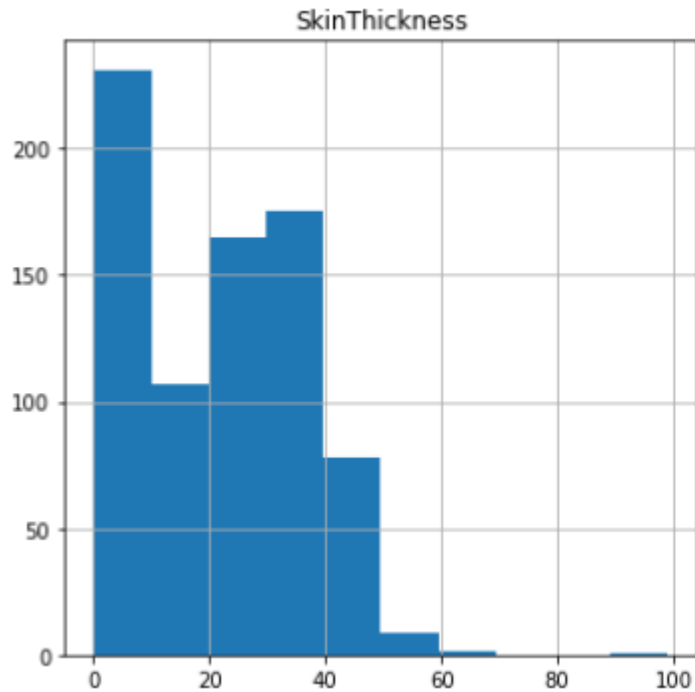
□ Plotting the data distribution plots

```
p = df.hist(figsize = (20,20))
```



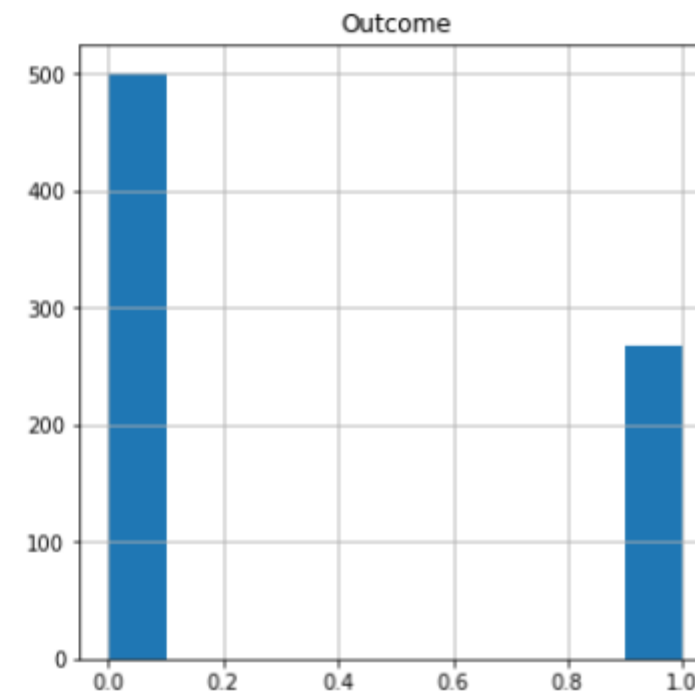
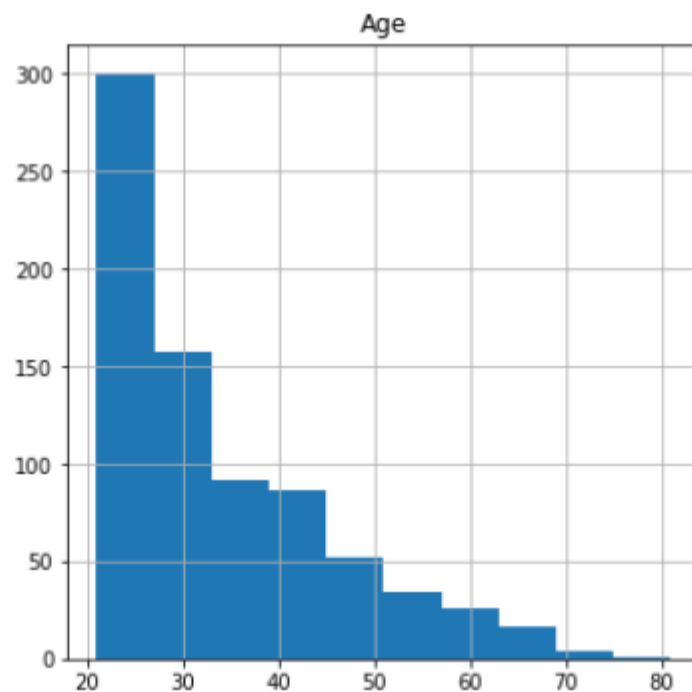
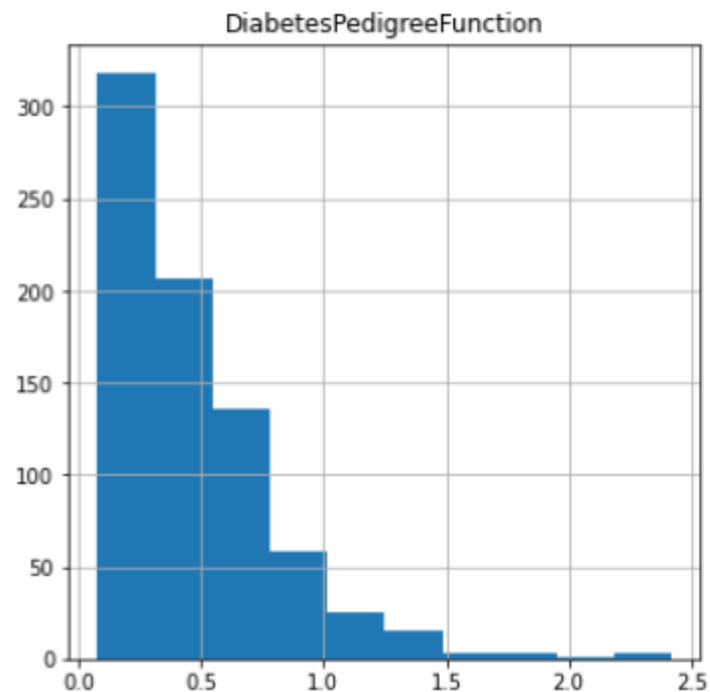
Data Preparation

□ Plotting the data distribution plots



Data Preparation

□ Plotting the data distribution plots



Data Preparation

❑ Handling the outliers

```
# Replace the 0 values with the NAN
df[
    ['Glucose', 'BloodPressure', 'SkinThickness', 'Insulin', 'BMI']
] =
df[
    ['Glucose', 'BloodPressure', 'SkinThickness', 'Insulin', 'BMI']].replace(0, np.NaN)
```

```
# check the number of missing values
df.isnull().sum()
```

Pregnancies	0
Glucose	5
BloodPressure	35
SkinThickness	227
Insulin	374
BMI	11
DiabetesPedigreeFunction	0
Age	0
Outcome	0
dtype: int64	

Data Preparation

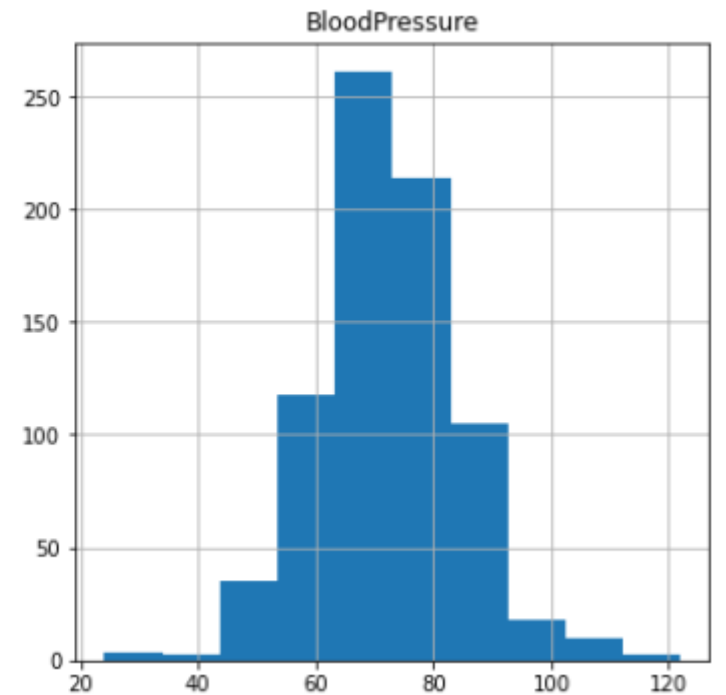
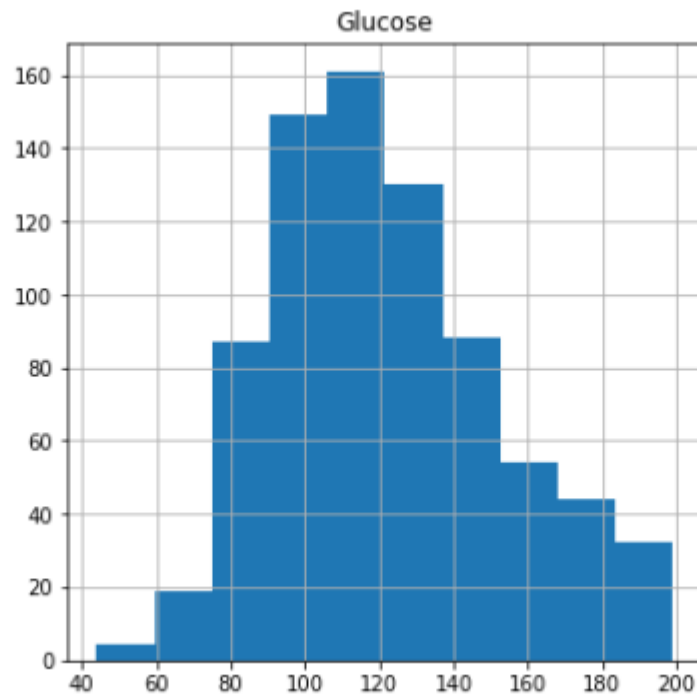

❑ Handling the outliers

```
# Replace the missing values
df['Glucose'].fillna(df['Glucose'].mean(), inplace = True)
df['BloodPressure'].fillna(df['BloodPressure'].mean(), inplace = True)
df['SkinThickness'].fillna(df['SkinThickness'].median(), inplace = True)
df['Insulin'].fillna(df['Insulin'].median(), inplace = True)
df['BMI'].fillna(df['BMI'].median(), inplace = True)
```

Data Preparation

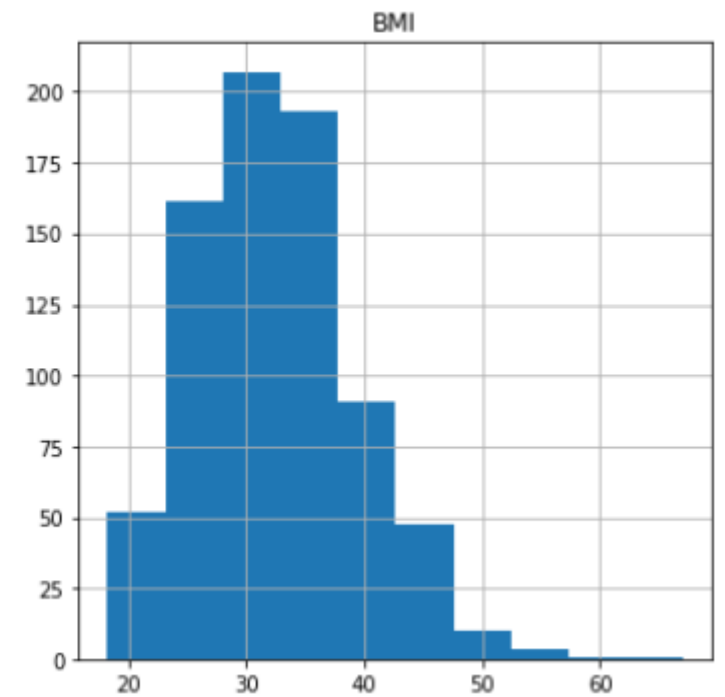
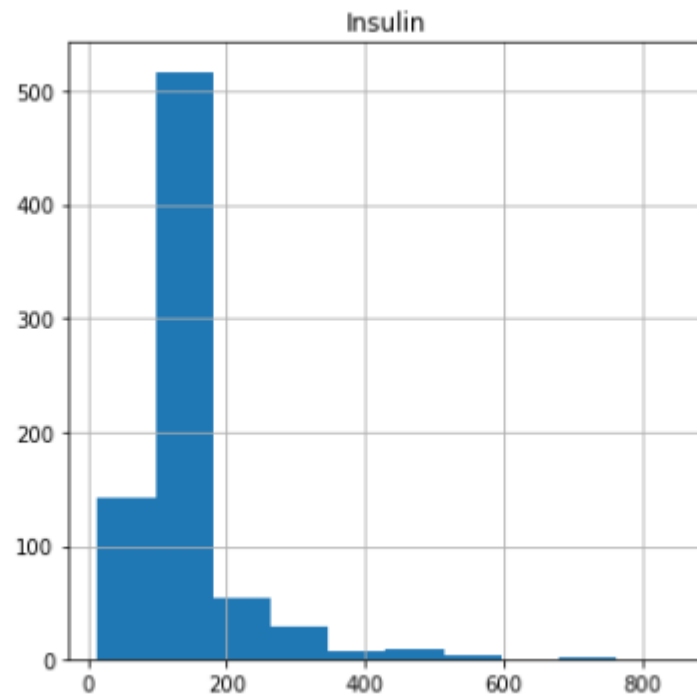
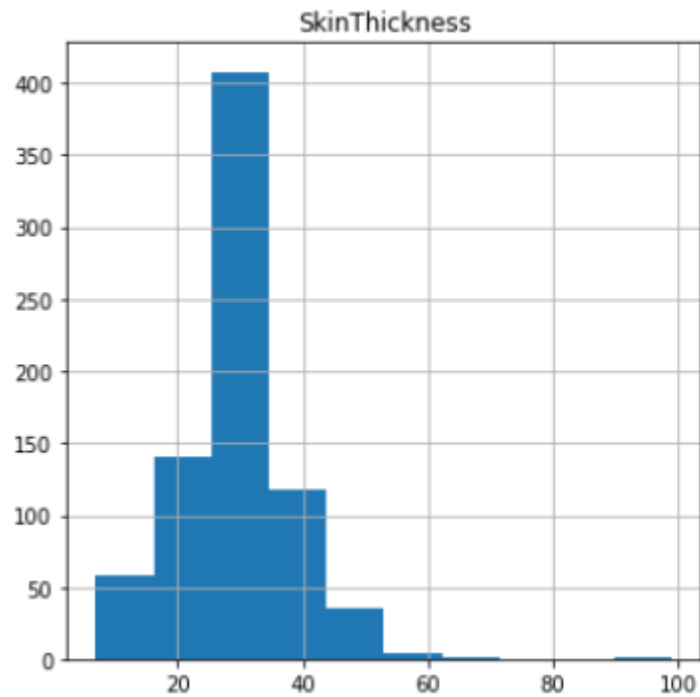
□ Plotting the data distribution plots

```
p = df.hist(figsize = (20,20))
```



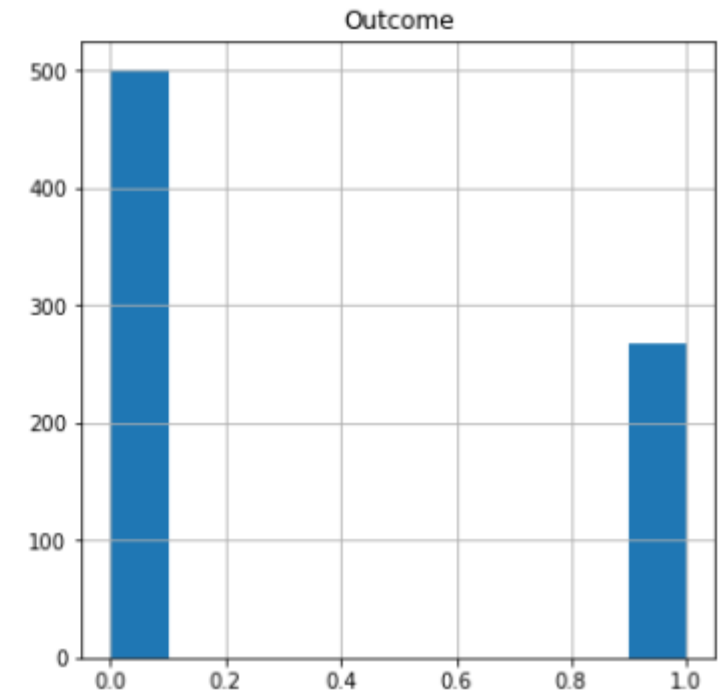
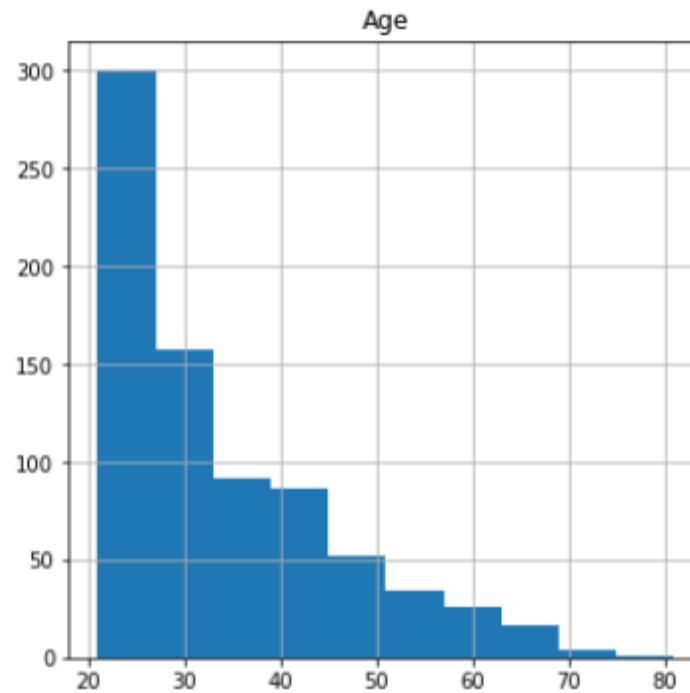
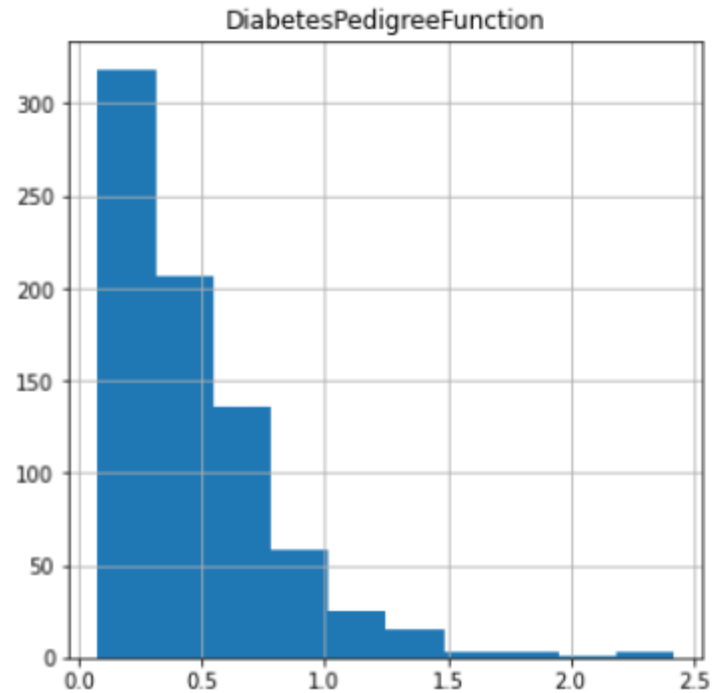
Data Preparation

□ Plotting the data distribution plots



Data Preparation

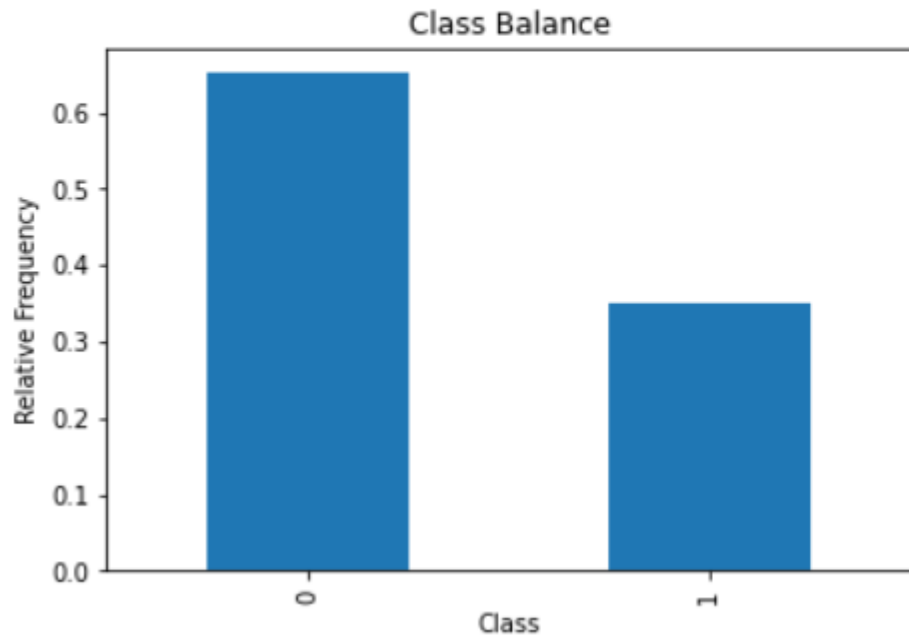
□ Plotting the data distribution plots



EDA

- ❑ Check that how well our outcome column is balanced

```
# Plot value counts of ~"Outcome"~  
df["Outcome"].value_counts(normalize=True).plot(  
    kind="bar", xlabel="Class", ylabel="Relative Frequency", title="Class Balance"  
);
```




EDA

- ❑ Correlation between all the features

```
plt.figure(figsize=(12,10))  
p = sns.heatmap(df.corr(), annot=True, cmap = 'RdYlGn')
```

EDA

❑ Correlation between all the features



Model Building

Algorithms Used in this Project

❑ Decision Tree

```
from sklearn.tree import DecisionTreeClassifier
```

❑ Random Forest

```
from sklearn.ensemble import RandomForestClassifier
```

❑ Support Vector Machine

```
from sklearn.svm import SVC
```

Model Building | Split Data

```
# Create your feature matrix X and target vector y
target = "Outcome"
X = df.drop(columns=target)
y = df[target]
```

```
# Features Scaling
sc = StandardScaler()
X = sc.fit_transform(X)
print(X)
```

```
# Divide the dataset
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42
)

print("X_train shape:", X_train.shape)
print("y_train shape:", y_train.shape)
print("X_test shape:", X_test.shape)
print("y_test shape:", y_test.shape)
```

```
X_train shape: (614, 8)
y_train shape: (614,)
X_test shape: (154, 8)
y_test shape: (154,)
```

Model Building | Decision Tree

```
# Instanciate the model  
DTClassifier = DecisionTreeClassifier()
```

```
# Fit the model  
DTClassifier.fit(X_train, y_train)
```

```
DecisionTreeClassifier()
```

```
# Prediction and accuracy  
predictions = DTClassifier.predict(X_test)  
print("Accuracy Score =", format(accuracy_score(y_test, predictions)))
```

```
Accuracy Score = 0.7662337662337663
```

Model Building | Random Forest

```
# Instanciate the model  
RFClassifier = RandomForestClassifier(n_estimators=600)
```

```
# Fit the model  
RFClassifier.fit(X_train, y_train)  
  
RandomForestClassifier(n_estimators=600)
```

```
# Prediction and accuracy  
predictions = RFClassifier.predict(X_test)  
print("Accuracy_Score =", format(accuracy_score(y_test, predictions)))  
  
Accuracy_Score = 0.7402597402597403
```

Model Building | Support Vector Machine

```
# Instanciate the model  
svClassifier = SVC()
```

```
# Fit the model  
svClassifier.fit(X_train, y_train)
```

```
SVC()
```

```
# Prediction and accuracy  
svc_pred = svc_model.predict(X_test)  
print("Accuracy Score =", format(accuracy_score(y_test, svc_pred)))
```

```
Accuracy Score = 0.7272727272727273
```


Communicate Results

❑ Decision Tree

```
# Display the confusion matrix
cm = confusion_matrix(predictions2, y_test)
cm

array([[79, 16],
       [20, 39]], dtype=int64)
```

❑ Random Forest

```
# Display the confusion matrix
cm = confusion_matrix(predictions, y_test)
cm

array([[78, 19],
       [21, 36]], dtype=int64)
```

❑ Support Vector Machine

```
# Display the confusion matrix
cm = confusion_matrix(svc_pred, y_test)
cm

array([[81, 24],
       [18, 31]], dtype=int64)
```

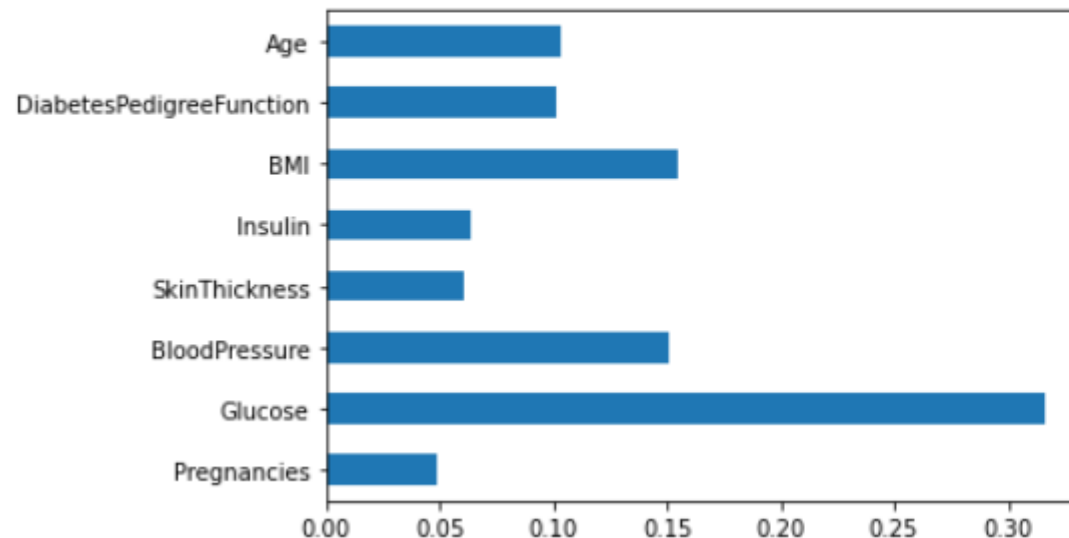
Communicate Results

Therefore Decision Tree is the best model for this prediction since it has an accuracy_score of 0.76

❑ Features importances

```
DTClassifier.feature_importances_  
array([0.04885521, 0.31590208, 0.15113023, 0.0606083 , 0.06355162,  
       0.15468574, 0.1017383 , 0.10352852])  
  
(pd.Series(DTClassifier.feature_importances_, index=df.columns[:8]).plot(kind='barh'))
```

<AxesSubplot:>



Communicate Results

Here from the above graph, it is clearly visible that Glucose as a feature is the most important in this dataset.

❑ Saving the model

```
filename = 'finalized_model.pkl'  
pickle.dump(DTClassifier, open(filename, 'wb'))
```

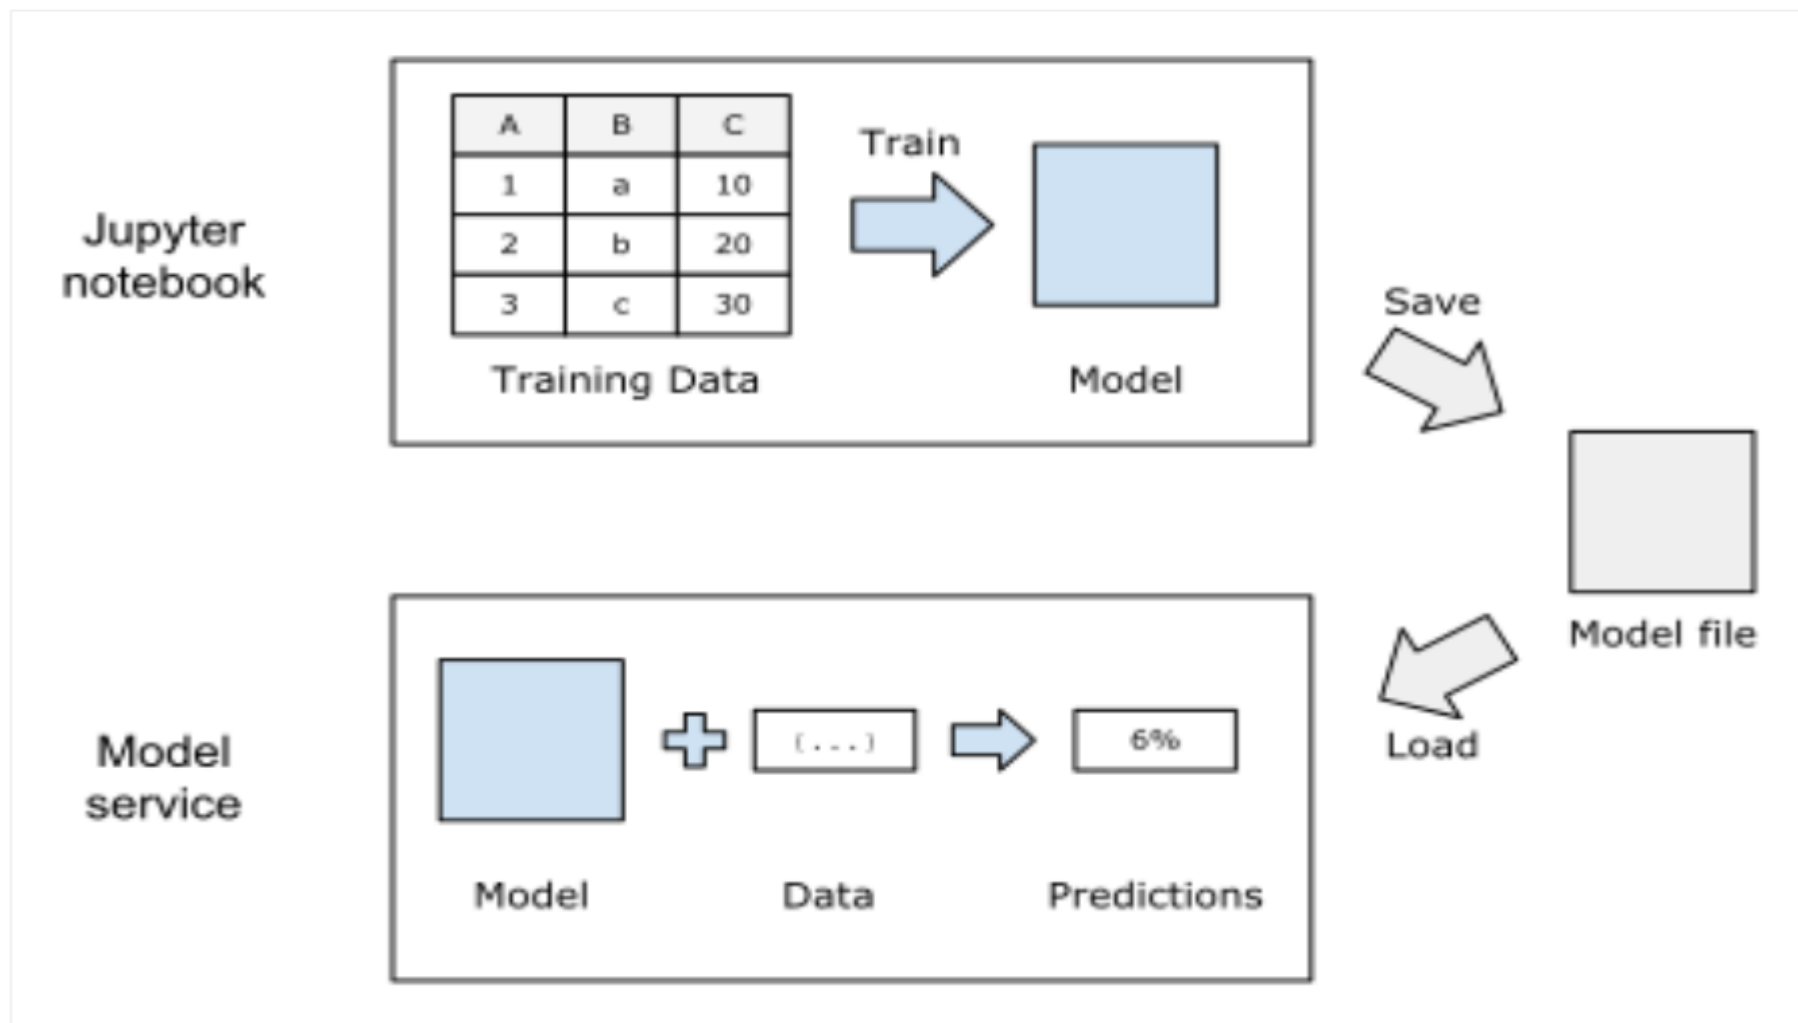
Operationalize

Here from the above graph, it is clearly visible that Glucose as a feature is the most important in this dataset.

❑ Saving the model

```
filename = 'finalized_model.pkl'  
pickle.dump(DTClassifier, open(filename, 'wb'))
```

Operationalize













Operationalize

Technologies Used in this Project



Operationalize

 .git	24/04/2022 17:21	File folder	
 image	24/04/2022 13:30	File folder	
 static	24/04/2022 17:14	File folder	
 templates	24/04/2022 13:30	File folder	
 app	24/04/2022 13:14	JetBrains PyCharm C...	2 KB
 DiabeteClassification	23/04/2022 21:17	IPYNB File	238 KB
 finalized_model.pkl	23/04/2022 21:17	PKL File	16 KB
 Procfile	24/04/2022 12:54	File	1 KB
 README	24/04/2022 13:30	MD File	1 KB
 requirements	24/04/2022 12:54	Text Document	1 KB

Operationalize

```
$ git push -f heroku
Enumerating objects: 12, done.
Counting objects: 100% (12/12), done.
Delta compression using up to 4 threads
Compressing objects: 100% (8/8), done.
Writing objects: 100% (8/8), 1.32 MiB | 556.00 KiB/s, done.
Total 8 (delta 3), reused 0 (delta 0), pack-reused 0
remote: Compressing source files... done.
remote: Building source:
remote:
remote: -----> Building on the Heroku-20 stack
remote: -----> Using buildpack: heroku/python
remote: -----> Python app detected
remote: -----> No Python version was specified. Using the same version as the la
st build: python-3.10.4
remote:       To use a different version, see: https://devcenter.heroku.com/art
icles/python-runtimes
remote: -----> No change in requirements detected, installing from cache
remote: -----> Using cached install of python-3.10.4
remote: -----> Installing pip 22.0.4, setuptools 60.10.0 and wheel 0.37.1
remote: -----> Installing SQLite3
remote: -----> Installing requirements with pip
remote: -----> Discovering process types
remote:       Procfile declares types -> web
remote:
remote: -----> Compressing...
remote:       Done: 191.3M
remote: -----> Launching...
remote:       Released v4
remote:       https://diabetepredicto.herokuapp.com/ deployed to Heroku
remote:
remote: Verifying deploy... done.
```


Diabetes Predictor

Number of Pregnancies eg. 0

Glucose (mg/dL) eg. 80

Blood Pressure (mmHg) eg. 80

Skin Thickness (mm) eg. 20

Insulin Level (IU/mL) eg. 80

Body Mass Index (kg/m²) eg. 23.1

Diabetes Pedigree Function eg. 0.52

Age (years) eg. 34

Predict

Operationalize

Prediction: Wow ! You DON'T have diabetes.



Operationalize

You can play with the web app by follow this link

<https://diabetepredicto.herokuapp.com/>

The GitHub repository is attached below

<https://github.com/Dilane-Kamga/DiabetesPredictionApp.git>

MERCI
Pour votre attention